

Welcome to CSE 361!

Course Overview

CSE 361s: Introduction to Computer Systems

Instructor:


I-Ting Angelina Lee

Overview

- What this course is about
- Logistics and administrivia

What This Course Is About

- Understanding the “complete system:” the big picture



What is System Software?

- Operating System (OS)

- Kernel, scheduler
- File system
- BIOS, drivers

- Utility Software

- compiler, linker, assembler, debugger
- libraries that provide basic facilities:
dynamic memory allocations, file I/O, shell
- other system calls that provide system-level services

Course Topics

- Part I: how an application program interacts with the hardware
 - program and data representation, memory hierarchy, and how these things relating to the performance of your program.
- Part II: how an application program interacts with system software and how system software serve the application and the hardware
 - linking, process, exceptional control flow, virtual memory
- Part III: how programs interact with each other
 - processes
 - threading and synchronization

A Very Loose Description of Course

*“Everything you need to know
about system software and
hardware to write (or debug)
[fast/correct/secure] software.”*

Course Theme:

Abstraction Is Good But Don't Forget Reality

- **Most CS courses emphasize abstraction**
 - Abstract data types
 - Asymptotic analysis
- **These abstractions have limits**
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- **Useful outcomes from taking this class**
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
 - Prepare for later “systems” classes in CS & CoE
 - Compilers, Operating Systems, Wireless Networks, Network Security, System Security, Computer Architecture, Embedded Systems, etc.

The Five Great Realities

Code Puzzle #1

```
x = ...    // initialize to some numeric value  
assert((x*x) >= 0);
```

- Q: Does this assertion succeed always?
- A: depending on its data type!

Code Puzzle #2

```
float f1 = ...;  
float f2 = ...;  
float f3 = ...;  
  
assert( (f1+f2)+f3 == f1+(f2+f3) );
```

- Q: Does this assertion succeed always?
- A: depending on the values of floats!

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$

■ Float's: Yes!

■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

■ Unsigned & Signed Int's: Yes!

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Code Puzzle #3:

```
typedef struct {
    int a[2];
    long l;
} struct_t;

long fun(int i) {
    volatile struct_t s;
    s.l = 999;
    s.a[i] = 1000; /* Possibly out of bounds */
    return s.l;
}
```

■ Q: What does function 'fun' return?

■ A: depending on its input!


fun(0)	→	999
fun(1)	→	999
fun(2)	→	1000
fun(3)	→	4294967296999
fun(4)	→	Segmentation fault

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    long l;  
} struct_t;
```

fun(0)	→	999
fun(1)	→	999
fun(2)	→	1000
fun(3)	→	4294967296999
fun(4)	→	Segmentation fault

Explanation:



Result is system specific!

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

- **Memory referencing bugs especially pernicious**

- Effects are distant in both time and space

- **Memory is not unbounded**

- It must be allocated and managed
- Many applications are memory dominated

- **Memory performance is not uniform**

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Bug may manifest much later down the execution
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. gdb, Valgrind)

Code Puzzle #4

```
void copyij(int n,  
            int src[n][n],  
            int dst[n][n])  
{  
    int i,j;  
    for (i = 0; i < n; i++)  
        for (j = 0; j < n; j++)  
            dst[i][j] = src[i][j];  
}
```

```
void copyji(int n,  
            int src[n][n],  
            int dst[n][n])  
{  
    int i,j;  
    for (j = 0; j < n; j++)  
        for (i = 0; i < n; i++)  
            dst[i][j] = src[i][j];  
}
```



- Q: Do these two programs behave similarly?
- A: Their performance differ greatly!

2.7ms*

45ms*

17 times performance difference!

- Hierarchical memory organization
- Performance depends on access patterns

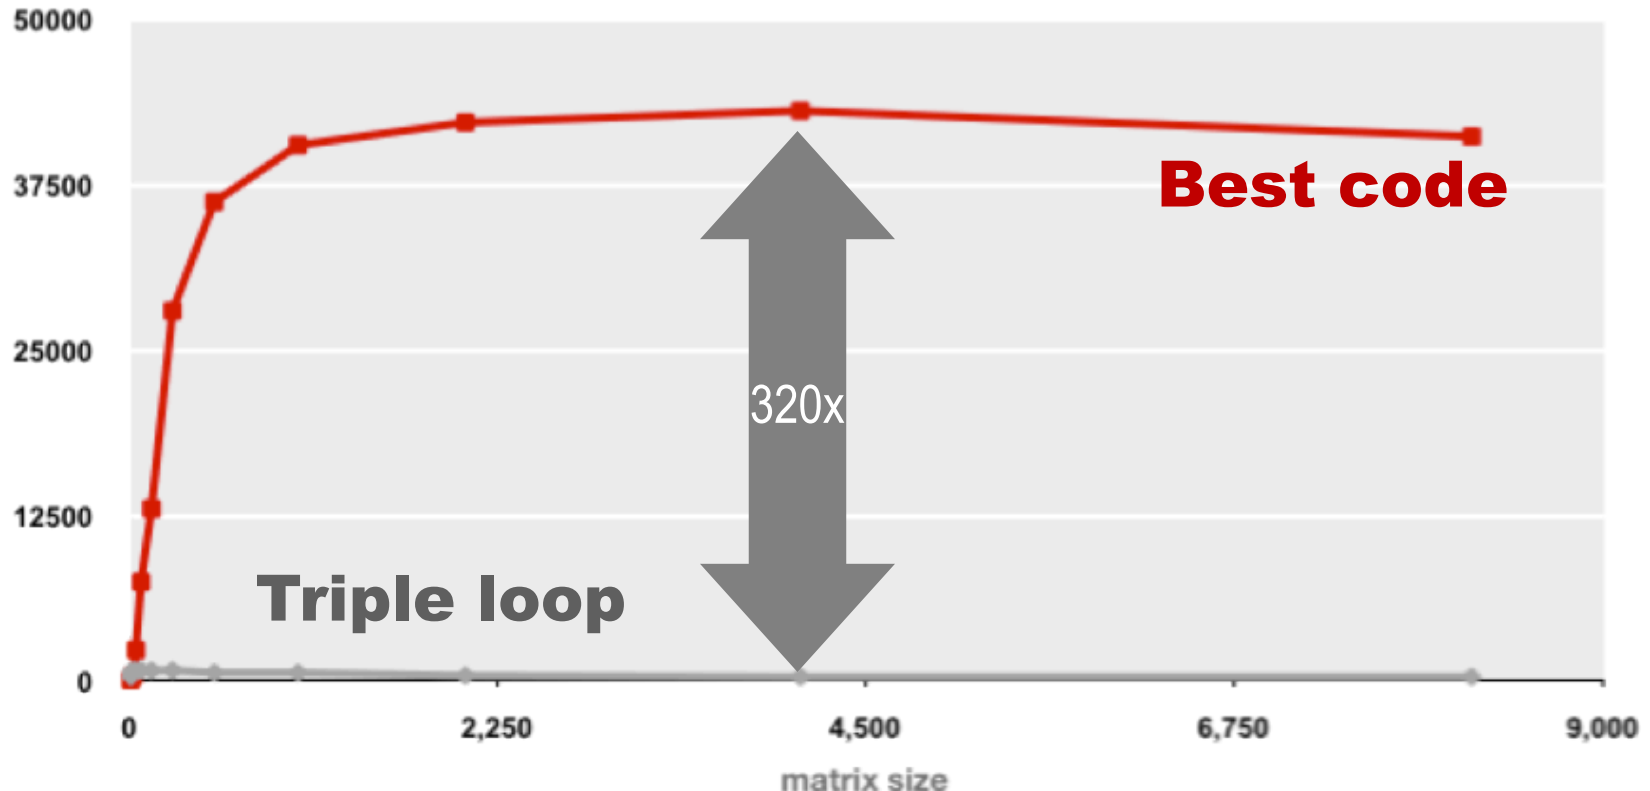
*on a 1.2 GHz Intel Xeon Processor

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Example Matrix Multiplication

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)
Gflop/s




- Standard desktop computer, vendor compiler, using optimization flags
- Both implementations have **exactly** the same operations count ($2n^3$)
- What is going on?

MMM Plot: Analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz

Gflop/s



- Reason for 20x: Blocking or tiling, loop unrolling, array scalarization, instruction scheduling, search to find best choice
- Reason for another 16x: parallelism

Great Reality #5:

Parallelism / Concurrency Matters

■ Instruction-Level Parallelism:

- Hardware schedules instructions out of order
- Vectorization unit

■ The notion of processes:

- The OS manages the sharing of resources and maintains the illusion that your program is the only one running.

■ Parallelism among threads:

- Utilize the underlying hardware to finish a computation faster / process more data in shorter amount of time.
- Manage sharing of resources at the user level.
- Complex thread interaction can lead to programming errors that are difficult to detect / debug.
- More complex performance issues.

Course Perspective

- Most Systems Courses are Builder-Centric
 - Computer Architecture
 - Design pipelined processor in Verilog
 - Operating Systems
 - Implement sample portions of operating system
 - Compilers
 - Write compiler for simple language
 - Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is **Programmer-Centric**

- Show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to:
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers, multithreading
- Not just a course for dedicated hackers
 - *It a course for everyone who wants to be a power programmer!*

Know How and When to Use Your Tools

- linux: command line
 - git: (command line)
 - text editor for coding: (vi / vim / emacs)
 - compilation: makefile, gcc
 - debugging: gdb, valgrind
-
- Exercises to familiarize yourself with basic linux tools and intro to C will be released sometime this week.

Logistics and Administrvia

CSE 361 Course Staff

- Instructor: I-Ting Angelina Lee
 - Office hours: after class on Mon/Wed, 5:20-6:20pm or by appointment

- TAs:

Zihao Chen,
Noah Goldstein (head TA),
Yuchen Han,
Clayton Knittel,

Michael Liu,
Connor Monahan,
Richard Wu,
Charles Yang,
Yiheng Yao

CSE 361 Online

- Webpage:

- <https://www.cse.wustl.edu/~angelee/cse361>

- Syllabus, textbook, etc.
 - Course schedule (lab due dates, exam dates)
 - Lab hours (TBD)



- Piazza:

- <https://piazza.com/wustl/fall2019/cse361/home>

- Q & A (**do not use email**)
 - Answer someone else's question!
 - Lecture notes, lab assignments, sample exams

Textbooks

- Randal E. Bryant and David R. O'Hallaron,
 - “Computer Systems: A Programmer’s Perspective, Third Edition” (CS:APP3e), Pearson, 2015
 - <http://csapp.cs.cmu.edu>
 - This book really matters for the course!
 - How to solve labs
 - Practice problems typical of exam problems



- Brian Kernighan and Dennis Ritchie,
 - “The C Programming Language, Second Edition”, Prentice Hall, 1988

Course Components

■ Lectures

- Lopata 101
- Higher level concepts, practice exercises (participate!)

■ Lab Hours

- starting 2nd week of class; check course page for hours

■ Labs (5) + one extra credit lab

- The heart of the course
- 2-3 weeks each except for Lab 5 (~3.5 weeks)
- Provide in-depth understanding of an aspect of systems

■ Exams (midterm + final)

- Test your understanding of concepts & principles
- Crib sheets allowed.

Grading

- Exams (45%):
 - 20% (midterm: Wednesday, Oct 16, in class)
 - 25% (final: Friday, Dec 13, 6-8pm)
 - If you have special arrangement with Cornerstone, or it conflicts with another class, please make PRIOR arrangements with me
- Labs (55%):
 - Lab grades are weighted by difficulty:
Lab 1: 10%, Lab 2-4: 11%, Lab 5 12%
 - Extra credit will be worth 1.5% of your overall grade.
 - You get 2-day extension per lab automatically
- Exercises:
 - Won't be graded; for your benefit only
- Grades will be posted through Canvas.

Facilities and Getting Help

■ Linuxlab:

- ssh into shell.cec.wustl.edu and do `qlogin`
- Linuxlab desktop sessions:
<https://linuxlab.seas.wustl.edu/equeue/>

■ You can also use:

- Your own Linux box
- Terminal or X11 on your Mac
- Cygwin or or secure shell from your Windows machine
- But these options won't be supported by class staff

Facilities and Getting Help (Cont)

- Go to Lab Hours for help on projects.
- Come to office hour for lecture materiel (though I can answer questions on labs, too).
- Do the exercises!
- We will hold only limited hours during the 2nd week and official lab hours will start on the 3rd week.
 - Will post on Piazza / webpage for 2nd week hours
 - Will update on Piazza / webpage for 3rd week hours right before the week starts

Academic Integrity

- What is cheating?
 - **Searching online for possible solution**
 - **Copying code** from previous terms or from the web
 - Only allowed to use code **we** supply
 - **Sharing code**: by copying, retyping, looking at, or supplying a file
 - **Coaching**: helping your friend to write a lab, line by line
 - Warning: once you see a solution, it is **very hard to un-see**
- What is NOT cheating?
 - Explaining how to use systems or tools
 - Helping others with high-level design issues

TO DO

- Make sure that you are signed up for Piazza
- Read Chap 1 of the textbook
- If you are not already familiar with Linux command lines / C, plan to work through exercise 1 before starting lab 1.
 - Exercise 1 will be posted soon on Piazza

Welcome
and Enjoy!