

Announcement

- **Limited Lab hours this week.**
 - See course webpage for time / location.
- **We will add more lab hours this coming week.**
- **GitHub repositories set up! Will update lab 1 and exercise 1 by tomorrow (will announce on Piazza).**

Bits, Bytes, and Integers (Cont.)

B&O Readings: 2.2-2.3

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Implement a pop_count function, cont

- How do you implement pop_count, that counts the number of bits set in a 4 byte memory?

ex: pop_count(0x000000FF) = 8

```
int pop_count(unsigned int x) {  
    int count = 0;  
    for(; x != 0; x &= ~(x & (-x))) {  
        count++;  
    }  
    return count;  
}
```

The expression $(x \& (-x))$ computes a mask with a single 1 set at the least-significant position where x has a bit 1 set .



Implement a pop_count function

- How do you implement pop_count, that counts the number of bits set in a 4 byte memory?

ex: pop_count(0x000000FF) = 8

```
#define MASK 0xF;
int count_arr[16] =
    {0,1,1,2,1,2,2,3,1,2,2,3,2,3,3,4}

int pop_count(unsigned int x) {
    int count = 0;
    while(x != 0) {
        count += count_arr[x&MASK];
        x = x >> 4;
    }
    return count;
}
```

We can check 4 bits at a time!



Recap: What We Learned Thus Far

- The right shift behaves differently depending on whether an expression is signed versus unsigned.

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition
 - How to detect overflow
 - Multiplication / Division
 - How to do these operations with shifts

Casting Between Signed vs. Unsigned in C

■ Constants

- *By default are considered to be signed integers*
- Unsigned if have “U” as suffix: `0U`, `4294967259U`

■ Casting

- Explicit casting between signed & unsigned same as U2T and T2U

```
int tx, ty;  
unsigned ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

Rule of Thumb: Keep bit representations and reinterpret!

- Implicit casting also occurs via assignments and procedure calls

```
tx = ux;  
uy = ty;
```

Use -Werror and -Wall compiler flag to catch this!

Casting Surprises

■ Expression Evaluation

- If there is a mix of unsigned and signed in single expression,
signed values implicitly cast to unsigned
- Including comparison operations <, >, ==, <=, >=

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- Similar to code found in FreeBSD's implementation of `getpeername`
- There are legions of smart people trying to find vulnerabilities in programs

Typical Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

Malicious Usage

```
/* Declaration of library function memcpy */  
void *memcpy(void *dest, void *src, size_t n);
```

```
/* Kernel memory region holding user-accessible data */  
#define KSIZE 1024  
char kbuf[KSIZE];  
  
/* Copy at most maxlen bytes from kernel region to user buffer */  
int copy_from_kernel(void *user_dest, int maxlen) {  
    /* Byte count len is minimum of buffer size and maxlen */  
    int len = KSIZE < maxlen ? KSIZE : maxlen;  
    memcpy(user_dest, kbuf, len);  
    return len;  
}
```

```
#define MSIZE 528  
  
void getstuff() {  
    char mybuf[MSIZE];  
    copy_from_kernel(mybuf, -MSIZE);  
    . . .  
}
```

You end up allowing the user to copy lots of data from kernel!

Recap: What We Learned About Casting

- **C allows one to cast from signed to unsigned and vice versa:**
 - Bit pattern is maintained
 - But reinterpreted
 - Can have unexpected effects: adding or subtracting 2^w
- **When an expression contains both signed and unsigned, it's implicitly treated as unsigned.**
- **Understanding these quirks in C allows you to write correct and secure code.**

When Should I Use Unsigned?

■ *Don't Use Just Because the Number are Nonnegative*

- Easy to make mistakes

```
unsigned i;  
for (i = cnt-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- Can be very subtle

```
#define DELTA sizeof(int)  
int i;  
for (i = CNT; i-DELTA >= 0; i-= DELTA)  
    . . .
```

■ *Do Use When Using Bits to Represent Sets*

- Logical right shift, no sign extension

Expression Evaluation Puzzles

■ Assuming int type (32 bits)

■ Constant ₁	Constant ₂	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483647-1	>	signed
2147483647U	-2147483647-1	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed

TMIN = -2147483647-1 (0x80000000)

TMAX = 2147483647 (0x7FFFFFFF)



Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition
 - How to detect overflow
 - Multiplication / Division
 - How to do these operations with shifts
- **Summary**

Extension

- When operating with types of different width, C automatically performs extension
- Converting from smaller to larger integer data type
 - Given w -bit integer X
 - Convert it to $w+k$ -bit integer X' **with same value**
- Two different kinds of extension:
 - zero extension: used for unsigned data types
(similar: \gg uses logical right shift for unsigned values)
 - sign extension: used for signed data types
(similar: \gg uses arithmetic right shift for signed values)

```
unsigned short sx = 361;  
unsigned int x = sx; /* use zero extension */  
  
short sy = -361;  
int y = sy; /* use sign extension */
```


Zero Extension for Unsigned type

■ What It Does:


- Given w -bit unsigned integer X
- Convert it to $w+k$ -bit unsigned integer X' with same value

■ Rule:

- Prepend k bits of 0:
- $X' = \underbrace{0, \dots, 0}_{k \text{ copies}}, x_{w-1}, x_{w-2}, \dots, x_0$

$\underbrace{\hspace{2cm}}$
 k copies

- Easy to see that the extension preserves the value: added bits don't contribute any weight to the value.



Sign Extension

■ What It Does:


- Given w -bit signed integer X
- Convert it to $w+k$ -bit unsigned integer X' with same value

■ Rule:

- Make k copies of the sign bit:
- $X' = \underbrace{X_{w-1}, \dots, X_{w-1}}_{k \text{ copies of MSB}}, X_{w-1}, X_{w-2}, \dots, X_0$

k copies of MSB

- The extension preserves the value!




Sign Extension Preserves the Value

- **X is positive:**

- easy to see: 0 bits don't add weight

- **X is negative:**



We can show that sign extension does not change the value by inducting on k .


Truncation

■ Task:

- Given w -bit signed integer X
- Convert it to k -bit integer X' with same value (maybe...)

■ Rule:

- Drop high-order $w-k$ bits



■ Effect:

- Can change the value of X (overflow)
- Unsigned: mathematical mod on X
- Signed: reinterpret the bits (add -2^k if the most-significant bit is 1)

Code Puzzle

- What is the output of the following code?
Assume that int is 32 bits, short is 16 bits, and the representation is two's complement.

```
unsigned short y = 0xFFFF;  
int x = y;  
printf("%x", x); /* print in hexadecimal */
```

A) 0xFFFF B) 0xFFFFFFFF C) None of the above

Answer: A) 0xFFFF

DO NOT write code like this!



Recap: What We Learned Thus Far

- **Expanding (e.g., short int to int)**
 - Unsigned: zeros added
 - Signed: sign extension
 - Both yield expected result
- **Truncating (e.g., unsigned to unsigned short)**
 - Unsigned/signed: bits are truncated
 - Unsigned: keep the last k bits; like a mod operation
 - Signed: keep the last k bits, and reinterpret the bits as signed
 - For small numbers yields expected behavior; for large number, can change the value.
- **Know how an expression with mixed types is treated but avoid writing code like that to avoid unintended behavior!**

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition
 - How to detect overflow
 - Multiplication / Division
 - How to do these operations with shifts
- **Summary**

Integer Addition

- **Rule of thumb 1: Do the normal binary operations assuming enough bits, and chop off the extra bits that cannot fit.**
- **Rule of thumb 2: The hardware does not care about whether the variables are signed versus unsigned; the operations are the same for both.**

Unsigned Addition

$$0 \leq u, v \leq 2^w - 1$$

$$0 \leq u + v \leq 2^{w+1} - 2$$

Operands: w bits




+ v



True Sum: $w+1$ bits

true $u + v$



Discard Carry: w bits

$\text{UAdd}_w(u, v)$



■ Standard Addition Function

- Ignores carry output

How to Detect Overflow in UAdd?

Hint: try performing UAdd with 4-bit values.

- **What's the range of value that a 4-bit variable can represent?**
- **How does one interpret the result with overflow?**

Detecting Overflow in Unsigned Addition

■ When overflow:

- Assume w -bit operands
- If overflow, true sum $\geq 2^w$ but can overflow by 1 bit only
- $\text{Uadd}(u,v) = \text{true sum mod } 2^w$

$$= u + v - 2^w$$

$$= u + \underbrace{(v - 2^w)}_{< 0} \quad \text{or} \quad v + \underbrace{(u - 2^w)}_{< 0}$$

- To detect overflow in UAdd, check if $\text{UAdd}(u,v) < u$ or $< v$

Two's Complement Addition


$$-2^{w-1} \leq u, v \leq 2^{w-1}-1$$

$$-2^w \leq u + v \leq 2^w - 2$$

Operands: w bits



True Sum: $w+1$ bits



Discard Carry: w bits



■ TAdd and UAdd have Identical Bit-Level Behavior

```
int s, t, u, v;  
... /* initialize their values */  
s = (int) ((unsigned) u + (unsigned) v);  
t = u + v;  
assert(s == t); /* always true! */
```

Same bit pattern, different interpretation for sign vs. unsigned.


How to Detect Overflow in TAdd?

Hint: try performing TAdd with 4-bit signed values.


- **What's the range of value that a 4-bit signed variable can represent?**
- **Try adding two largest values together**
 - $0111 + 0111 = 1110$ (-2)
 - Overflow to the MSB
- **Try adding two smallest values together**
 - $1000 + 1000 = 10000 \rightarrow 0000$ (0)
 - Overflow to a bit that gets truncated
 - The MSB must be 0



TAdd Overflow




TAdd Overflow



■ Positive overflow:

- Adding two positive values, where $(u + v) > 2^{w-1}-1$ (TMax)
- w^{th} bit contributes to true sum weight of 2^{w-1} but to TAdd sum -2^{w-1}
- TAdd sum = true sum - 2^w (negative)
 $\underbrace{\hspace{1cm}}_{< (2^w-1)}$

TAdd Overflow



■ Negative overflow:

- Adding two negative values, where $(u + v) < -2^{w-1}$ (TMin)
- Missing the carry $(w+1)^{th}$ bit (which would have contributed weight -2^w)
- TAdd sum = true sum + 2^w (positive)

$$\underbrace{\hspace{1cm}}_{< (2^w - 1)}$$

Detecting Overflow in Two's Complement Addition

■ Positive overflow:

- the carry-bit overflow into the most-significant bit (MSB)
- $\text{true sum} \geq 2^{w-1} - 1$, MSB contributes negative weight instead of positive
- $\text{TAdd}(u,v) = (u+v) - 2^w$ (which results a negative value)

■ Negative overflow:

- If $\text{true sum} < -2^{w-1}$, the carry-bit overflow into the bit that got truncated
- The $(w+1)^{\text{th}}$ bit would have contributed -2^w weight
- $\text{TAdd}(u,v) = (u+v) + 2^w$ (which results positive value)

■ To detect overflow in Tadd, check if signs of input operands and output differ.

Recap: What We Learned Thus Far

- For w -bit operands, need $w+1$ bits for true sum
- For fixed-width integer addition, do the usual addition and truncate extra bits
- For unsigned addition
 - Check for overflow by checking if the output is smaller than either input
- For two's complement addition
 - Can only overflow when both operands have the same sign
 - Check for overflow by checking if the signs of inputs and output differ
- Knowing when overflow might occur and how to check for them enables you to write correct code.

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition
 - How to detect overflow
 - Multiplication / Division
 - How to do these operations with shifts
- **Summary**

Integer Multiplication

- **Rule of thumb 1: Do the normal binary operations assuming enough bits, and chop off the extra bits that cannot fit.**
- **Rule of thumb 2: The hardware does not care about whether the variables are signed versus unsigned; the operations are the same for both.**
- **Same Rules as Integer Addition!**

Unsigned Multiplication in C

Operands: w bits



True Product: $2*w$ bits



Discard w bits: w bits



■ Standard Multiplication Function

- Ignores high order w bits

■ Implements Modular Arithmetic

$$\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$$

Signed Multiplication in C

Operands: w bits

u 

$*$ v 

True Product: $2*w$ bits

$u \cdot v$ 

Discard w bits: w bits

$\text{TMult}_w(u, v)$ 

■ Standard Multiplication Function

- Ignores high order w bits
- Same treatment as unsigned, just reinterpret the bits

Power-of-2 Multiply with Shift


■ Operation

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned

Operands: w bits

True Product: $w+k$ bits

Discard k bits: w bits



■ Examples

- $u \ll 3 \quad == \quad u * 2^3$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Power-of-2 Multiply with Shift Example

- Q: How do you computing $X \cdot 6$ by using left shift?



Power-of-2 Multiply with Shift Example

- Q: How do you computing $X \cdot 6$ by using left shift?

$6 = 0\dots0110$ (in binary)

$$\begin{aligned}x \cdot 6 &= x \cdot (2^2 + 2^1) \\ &= x \ll 2 + x \ll 1\end{aligned}$$

Or, equivalently (assuming no overflow),


$$\begin{aligned}x \cdot 6 &= x \cdot (2^3 - 2^1) \\ &= x \ll 3 - x \ll 1\end{aligned}$$



Unsigned Power-of-2 Divide with Shift

■ Quotient of Unsigned by Power of 2

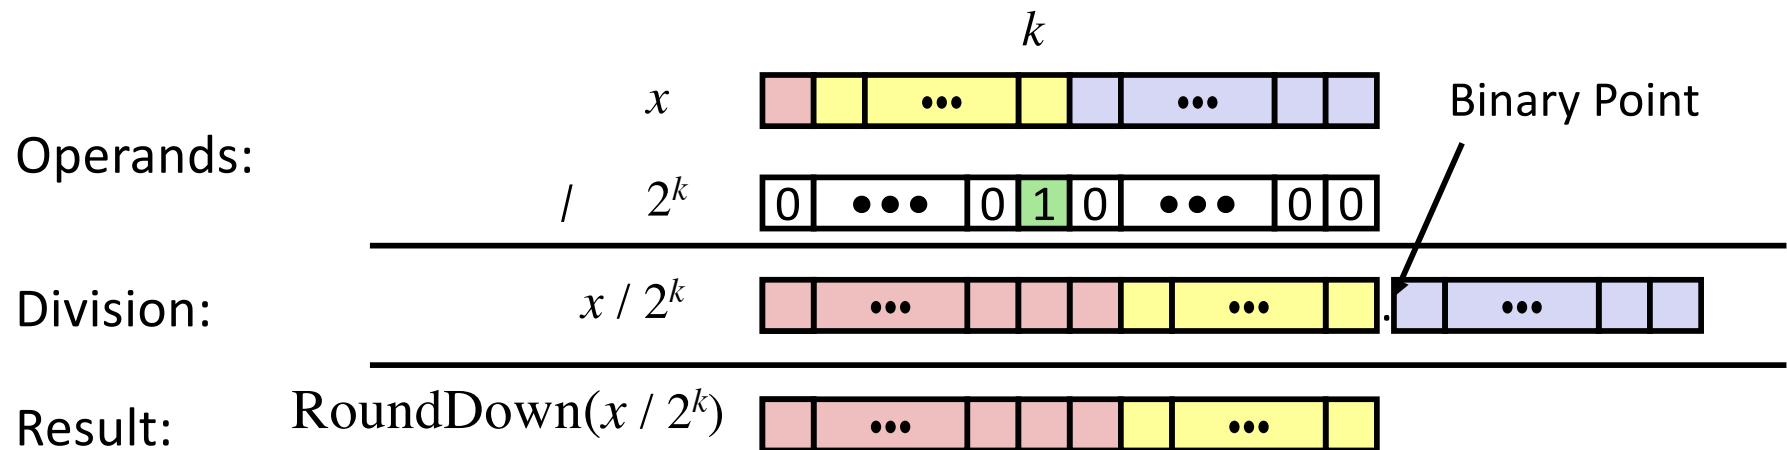
- $u \gg k$ gives $\lfloor u / 2^k \rfloor$
- Uses logical shift



Signed Power-of-2 Divide with Shift

■ Quotient of Signed by Power of 2

- $x \ggg k$ gives $\lfloor x / 2^k \rfloor$
- Uses arithmetic shift
- Rounds wrong direction when $x < 0$ (normal division rounds towards 0)



Recap: What We Learned Thus Far

- **Integer Multiplication:**
 - For w -bit operands, need $2w$ bits for true product
 - Signed vs unsigned values are treated the same way
- **Multiplication by 2^k can be done with left shift**
- **Division by 2^k can be done with right shift**
 - Unsigned: logical shift
 - Signed: arithmetic shift
 - Watch out: for negative numbers, round away from zero!
- **Use $2w$ -bit integer data type for w -bit multiplications to avoid overflow.**
- **Whenever possible, use shifts for multiplication / division**

Today: Bits, Bytes, and Integers


- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition
 - How to detect overflow
 - Multiplication / Division
 - How to do these operations with shifts
- **Summary**

Code Security Example


■ SUN XDR library

- Widely used library for transferring data between machines

```
void* copy_elements(void *ele_src[], int ele_cnt, size_t ele_size);
```



`malloc(ele_cnt * ele_size)`



"In this array I've got pointers to 4 chunks of data. I'd like you to allocate a block of memory and store all these chunks in that block."

XDR Code

```
void* copy_elements(void *ele_src[], int ele_cnt, size_t ele_size) {  
    /*  
     * Allocate buffer for ele_cnt objects, each of ele_size bytes  
     * and copy from locations designated by ele_src  
     */  
    void *result = malloc(ele_cnt * ele_size);  
    if (result == NULL)  
        /* malloc failed */  
        return NULL;  
    void *next = result;  
    int i;  
    for (i = 0; i < ele_cnt; i++) {  
        /* Copy object i to destination */  
        memcpy(next, ele_src[i], ele_size);  
        /* Move pointer to next memory region */  
        next += ele_size;  
    }  
    return result;  
}
```

XDR Vulnerability on 32-bit System

```
malloc(ele_cnt * ele_size)
```

■ What if:

- `ele_cnt` = $2^{20} + 1$
- `ele_size` = 4096 = 2^{12}
- Allocation = ??



XDR Vulnerability on 32-bit System

```
malloc(ele_cnt * ele_size)
```

■ What if:

- $\text{ele_cnt} = 2^{20} + 1$

- $\text{ele_size} = 4096 = 2^{12}$

- $\text{Allocation} = 2^{12} (2^{20} + 1) = 2^{32} + 2^{12}$
 $= 4096 \text{ bytes}$ (just shy of the 4.3 billion needed)
You're going to overwrite a lot of data in your program.

Integer C Puzzles

Initialization

```
int x = foo();  
int y = bar();  
unsigned ux = x;  
unsigned uy = y;
```

1. $x < 0 \implies ((x*2) < 0)$
2. $x > 0 \implies ((x*2) > 0)$
3. $ux \geq 0$
4. $x \& 7 == 7 \implies (x \ll 30) < 0$
5. $ux > -1$
6. $x > y \implies -x < -y$
7. $x * x \geq 0$
8. $x > 0 \&\& y > 0 \implies x + y > 0$
9. $x \geq 0 \implies -x \leq 0$
10. $x \leq 0 \implies -x \geq 0$
11. $(x|-x) \gg 31 == -1$
12. $ux \gg 3 == ux/8$
13. $x \gg 3 == x/8$
14. $x \& (x-1) != 0$

Integer C Puzzles Answers

1. No (TMin can overflow)
2. No (0x0100000...0 shift becomes Tmin)
3. Yes (same bits reinterpreted)
4. Yes (the 3 LSB are all 1's, after shift, 0x1100...0)
5. No (actually it's never true, since -1 is evaluated as unsigned)
6. No (think of TMin ... the range of signed value is asymmetric)
7. No (overflow)

Integer C Puzzles Answers, Cont

- 8. No (overflow 1. No (TMin can overflow)
- 9. Yes (the range is asymmetric, but for every positive value representable, its negative value is also within the range)
- 10. No (again, TMin)
- 11. No (counter example: 0)
- 12. Yes (since it always rounds towards 0)
- 13. No ($x/8$ will round towards 0 if $x < 0$)
- 14. No (simple counter example: 0)