

Announcement

- **Locations for lab hours are finalized (check course webpage).**
- **Lab 1 due this Friday (with the 2-day extension, that will be Sunday night).**

Floating Points (Cont)

B&O Readings: 2.4

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- **Floating point operations and rounding**
- Floating point in C

Floating Point Multiplication

- $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$

- **Exact Result:** $(-1)^s M 2^E$

- Sign s: $s1 \wedge s2$
- Mantissa M: $M1 \times M2$
- Exponent E: $E1 + E2$

- **Fixing**

- If $M \geq 2$, shift M right, increment E
- If E out of range, overflow
- Round M to fit `frac` precision

- **Implementation**

- Biggest chore is multiplying the Mantissas


Floating Point Addition

■ $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

- Assume $E1 > E2$

■ **Exact Result:** $(-1)^s M 2^E$

- Sign s , mantissa M :
 - Result of signed align & add
- Exponent E : $E1$




■ Fixing

- If $M \geq 2$, shift M right, increment E
- if $M < 1$, shift M left k positions, decrement E by k
- Overflow if E out of range
- Round M to fit `frac` precision

Floating Point Operations: Basic Idea

$$V = (-1)^s \cdot M \cdot 2^E$$



- $x \oplus_f y = \text{Round}(x + y)$
 - E could be very different
 - the binary point needs to line up
- $x \otimes_f y = \text{Round}(x \times y)$
 - need to ensure that the resulting exponent is still in range
- **Basic idea**
 - First **compute exact result**
 - Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly **round to fit into** frac

IEEE Rounding Modes

■ Rounding Modes (illustrate with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	−\$1.50
■ Towards zero	\$1	\$1	\$1	\$2	−\$1
■ Round down ($-\infty$)	\$1	\$1	\$1	\$2	−\$2
■ Round up ($+\infty$)	\$2	\$2	\$2	\$3	−\$1
■ Nearest Even (default)	\$1	\$2	\$2	\$2	−\$2

■ Round to nearest Even:

- When more than halfway, round up; when less than halfway, round down.
- When exactly halfway between two possible values, round it so that least significant digit is even
- The default rounding mode.
- Why? So that we don't introduce statistical bias.
- All others are statistically biased

Rounding Binary Numbers

- **When exactly halfway between two possible values**

- Round so that least significant digit is even

- **Binary Fractional Numbers**

- “Even” when least significant bit is 0
- “Half way” when bits to right of rounding position = $100..._2$

- **Examples**

- Round to nearest $1/4$ (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded Value
$2 \frac{3}{32}$	$10.00\textcolor{red}{11}_2$	10.00_2	($<1/2$ —down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	10.01_2	($>1/2$ —up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$	11.00_2	($1/2$ —up)	3
$2 \frac{5}{8}$	$10.10\textcolor{red}{100}_2$	10.10_2	($1/2$ —down)	$2 \frac{1}{2}$

Mathematical Properties of FP Add

- **Commutative?** Yes
- **Associative?** No
 - Overflow and inexactness of rounding
 - $(3.14+1e10) - 1e10 = 0$, $3.14+(1e10-1e10) = 3.14$
- **0 is additive identity?** Yes
- **Every element has additive inverse?** Almost
 - Yes, except for infinities & NaNs
- **Monotonicity** Almost
 - $a \geq b \Rightarrow a+c \geq b+c$
 - Except for infinities & NaNs

Mathematical Properties of FP Mult

- **Multiplication Commutative?** Yes
- **Multiplication is Associative?** No
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \text{inf}$, $1e20 * (1e20 * 1e-20) = 1e20$
- **1 is multiplicative identity?** Yes
- **Multiplication distributes over addition?** No
 - Possibility of overflow, inexactness of rounding
 - $1e20 * (1e20 - 1e20) = 0.0$, $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$
- **Monotonicity** Almost
 - $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c$
 - Except for infinities & NaNs

Recap: What We Learned Thus Far

- Due to inexactness of rounding and possibility of overflow, floating point operations are NOT associative. Although they are commutative and generally maintains monotonicity (except when you have +/- infinity / NaN involved).
- Watch out:
 - When reordering the order of floating point operations, you may not get the same result!
 - Compiler never does this!

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Rounding, addition, multiplication
- **Floating point in C**

Float and Double in C

- Designed to represent reals and fractional numbers
- Still fixed width, like the integer data types:
 - float: 4 bytes
 - double: 8 bytes

What Does This Code Print?

```
float f = 0.3;  
printf("%.20f\n", f);  
printf("%.20f\n", 0.1+0.2);
```

- A) 0.3 for both
- B) Not exactly 0.3, but two printouts show the same values
- C) Not exactly 0.3, and two printouts differ

```
0.30000001192092895508  
0.3000000000000000004441
```

Answer: C)

By default, real constants have type double, which has better precision (unless suffixed with `f` or `F`).



Floating Point in C

■ C Guarantees Two Levels

- `float` single precision
- `double` double precision

■ Conversions/Casting

- Casting between `int`, `float`, and `double` changes bit representation
- `double/float → int`
 - Truncates fractional part
 - Like rounding toward zero
 - Not defined when out of range or NaN: Generally sets to TMin
- `int → double`
 - Exact conversion, as long as `int` has ≤ 53 bit word size
- `int → float`
 - Will round according to rounding mode

Floating Point Puzzles

■ For each of the following C expressions, either:

- Argue that it is true for all argument values
- Explain why not true

```
int x = ...;  
long l = ...;  
float f = ...;  
double d = ...;
```

Assume neither
d nor f is NaN

A. $x == (\text{int})(\text{float})\ x$

B. $x == (\text{int})(\text{double})\ x$

C. $f == (\text{float})(\text{double})\ f$

D. $l == (\text{long})(\text{double})\ l$

E. $d == (\text{double})(\text{float})\ d$

F. $f == -(-f);$

G. $2/3 == 2/3.0$

H. $d < 0.0 \Rightarrow ((d*2) < 0.0)$

I. $d > f \Rightarrow -f > -d$

J. $d * d \geq 0.0$

K. $(f+d)-f == d$

Floating Point Puzzles Answers

- A. No, float has 23 frac bits, and x has 32 bits
- B. Yes, double has greater precision and range (52 frac bits)
- C. Yes, since double has a wider range and prec.
- D. No, long has 64 bits and double has only 52 frac bits
- E. No, can lose precision / overflow to infinity
- F. Yes, it simply negates the sign bit
- G. No, result of $2/3$ would be int so you get 0, $2/3.0$ will be a double
- H. Yes, even if it overflows, it will overflow to $-\text{inf}$
- I. Yes, monotonicity, but also, note that we won't lost anything; unlike int, the negative and positive range representable by float or double are the same
- J. Yes, though it may overflow to $+\text{INF}$, but that's still > 0
- K. No, floating ops are not associative; if f is a really large number and d is really small, $f+d$ will be rounded to about f , and you may get 0 on the left.



Machine-Level Programming I: Basics

B&O Readings: 2.1, 3.1-3.5

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Byte-Oriented Memory Organization



- **Programs refer to data by address**
 - Conceptually, envision it as a very large array of bytes
 - In reality, it's not, but can think of it that way
 - An address is like an index into that array
 - and, a *pointer* variable stores an address
- **Note: system provides private address spaces to each “process”**
 - Think of a process as a program being executed
 - So, a program can clobber its own data, but not that of others

Machine Words

- **Any given computer has a “Word Size”**
 - Nominal size of integer-valued data
 - Or, the size of an address
- Until recently, most machines used 32 bits (4 bytes) as word size
 - Limits addresses to 4GB (2^{32} bytes)
- These days, most machines have 64-bit word size
 - Potentially, could have 18 PB (petabytes) of addressable memory
 - That's 18.4×10^{15}
- Machines support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes


Word-Oriented Memory Organization

■ Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit word) or 8 (64-bit word)

■ Pointers in C:



```
int x = 100;  
int *y = &x; //y store addr of x  
*y = *y + 1; //x is now 101
```



Byte Ordering

Question: how are the bytes within a multi-byte word ordered in memory?

- Example: Variable x has 4-byte value of 0x01234567
- Address given by &x is 0x100



C) Both are valid

Answer: Both