

# Announcement

## ■ Short Office Hour today: until 6pm

## ■ Lab 1 regrade

- if you failed to test your code on linuxlab machines using driver.pl then:
- TEST your code on linuxlab machine using driver.pl and see if you lose points due to violation of programming rules
- If you are surprised by this, you can fix your code and push the changes.
- We will regrade your lab1, and assign you:  
min(old code running with btest,  
the correctness portion of new code running with driver.pl)  
(Your lab1 score will incur some penalty, but it's better than 0.)
- NOTE that this won't help you
  - if you have gotten a decent score already running with driver.pl the first time around; or
  - If you have not implemented the correct functionalities and gotten 0 with btest the first time around

# Today

## ■ Procedures

- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- Illustration of Recursion

# Recursive Function

```
/* Recursive popcount */  
long pcount_r(unsigned long x) {  
    if (x == 0)  
        return 0;  
    else  
        return (x & 1)  
            + pcount_r(x >> 1);  
}
```

```
pcount_r:  
    movl    $0, %eax  
    testq   %rdi, %rdi  
    je      .L6  
    pushq   %rbx  
    movq    %rdi, %rbx  
    andl    $1, %ebx  
    shrq    %rdi  
    call    pcount_r  
    addq    %rbx, %rax  
    popq    %rbx  
.L6:  
    ret
```

# Recursive Function Terminal Case

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rdi	x	Argument
%rax	Return value	Return value


# Recursive Function Register Save

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rdi	x	Argument

Why? Because %rbx is callee-save.




# Recursive Function Call Setup

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rdi	x >> 1	Rec. argument
%rbx	x & 1	Callee-saved




# Recursive Function Call

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rbx	x & 1	Callee-saved
%rax	Recursive call return value	

We have our local state in %rbx, and the recursive pcount\_r has the right argument.




# Recursive Function Result

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rbx	x & 1	Callee-saved
%rax	Return value	






# Recursive Function Completion

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rax	Return value	Return value



# Observations About Recursion

## ■ Handled Without Special Consideration

- Stack frames mean that each function call has private storage
  - Saved registers & local variables
  - Saved return pointer
- Register saving conventions prevent one function call from corrupting another's data
  - Unless the C code explicitly does so (e.g., buffer overflow)
- Stack discipline follows call / return pattern
  - If P calls Q, then Q returns before P
  - Last-In, First-Out

## ■ Also works for mutual recursion

- P calls Q; Q calls P


# x86-64 Procedure Summary

## ■ Important Points

- Stack is the right data structure for procedure call / return
  - If P calls Q, then Q returns before P

## ■ Recursion (& mutual recursion) handled by normal calling conventions

- Can safely store values in local stack frame and in callee-saved registers
- Put function arguments at top of stack
- Result return in **%rax**




# x86-64 Procedure Summary

## ■ Important Points

- Stack is the right data structure for procedure call / return
  - If P calls Q, then Q returns before P

## ■ Recursion (& mutual recursion) handled by normal calling conventions

- Can safely store values in local stack frame and in callee-saved registers
- Put function arguments at top of stack
- Result return in **%rax**
- **%rbp** may be optionally used as frame pointer when the compiler does not know the stack frame size.





# Machine-Level Programming IV: Compound Data Types

B&O Readings: 3.8-3.9

CSE 361: Introduction to Systems Software

**Instructor:**

I-Ting Angelina Lee

# Today: Compound Types in C

## ■ Arrays

- One-dimensional arrays
- Multi-dimensional / nested arrays
- Multi-level arrays

## ■ Structures

- Allocation
- Access
- Alignment





## ■ Unions

# Array Allocation

## ■ Basic Principle

$T \ A[L];$

- Array of data type  $T$  and length  $L$
- Contiguously allocated region of  $L * \text{sizeof}(T)$  bytes in memory




# Array Access

## ■ Basic Principle

$T$  **A**[ $L$ ] ;

- Array of data type  $T$  and length  $L$
- Identifier **A** can be used as a pointer to array element 0: Type  $T^*$



## ■ Reference      Type      Value

`val[4]`

`val`

`val+1`

`&val[2]`

`val[5]`

`*(val+1)`


`val + i`



# Array Example


```
#define ZIP_LEN 5
typedef int zipcode[ZIP_LEN];

zipcode wustl = { 6, 3, 1, 3, 0 };
zipcode duke  = { 2, 7, 7, 0, 8 };
zipcode brown = { 0, 2, 9, 1, 2 };
```



- Declaration “`zipcode wustl`” equivalent to “`int wustl[5]`”
- Example arrays were allocated in successive 20 byte blocks
  - Not guaranteed to happen in general

# Array Accessing Example



```
int get_digit  
  (zipcode z, int digit){  
    return z[digit];  
  }
```

```
# %rdi = z  
# %rsi = digit  
movl (%rdi,%rsi,4), %eax # z[digit]
```


- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at `%rdi + %rsi*4`
- Use memory reference `(%rdi,%rsi,4)`

# Array Loop Example

```
void zincr(zipcode z) {  
    size_t i;  
    for (i = 0; i < ZIP_LEN; i++)  
        z[i]++;  
}
```

```
# %rdi = z  
movl    $0, %eax           # i = 0  
jmp     .L3                # goto middle  
.L4:                        # loop:  
    addl    $1, (%rdi,%rax,4) # z[i]++  
    addq    $1, %rax        # i++  
.L3:                        # middle  
    cmpq    $4, %rax        # i:4  
    jbe     .L4             # if <=, goto loop  
ret
```

# Array Accessing Puzzle



```
... foo(zipcode z, int digit){  
    ...  
}
```


```
# %rdi = z  
# %rsi = digit  
leaq (%rdi,%rsi,4), %rax
```

- Q: What does foo return?
- A: `&(z[digit]) ;`



# Understanding Pointers & Arrays #1


Decl	<i>A<sub>n</sub></i>			<i>*A<sub>n</sub></i>		
	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>						
<code>int *A2</code>						



- **Cmp:** Compiles (Y/N)
- **Bad:** Possible bad pointer reference (Y/N)
- **Size:** Value returned by `sizeof`

# Understanding Pointers & Arrays #1



Decl	<i>A<sub>n</sub></i>			<i>*A<sub>n</sub></i>		
	Cmp	Bad	Size	Cmp	Bad	Size
int A1[3]	Y	N	12	Y	N	4
int *A2	Y	N	8	Y	Y	4



- Cmp: Compiles (Y/N)
- Bad: Possible bad pointer reference (Y/N)
- Size: Value returned by `sizeof`

# Understanding Pointers & Arrays #2



Decl	<i>A<sub>n</sub></i>			<i>*A<sub>n</sub></i>			<i>**A<sub>n</sub></i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	<b>Y</b>	<b>N</b>	<b>12</b>	<b>Y</b>	<b>N</b>	<b>4</b>			
<code>int *A2[3]</code>									



- **Cmp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**

# Understanding Pointers & Arrays #2

Decl	<i>A<sub>n</sub></i>			<i>*A<sub>n</sub></i>			<i>**A<sub>n</sub></i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8	Y	Y	4



- **Cmp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**



# Today: Compound Types in C

## ■ Arrays

- One-dimensional arrays
- Multi-dimensional / nested arrays
- Multi-level arrays


## ■ Structures

- Allocation
- Access
- Alignment

## ■ Unions

# Nested Array Example

```
#define COUNT 4
typedef int zipcode[ZIP_LEN];
zipcode stlz[COUNT] =
    {{6, 3, 1, 3, 0},
     {6, 3, 1, 0, 3 },
     {6, 3, 1, 3, 9 },
     {6, 3, 1, 0, 9 }};
```



- **“`zipcode stlz[4]`” equivalent to “`int stlz[4][5]`”**
  - Variable `stlz`: array of 4 elements, allocated contiguously
  - Each element is an array of 5 `int`'s, allocated contiguously
- **“Row-Major” ordering of all elements guaranteed**

# Multidimensional (Nested) Arrays

## ■ Declaration

`T A[R][C];`


- 2D array of data type  $T$
- $R$  rows,  $C$  columns
- Type  $T$  element requires  $K$  bytes

## ■ Array Size

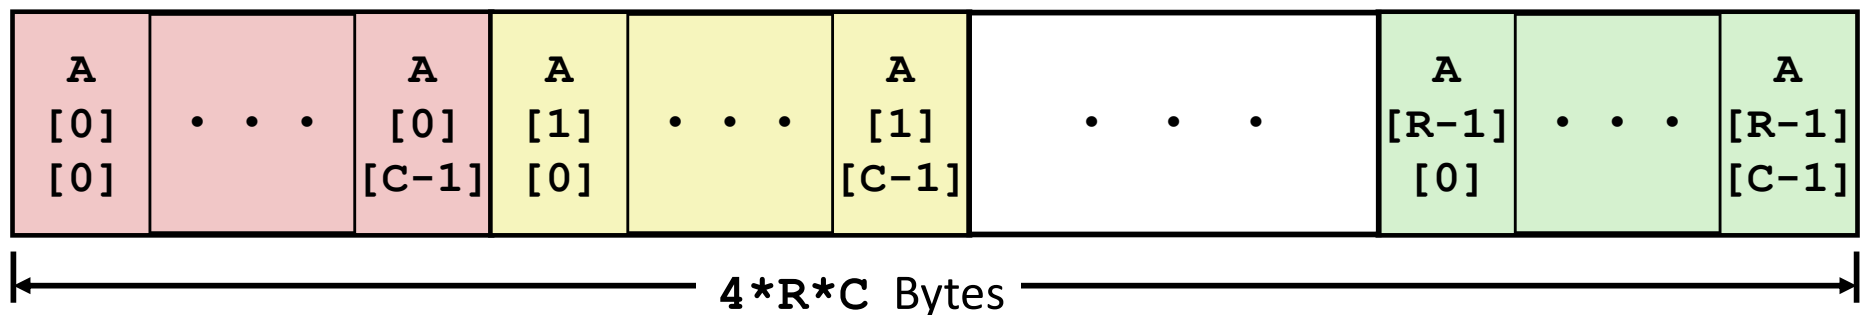
- $R * C * K$  bytes

## ■ Arrangement

- Row-Major Ordering



`int A[R][C];`




# Nested Array Row Access

## ■ Row Vectors

- $A[i]$  is array of  $C$  elements
- Each element of type  $T$  requires  $K$  bytes
- Starting address  $A + i * (C * K)$

```
int A[R][C];
```




# Nested Array Element Access

## ■ Array Elements

- $A[i][j]$  is element of type  $T$ , which requires  $K$  bytes
- Address  $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```




# Multi-Level Array Example

```
#define COUNT 4
typedef int zipcode[ZIP_LEN];
zipcode wustl = { 6, 3, 1, 3, 0 };
zipcode duke = { 2, 7, 7, 0, 8 };
zipcode brown = { 0, 2, 9, 1, 2 };
```

```
#define UCOUNT 3
int *univ[UCOUNT] = {duke,wustl,brown};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
  - 8 bytes
- Each pointer points to array of `int`'s



# Element Access in Multi-Level Array

```
int get_univ_digit
(int *univ, size_t index, size_t digit)
{
    return univ[index][digit];
}
```

```
salq    $2, %rdx          # 4*digit
addq    %rdi(,%rsi,8), %rdx # p = univ[index] + 4*digit
movl    (%rdx), %eax       # return *p
ret
```


## ■ Computation

- Element access **Mem[Mem[univ+8\*index]+4\*digit]**
- Must do two memory reads
  - First get pointer to row array
  - Then access element within array

# Array Element Accesses

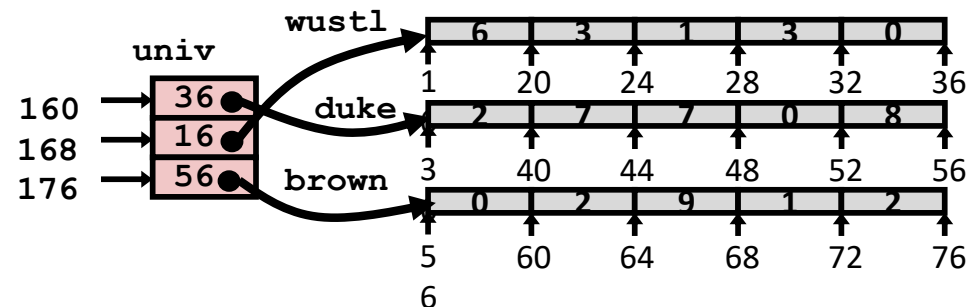
## Nested array

```
int get_stl_zip
(size_t index, size_t digit)
{
    return stlz[index][digit];
}
```



## Multi-level array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



Accesses looks similar in C, but address computations very different:

`Mem[stlz+20*index+4*digit]`   `Mem[Mem[univ+8*index]+4*digit]`



# Ex: N X N Matrix Code (2D Nested Array)

## ■ Variable dimensions, explicit indexing

- Traditional way to implement arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))  
/* Get element a[i][j] */  
int vec_ele(size_t n, int *a,  
            size_t i, size_t j){  
    return a[IDX(n,i,j)];  
}
```

## ■ Variable dimensions, implicit indexing

- Now supported by gcc

```
/* Get element a[i][j] */  
int var_ele(size_t n, int a[n][n],  
            size_t i, size_t j) {  
    return a[i][j];  
}
```

## ■ Fixed dimensions

- Know value of N at compile time

```
#define N 16  
/* Get element a[i][j] */  
int fix_ele(int a[N][N],  
            size_t i, size_t j){  
    return a[i][j];  
}
```

# n X n Matrix Access

## ■ Array Elements

- Address  $A + i * (C * K) + j * K$
- $C = n, K = 4$
- Must perform integer multiplication

```
/* Get element a[i][j] */  
int var_ele(size_t n, int a[n][n], size_t i, size_t j) {  
    return a[i][j];  
}
```

```
# n in %rdi, a in %rsi, i in %rdx, j in %rcx  
imulq    %rdx, %rdi          # n*i  
leaq     (%rsi,%rdi,4), %rax  # a + 4*n*i  
movl     (%rax,%rcx,4), %eax  # a + 4*n*i + 4*j  
ret
```

# 16 X 16 Matrix Access

## ■ Array Elements

- Address  $\mathbf{A} + i * (\mathbf{C} * \mathbf{K}) + j * \mathbf{K}$
- $\mathbf{C} = 16, \mathbf{K} = 4$
- Size known at compile time; may be able to compute index offset more efficiently

```
#define N 16
/* Get element a[i][j] */
int fix_ele(int a[N][N], size_t i, size_t j) {
    return a[i][j];
}
```

```
# a in %rdi, i in %rsi, j in %rdx
salq    $6, %rsi           # 64*i
addq    %rsi, %rdi         # a + 64*i
movl    (%rdi,%rdx,4), %eax # M[a + 64*i + 4*j]
ret
```

# Recap: What We Learned Thus Far

## ■ Nested Array vs Multi-level Array:

- Nested array: allocated in contiguous memory, row-major order.
- Multi-level array: each level uses continuous memory but not across levels.

## ■ Compiler supports limited form of nested array

- The array dimensions must be known to the compiler, either as a compile-time constants or variables.
- Knowing the array size at compile time may lead to more efficient array index calculation.

# Today: Compound Types in C

## ■ Arrays

- One-dimensional arrays
- Multi-dimensional / nested arrays
- Multi-level arrays

## ■ Structures

- Allocation
- Access
- Alignment


## ■ Unions

# Struct in C

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};  
struct rec g;  
struct rec *r = &g;
```

or

```
typedef struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
} rec_t;  
rec_t g;  
rec_t *r = &g;
```



## ■ Concept

- Groups data of possibly different types into a single object
- Refer to members within structure by names
  - `r->a[2]`
  - `g.a[2]`