

Machine-Level Programming IV: Compound Data Types, Cont'd

B&O Readings: 3.8-3.9

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Today: Compound Types in C

■ Arrays

- One-dimensional arrays
- Multi-dimensional / nested arrays
- Multi-level arrays

■ Structures

- Allocation
- Access
- Alignment


■ Unions

Struct in C

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};  
struct rec g;  
struct rec *r = &g;
```

or

```
typedef struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} rec_t;  
rec_t g;  
rec_t *r = &g;
```




■ Concept

- Groups data of possibly different types into a single object
- Refer to members within structure by names
 - `r->a[2]`
 - `g.a[2]`

Structure Representation


```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} g;  
struct rec *r = &g;
```



- **Structure represented as block of memory**
 - Big enough to hold all of the fields
- **Fields ordered according to declaration**
 - Even if another ordering could yield a more compact representation
- **Compiler determines overall size + positions of fields**
 - Machine-level program has no understanding of the structures in the source code

Structure Access

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```



■ Accessing Structure Member

- Pointer indicates first byte of structure
- Access elements with offsets

```
void  
set_i(struct rec *r,  
      int val)  
{  
    r->i = val;  
}
```

```
# %rdi = r  
# %esi = val  
movslq %esi, %rsi      # sign extend val to (long)val  
movq    %rsi, 16(%rdi)  # store (long)val in r->i
```

Example: Struct Access


```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```

```
jmp     .L11          # goto test  
.L12:          # loop:  
    movslq 16(%rdi), %rax    # tmp = *(r+16)  
    movl   %esi, (%rdi,%rax,4) # *(r+4*tmp) = val  
    movq   24(%rdi), %rdi    # r = *(r+24)  
.L11:          # test:  
    testq  %rdi, %rdi       # Test r  
    jne    .L12             # if !=0 goto loop
```

■ C Code

```
void set_val  
(struct rec *r, int val){  
    while (...) {  
        ...;  
        ...;  
        ...;  
    }  
}
```

Register	Value
%rdi	r
%esi	val



Example: Struct Access


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```

```
jmp      .L11          # goto test  
.L12:      # loop:  
    movq    16(%rdi), %rax    # tmp = Mem[r+16]  
    movl    %esi, (%rdi,%rax,4) # Mem[r+4*tmp] = val  
    movq    24(%rdi), %rdi    # r = Mem[r+24]  
.L11:      # test:  
    testq   %rdi, %rdi      # Test r  
    jne     .L12            # if !=0 goto loop
```

■ C Code


```
void set_val  
(struct rec *r, int val){  
    while (r) {  
        long tmp = r->i;  
        r->a[tmp] = val;  
        r = r->next;  
    }  
}
```

Register	Value
%rdi	r
%esi	val



Structures & Alignment

■ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Alignment Principle

- Primitive data type requires K bytes: field must be ***K-byte aligned*** (i.e., its address is multiple of K).
- The struct itself must also be multiple of K, where K is the largest alignment requirement among its fields.

Specific Cases of Alignment (x86-64)

- **1 byte: `char`, ...**
 - no restrictions on address
- **2 bytes: `short`, ...**
 - lowest 1 bit of address must be 0_2
- **4 bytes: `int`, `float`, ...**
 - lowest 2 bits of address must be 00_2
- **8 bytes: `double`, `long`, `char *`, ...**
 - lowest 3 bits of address must be 000_2
- **16 bytes: `long double` (GCC on Linux)**
 - lowest 4 bits of address must be 0000_2

Satisfying Alignment with Structures

■ Within structure:

- Must satisfy each element's alignment requirement

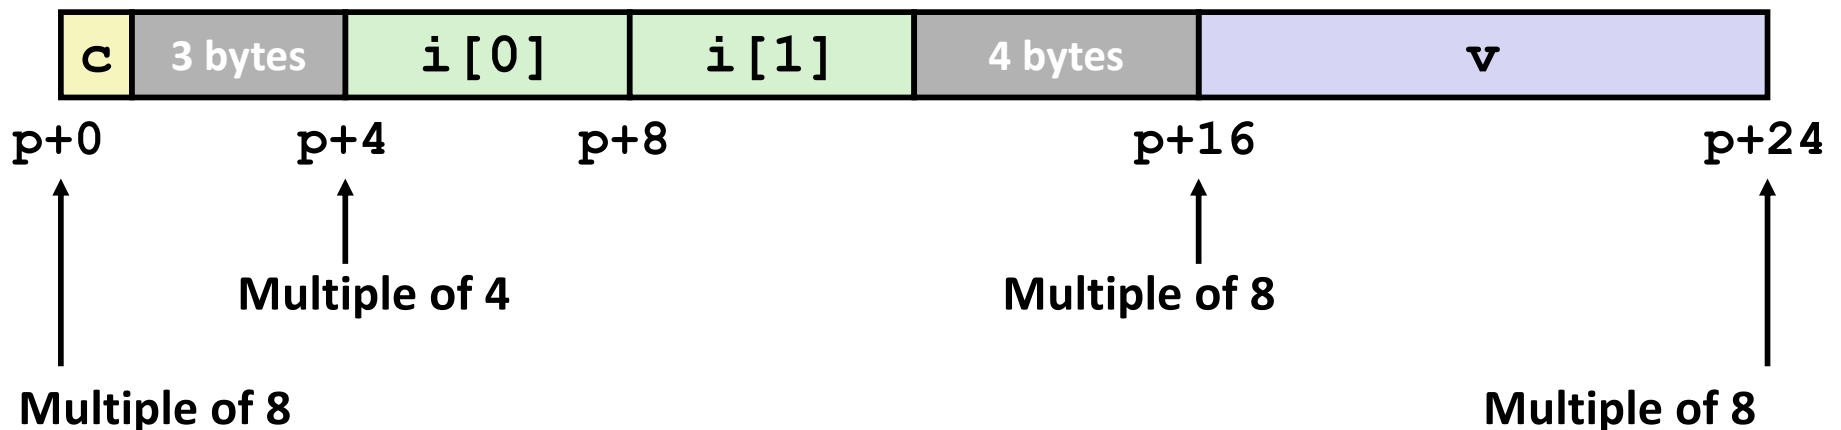
■ Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K (so an array of structs just works!)

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Example:

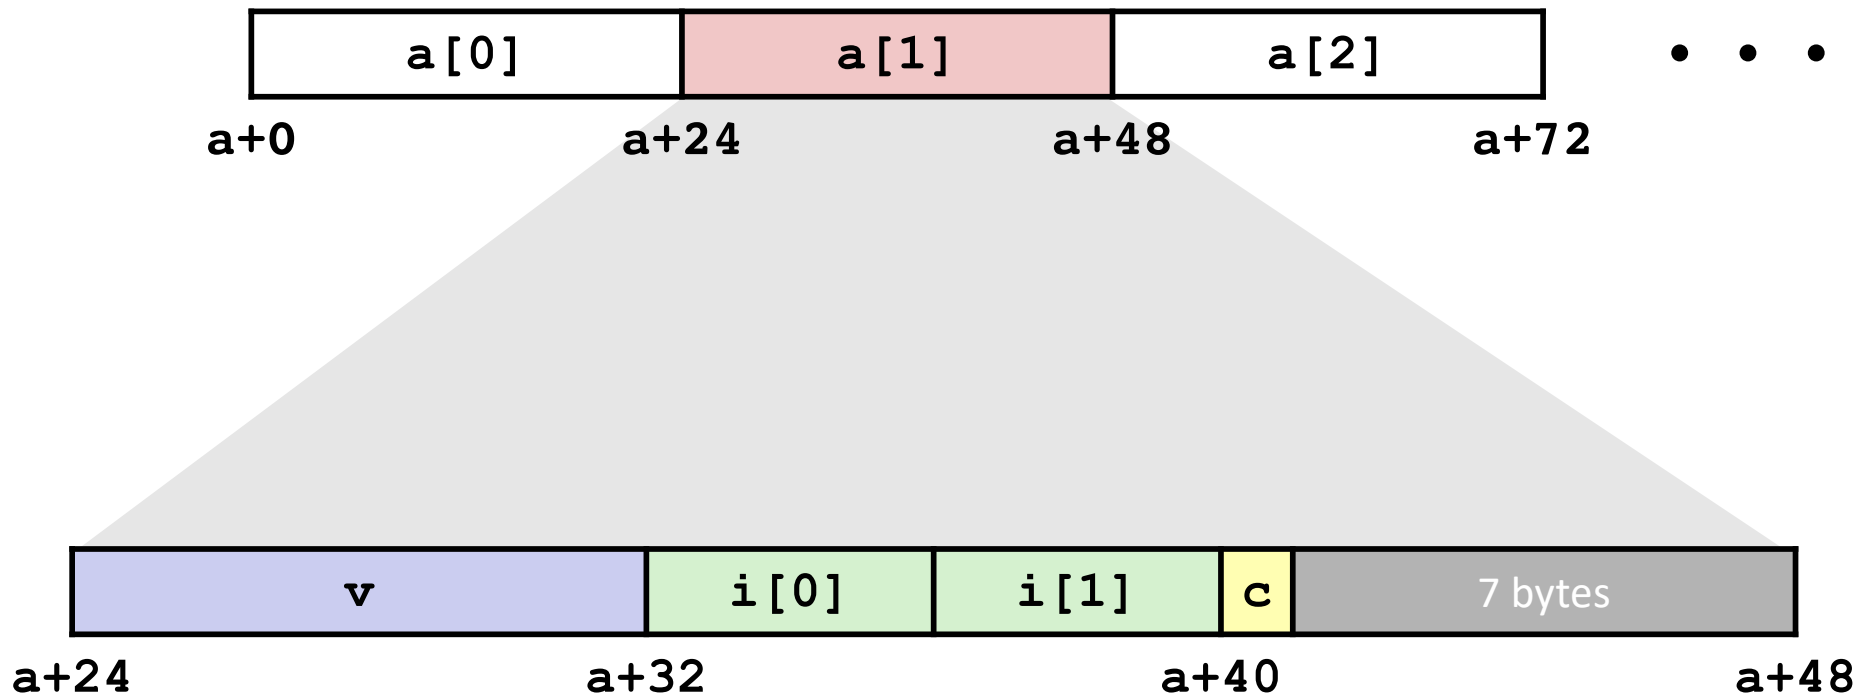
- K = 8, due to **double** element



Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

```
struct S3 {  
    short i;  
    float v;  
    short j;  
};
```

- Compute array offset
 - `sizeof(S3) = 12`, including alignment spacers
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`



Q: What's the assembly? (Note that `a` is declared as a global array)

```
Struct S3 a[10];  
short get_j(int idx){  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax    #3*idx  
movzwl a+8(,%rax,4),%eax    #a+8+12*idx
```

Saving Space

- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect: saving 4 bytes



Alignment Principles

■ Aligned Data

- For primitive data type that requires K bytes, its address must be a multiple of K.
- The whole struct must have size that is a multiple of K (where K is the largest alignment requirement among its fields).
- Required on some machines; advised on x86-64

■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory trickier when datum spans 2 pages

■ Compiler

- Inserts gaps in structure to ensure correct alignment of fields

Meeting Overall Alignment Requirement

- Q: What's the size of the following struct?


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



Meeting Overall Alignment Requirement

- Q: What's the size of the following struct?

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



- A: 24

Multiple of K=8



Recap: What We Learned Thus Far

- **The `struct` type in C:**

- The size and alignment requirement are determined by the largest alignment requirement amongst its fields (determined by the compiler).

- **Once we get to the assembly-level, the assembly instruction knows nothing about the compound type.**

- Access to a field in a struct is just like any memory reference.
- The compiler generates the memory reference with the address of the struct + the right offset.

Alignment Calculation Example

- Q: What's the alignment requirement for P3?

```
struct P1 { int i; char c; char d; long j; }
```

```
struct P2 { short w[3]; char c[3]; };
```

```
struct P3 { struct P2 a[2]; struct P1 t };
```

- A: 16
- B: 8
- C: 24
- D: None of the above



Alignment Calculation Example

■ Q: What's the alignment requirement for P3?

<code>struct P1 { int i; char c; char d; long j; }</code>	alignmt: 8 size: 16
offset 0 4 5 8	
<code>struct P2 { short w[3]; char c[3]; };</code>	alignmt: 2 size: 10
offset 0 6	
<code>struct P3 { struct P2 a[2]; struct P1 t };</code>	alignmt: 8 size: 40
offset 0 24	
:	

- A: 16 ■ B: 8 ■ C: 24 ■ D: None of the above

Answer: B



Today: Compound Types in C

■ Arrays

- One-dimensional arrays
- Multi-dimensional / nested arrays (next time)
- Multi-level arrays (next time)

■ Structures

- Allocation
- Access
- Alignment

■ Unions

Union in C

- Circumvent the type system of C
- Allowing a single object to be referenced according to multiple types
- Fields share the same memory location
- Refer to members within structure by names
 - `up->i[2]`
 - `(*up).i[2]`



```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} u;  
union U1 *up = &u;
```

Union Allocation

- Allocate according to largest element
- Can only use one field at a time


```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} u;  
union U1 *up = &u;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} s;  
struct S1 *p = &s;
```



Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



```
float bit2float(unsigned u)  
{  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as (float) u ?


```
unsigned float2bit(float f)  
{  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

Same as (unsigned) f ?



Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



```
float bit2float(unsigned u)  
{  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as (float) u?

```
unsigned float2bit(float f)  
{  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

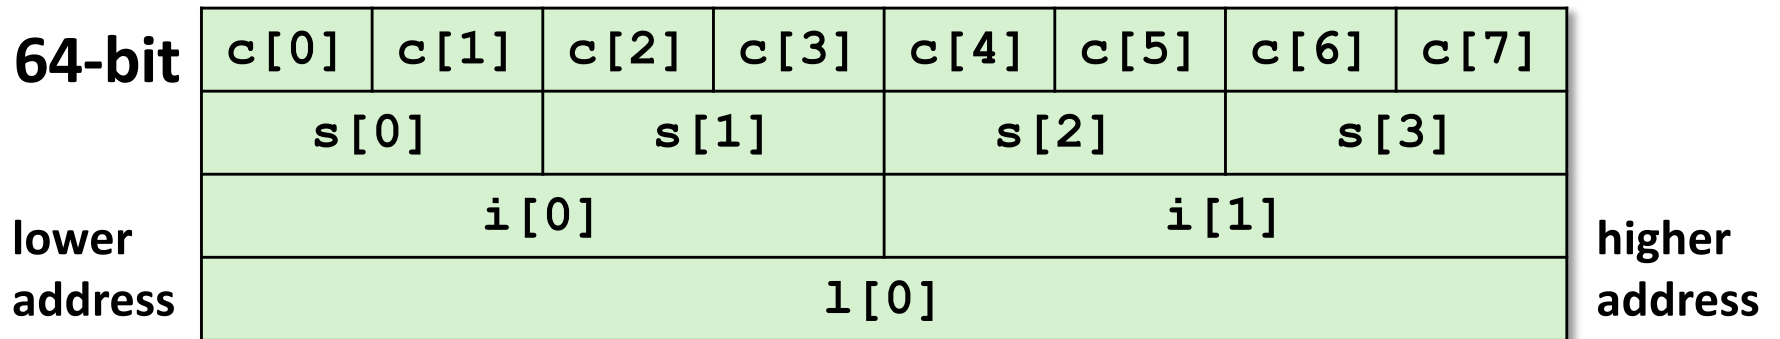
Same as (unsigned) f?

Answer: NO! With type cast, you change the bit representations.



Byte Ordering Puzzle

```
union {  
    unsigned char c[8];  
    unsigned short s[4];  
    unsigned int i[2];  
    unsigned long l[1];  
} dw;
```



What's the size of this union? 8 bytes

Byte Ordering Puzzle (Cont.)

Little Endian

	f0	f1	f2	f3	f4	f5	f6	f7	
	c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]	
	s[0]		s[1]		s[2]		s[3]		
lower address	i[0]				i[1]				higher address
	l[0]								

What are the output on x86-64?

Characters 0-7 == [0xf0, 0xf1, 0xf2, 0xf3, 0xf4, 0xf5, 0xf6, 0xf7]

Shorts 0-3 ==

Ints 0-1 ==

Long 0 ==



Byte Ordering Puzzle (Cont.)

Little Endian

	f0	f1	f2	f3	f4	f5	f6	f7	
	c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]	
	s[0]		s[1]		s[2]		s[3]		
lower address	i[0]				i[1]				higher address
	l[0]								

What are the output on x86-64?

Characters 0-7 == [0xf0, 0xf1, 0xf2, 0xf3, 0xf4, 0xf5, 0xf6, 0xf7]
Shorts 0-3 == [0xf1f0, 0xf3f2, 0xf5f4, 0xf7f6]
Ints 0-1 == [0xf3f2f1f0, 0xf7f6f5f4]
Long 0 == [0xf7f6f5f4f3f2f1f0]



Byte Ordering Revisited

■ Idea

- Short/long/quad words stored in memory as 2/4/8 consecutive bytes
- Which byte is most (least) significant?
- Can cause problems when exchanging binary data between machines

■ Big Endian

- Most significant byte has the lowest address
- Sparc

■ Little Endian

- Least significant byte has the lowest address
- Intel x86, ARM Android and IOS

Recap: What We Learned Thus Far

■ The `union` type in C:

- All fields share the same storage (unlike `struct`)
 - Every field has offset 0; some are padded
 - Same bits just reinterpreted!
 - **Byte ordering depends on systems (Little vs Big Endian)**
- The size and alignment for a union type are determined by the largest alignment requirement amongst its fields (determined by the compiler, just like `struct`).

- Once we get to the assembly-level, the assembly instruction knows nothing about the compound type.



Machine-Level Programming V: Advanced Topic

B&O Readings: 3.9-3.10

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

x86-64 Linux Memory Layout

not drawn to scale

0x00007FFFFFFF

■ Stack

- Runtime stack (8MB limit)
- E. g., local variables

■ Heap


- Dynamically allocated as needed
- When call `malloc()`, `calloc()`, `new()`

■ Data

- Statically allocated data
- E.g., global vars, **static** vars, string constants

■ Text / Shared Libraries

- Executable machine instructions
- Read-only



not drawn to scale

Memory Allocation Example

```
char big_array[1L<<24]; /* 16 MB */
char huge_array[1L<<31]; /* 2 GB */

int global = 0;

int useless() { return 0; }

int main ()
{
    void *p1, *p2, *p3, *p4;
    int local = 0;
    p1 = malloc(1L << 28); /* 256 MB */
    p2 = malloc(1L << 8); /* 256 B */
    p3 = malloc(1L << 32); /* 4 GB */
    p4 = malloc(1L << 8); /* 256 B */
    /* Some print statements ... */
}
```



Where does everything go?

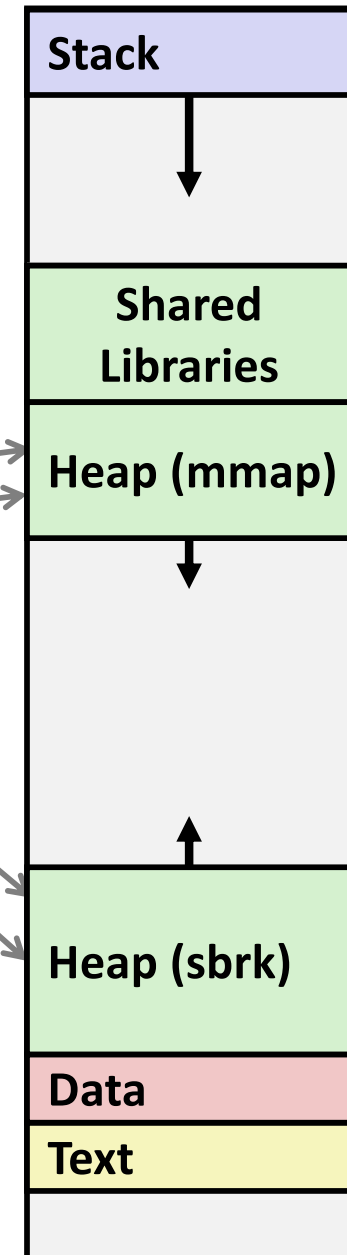
x86-64 Example Addresses

address range $\sim 2^{47}$

not drawn to scale

```
&local  
&p1  
p1  
p3  
p4  
p2  
big_array  
huge_array  
main()  
useless()
```

```
0x00007ffee4379f54  
0x00007ffee4379f58  
0x00007fffe7a19010  
0x00007fffe7a18010  
0x00000000081602120  
0x00000000081602010  
0x00000000080601060  
0x00000000000601060  
0x000000000004005c6  
0x000000000004005c0
```



Today

- Memory Layout
- **Buffer Overflow**
 - Vulnerability
 - Protection

Recall: Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    long l;  
} struct_t;  
  
long fun(int i) {  
    volatile struct_t s;  
    s.l = 999;  
    s.a[i] = 1000; /* Possibly out of bounds */  
    return s.l;  
}
```

fun(0)	→	999
fun(1)	→	999
fun(2)	→	1000
fun(3)	→	4294967296999
fun(4)	→	Segmentation fault


- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    long l;  
} struct_t;
```

<code>fun(0)</code>	<code>→</code>	999
<code>fun(1)</code>	<code>→</code>	999
<code>fun(2)</code>	<code>→</code>	1000
<code>fun(3)</code>	<code>→</code>	4294967296999
<code>fun(4)</code>	<code>→</code>	Segmentation fault

Explanation:



Result is system specific!

Such problems are a BIG deal

- **Generally called a “buffer overflow”**

- when exceeding the memory size allocated for an array

- **Why a big deal?**

- It's the #1 technical cause of security vulnerabilities
 - #1 overall cause is social engineering / user ignorance

- **Most common form**

- Unchecked lengths on string inputs
- Particularly for bounded character arrays on the stack
 - sometimes referred to as stack smashing