

Machine-Level Programming V: Advanced Topic (cont'd)

B&O Readings: 3.9-3.10

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Such problems are a BIG deal

- **Generally called a “buffer overflow”**

- when exceeding the memory size allocated for an array

- **Why a big deal?**

- It's the #1 technical cause of security vulnerabilities
 - #1 overall cause is social engineering / user ignorance

- **Most common form**

- Unchecked lengths on string inputs
 - Particularly for bounded character arrays on the stack
 - sometimes referred to as stack smashing

String Library Code

■ Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- No way to specify limit on number of characters to read (how big is the array `dest`??)

■ Similar problems with other library functions

- `strcpy`, `strcat`: Copy strings of arbitrary length
- `scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
void call_echo() {  
    echo();  
}
```

← btw, how big
is big enough?

```
unix>./bufdemo-nsp  
Type a string:01234567890  
01234567890
```

```
unix>./bufdemo-nsp  
Type a string:012345678901  
012345678901  
Segmentation fault (core dumped)
```

Why does the code seg fault with input of ≥ 12 characters?

Buffer Overflow Disassembly

echo:

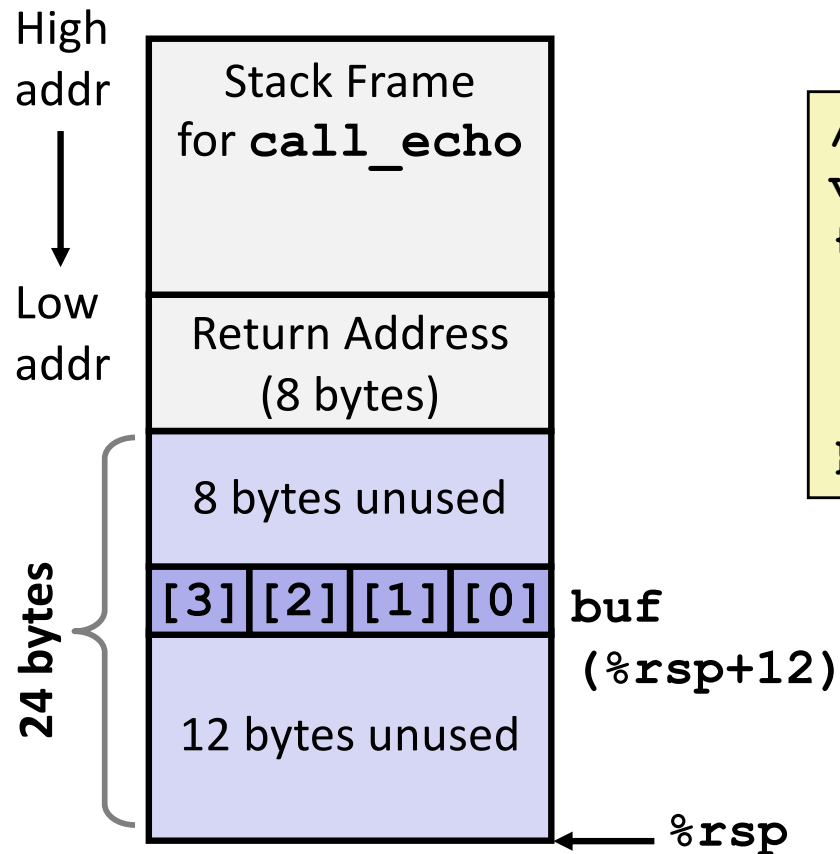
0000000000400694 <echo>:

400694:	48 83 ec 18	sub	\$0x18,%rsp
400698:	48 8d 7c 24 0c	lea	0xc(%rsp),%rdi
40069d:	e8 a4 ff ff ff	callq	400646 <gets>
4006a2:	48 8d 7c 24 0c	lea	0xc(%rsp),%rdi
4006a7:	e8 44 fe ff ff	callq	4004f0 <puts@plt>
4006ac:	48 83 c4 18	add	\$0x18,%rsp
4006b0:	c3	retq	

call_echo:

4006b1:	48 83 ec 08	sub	\$0x8,%rsp
4006b5:	b8 00 00 00 00	mov	\$0x0,%eax
4006ba:	e8 d5 ff ff ff	callq	400694 <echo>
4006bf:	48 83 c4 08	add	\$0x8,%rsp
4006c3:	c3	retq	

Buffer Overflow Stack



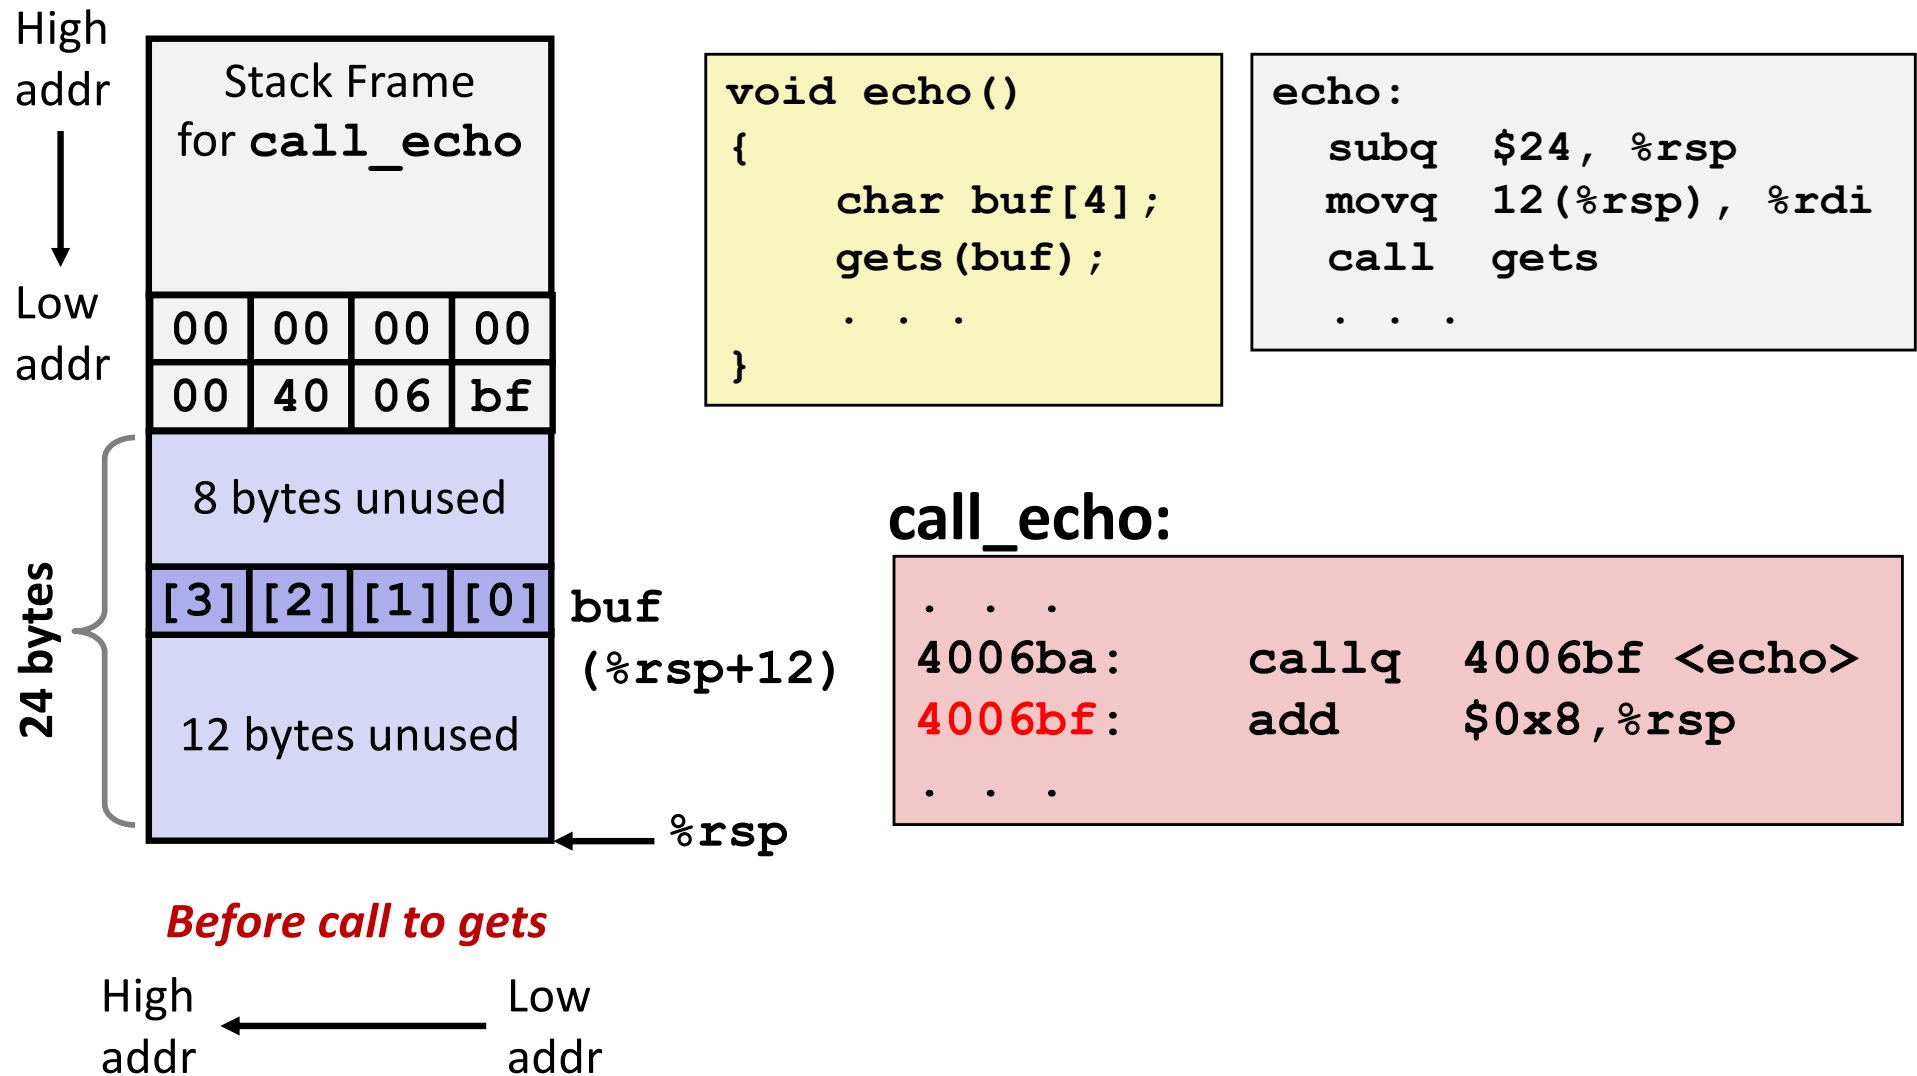
Before call to gets

High addr ← Low addr

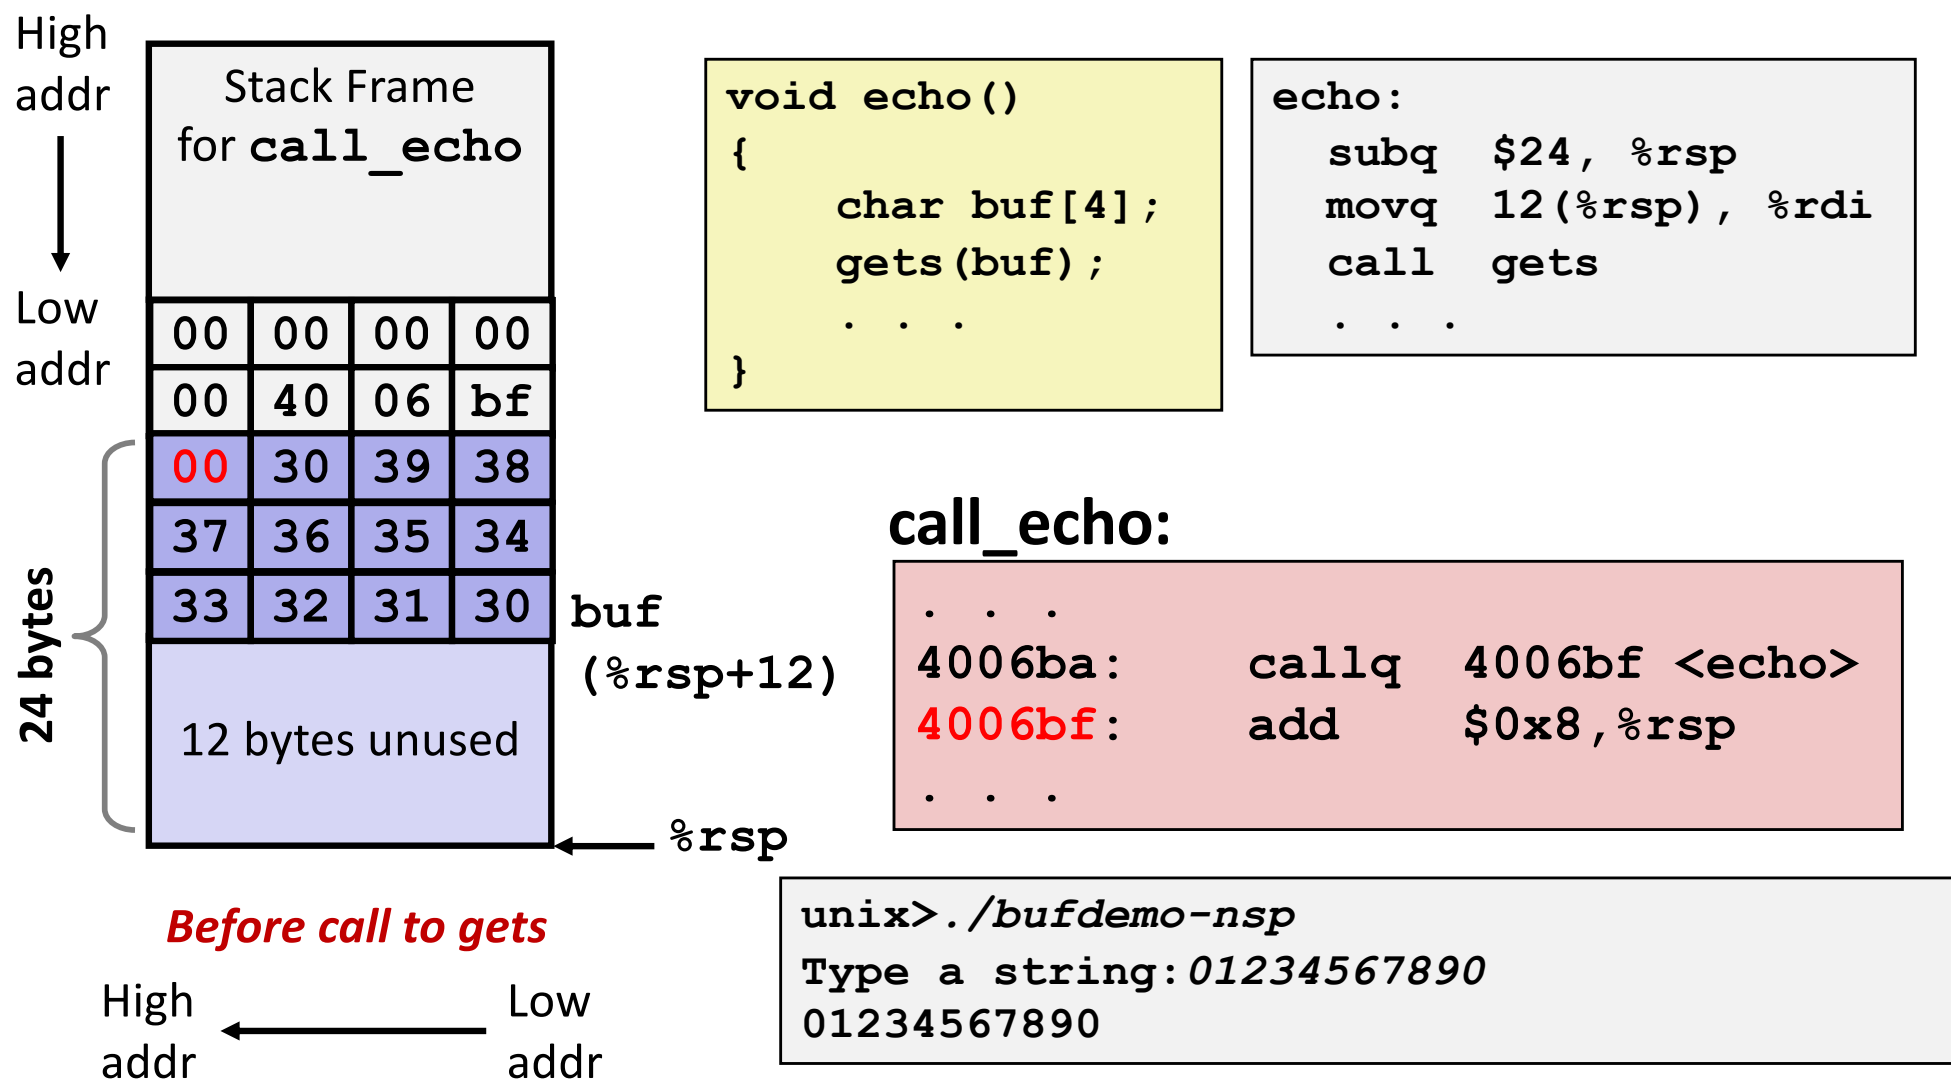
```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    subq    $24, %rsp
    movq    12(%rsp), %rdi
    call    gets
    . . .
```

Buffer Overflow Stack

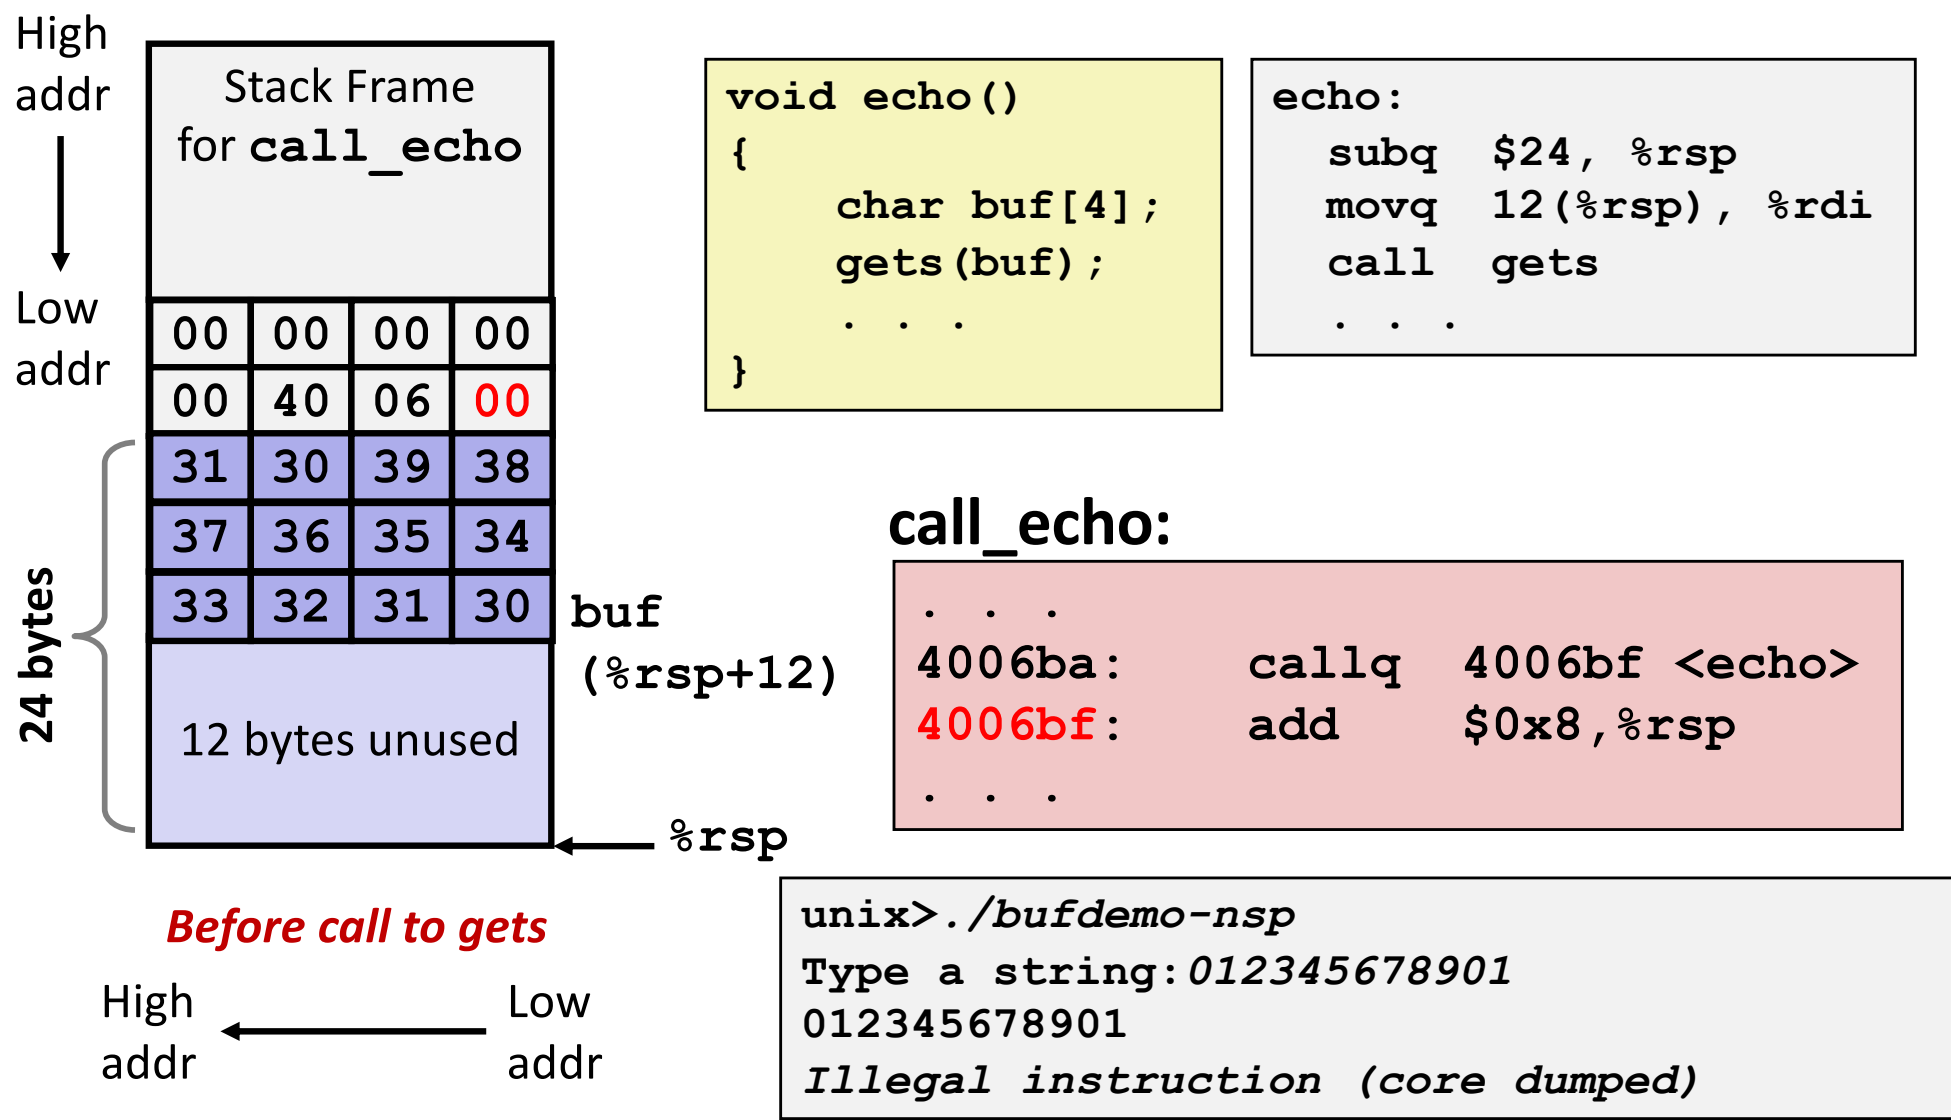


Buffer Overflow Stack



12 chars total: C string is null-terminated ('\0')!
Overflowed buffer, but did not corrupt state

Buffer Overflow Stack



Overflowed buffer and corrupted return address.

Stack Smashing Attacks

```
void P() {  
    Q();  
    ...  
}
```


return
address
A

```
int Q() {  
    char buf[64];  
    gets(buf);  
    ...  
    return ...;  
}
```

```
void S() {  
    /* Something  
       unexpected */  
    ...  
}
```

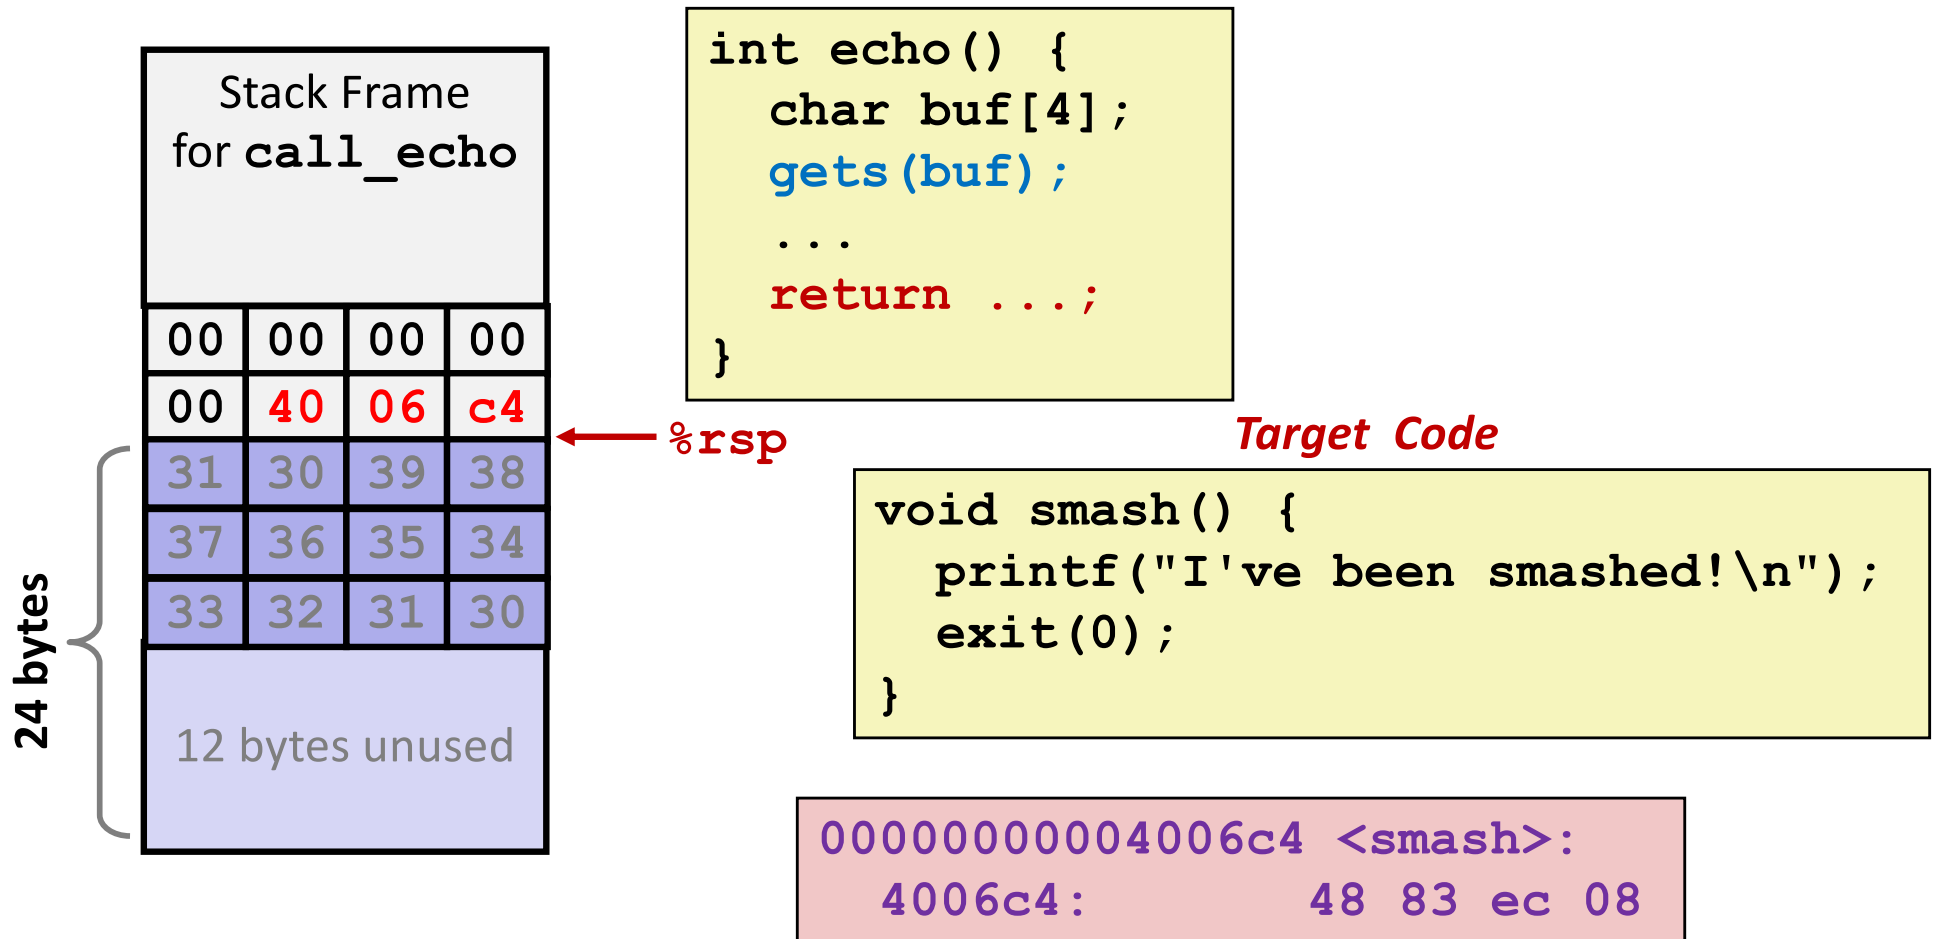
data written
by `gets()`

Stack after call to `gets()`



- Overwrite normal return address **A** with address of some other code **S**
- When **Q** executes `ret`, will jump to other code

Crafting Smashing String

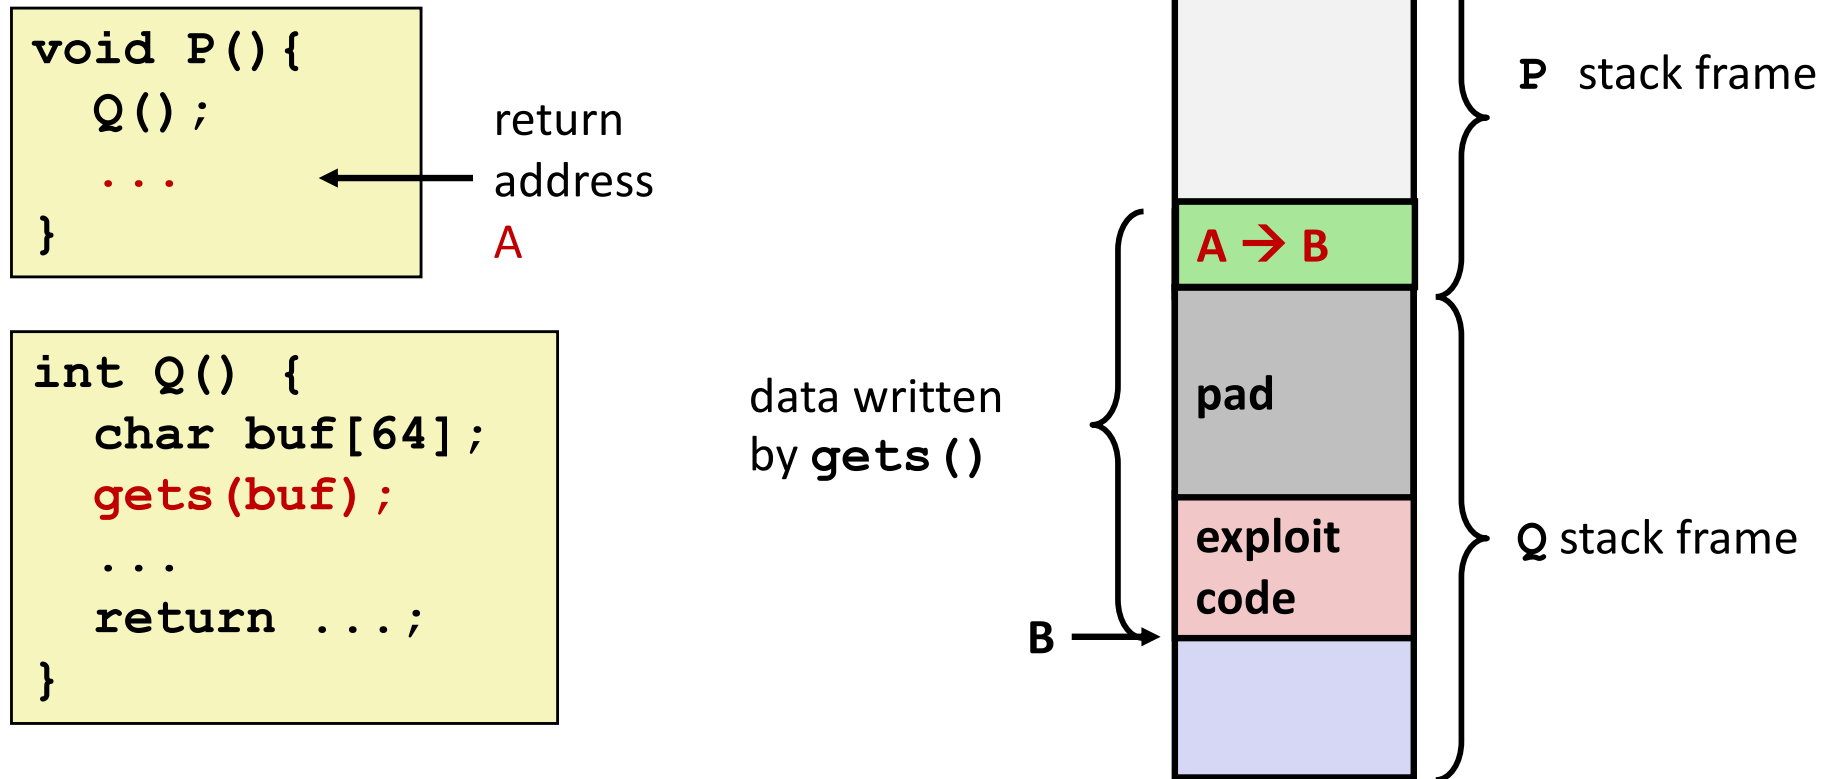


Attack String (Hex)

30 31 32 33 34 35 36 37 38 39 30 31 c4 06 40

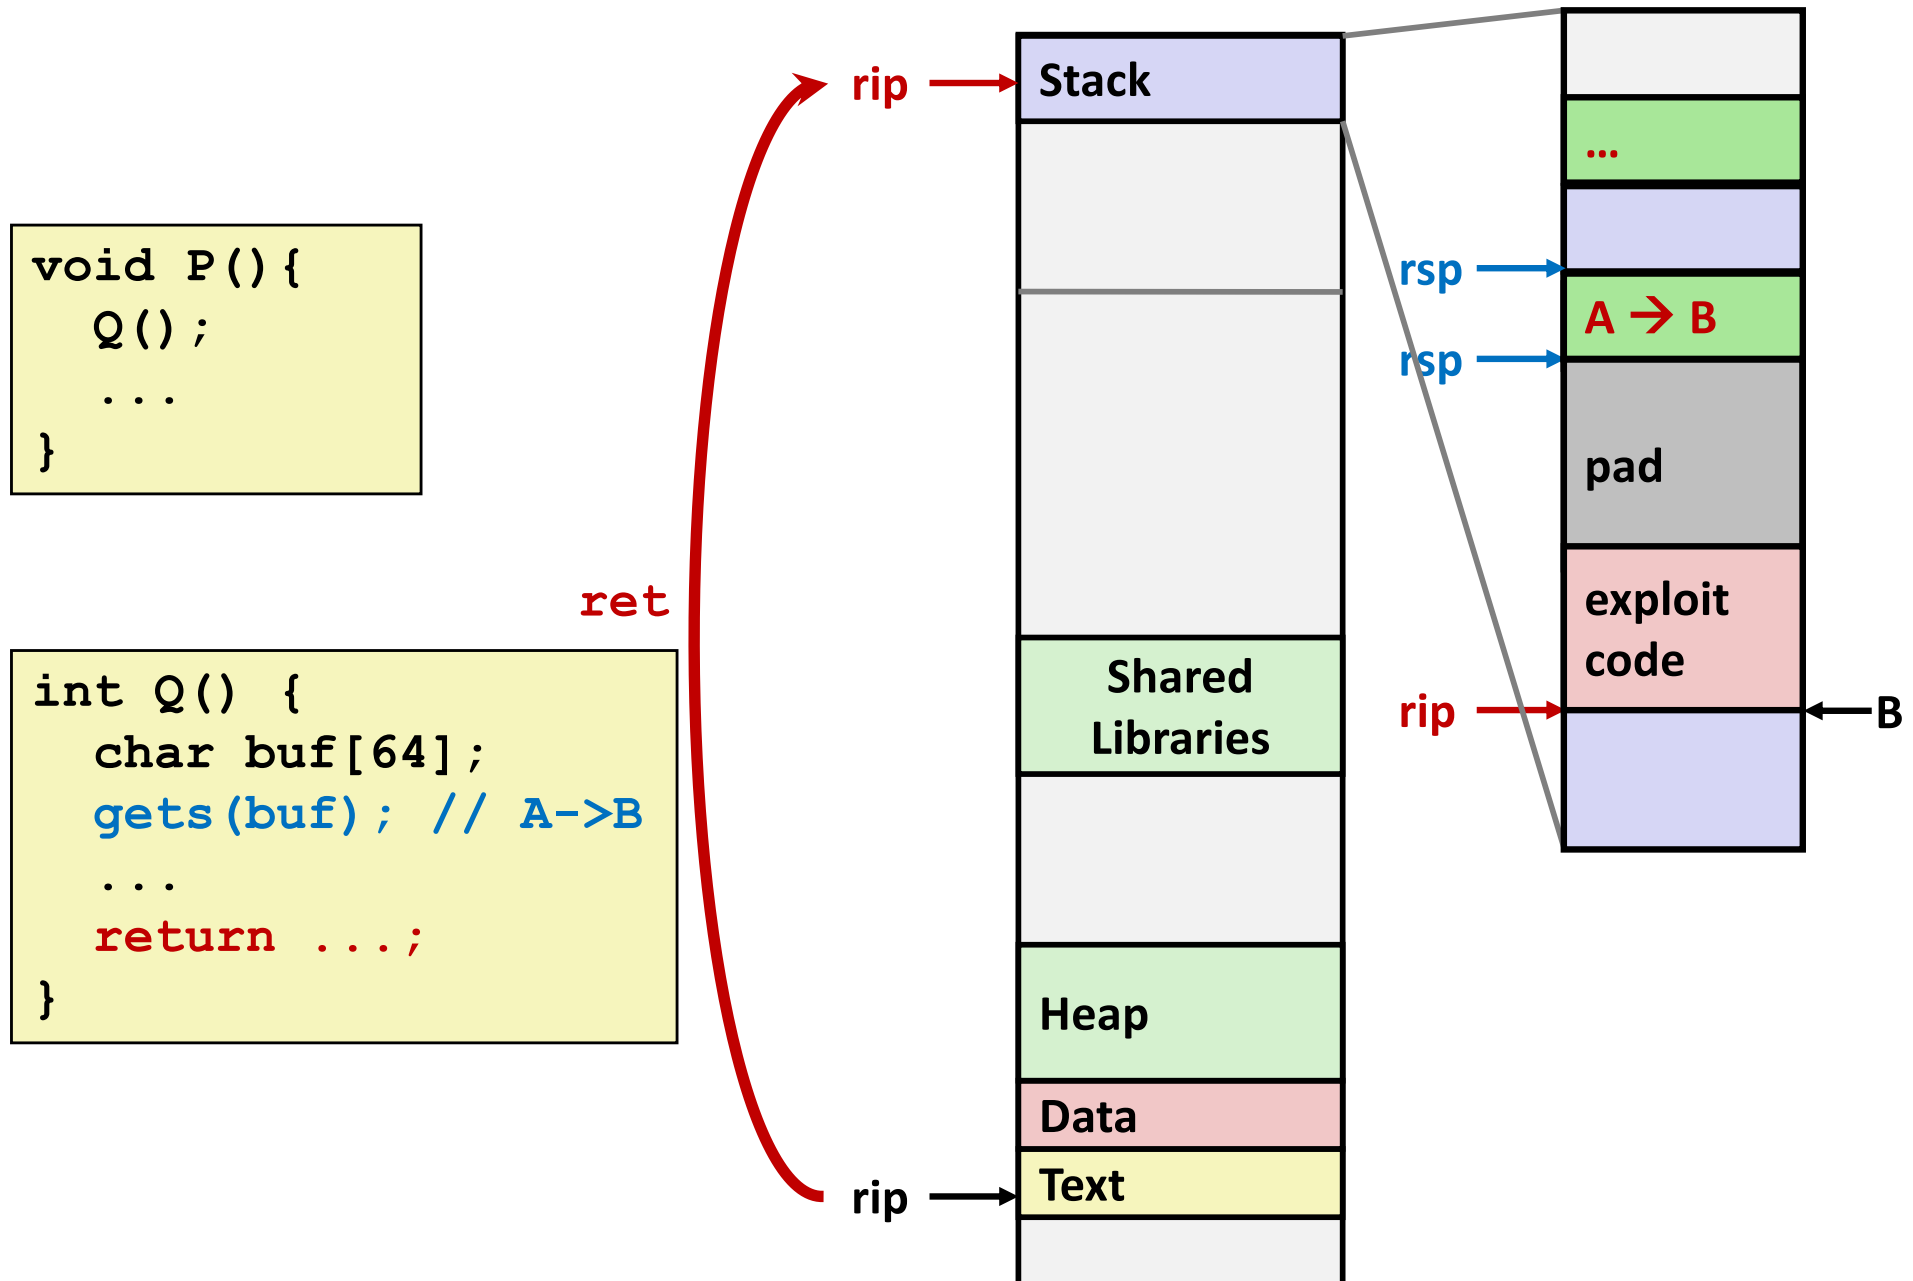
The attack string overwrites the return address to the desired target code.

Code Injection Attacks



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes `ret`, will jump to exploit code

How Does The Attack Code Execute?

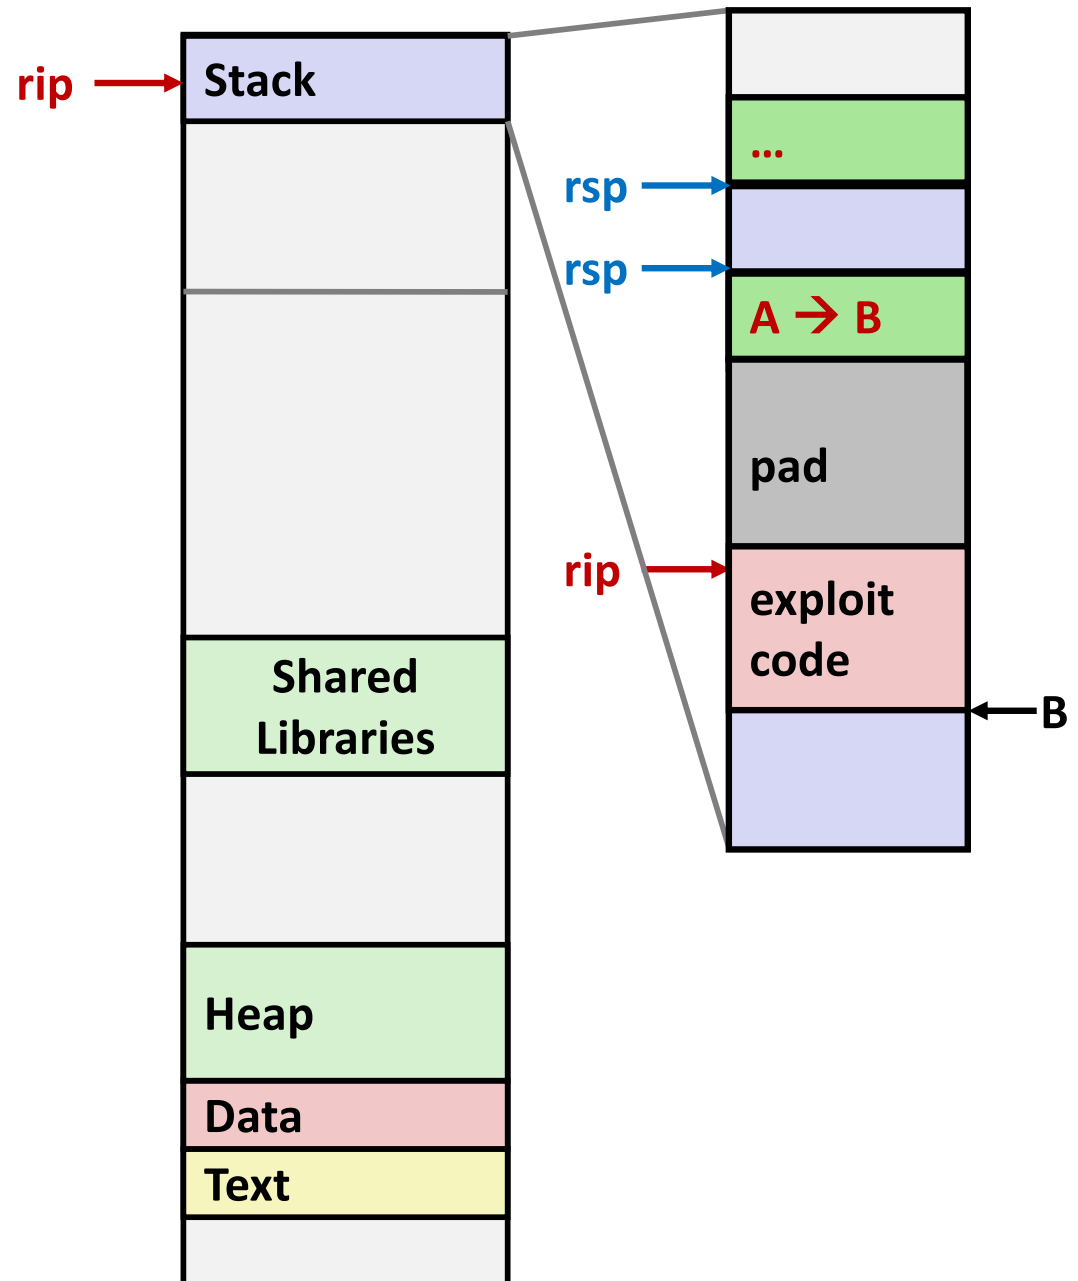


The Attacks Can Be Hidden

```
void P() {  
    Q();  
    ...  
}
```

```
int Q() {  
    char buf[64];  
    gets(buf); // A->B  
    ...  
    return ...;  
}
```

As part of the exploit code,
move %rsp back to the next
return address then call return.

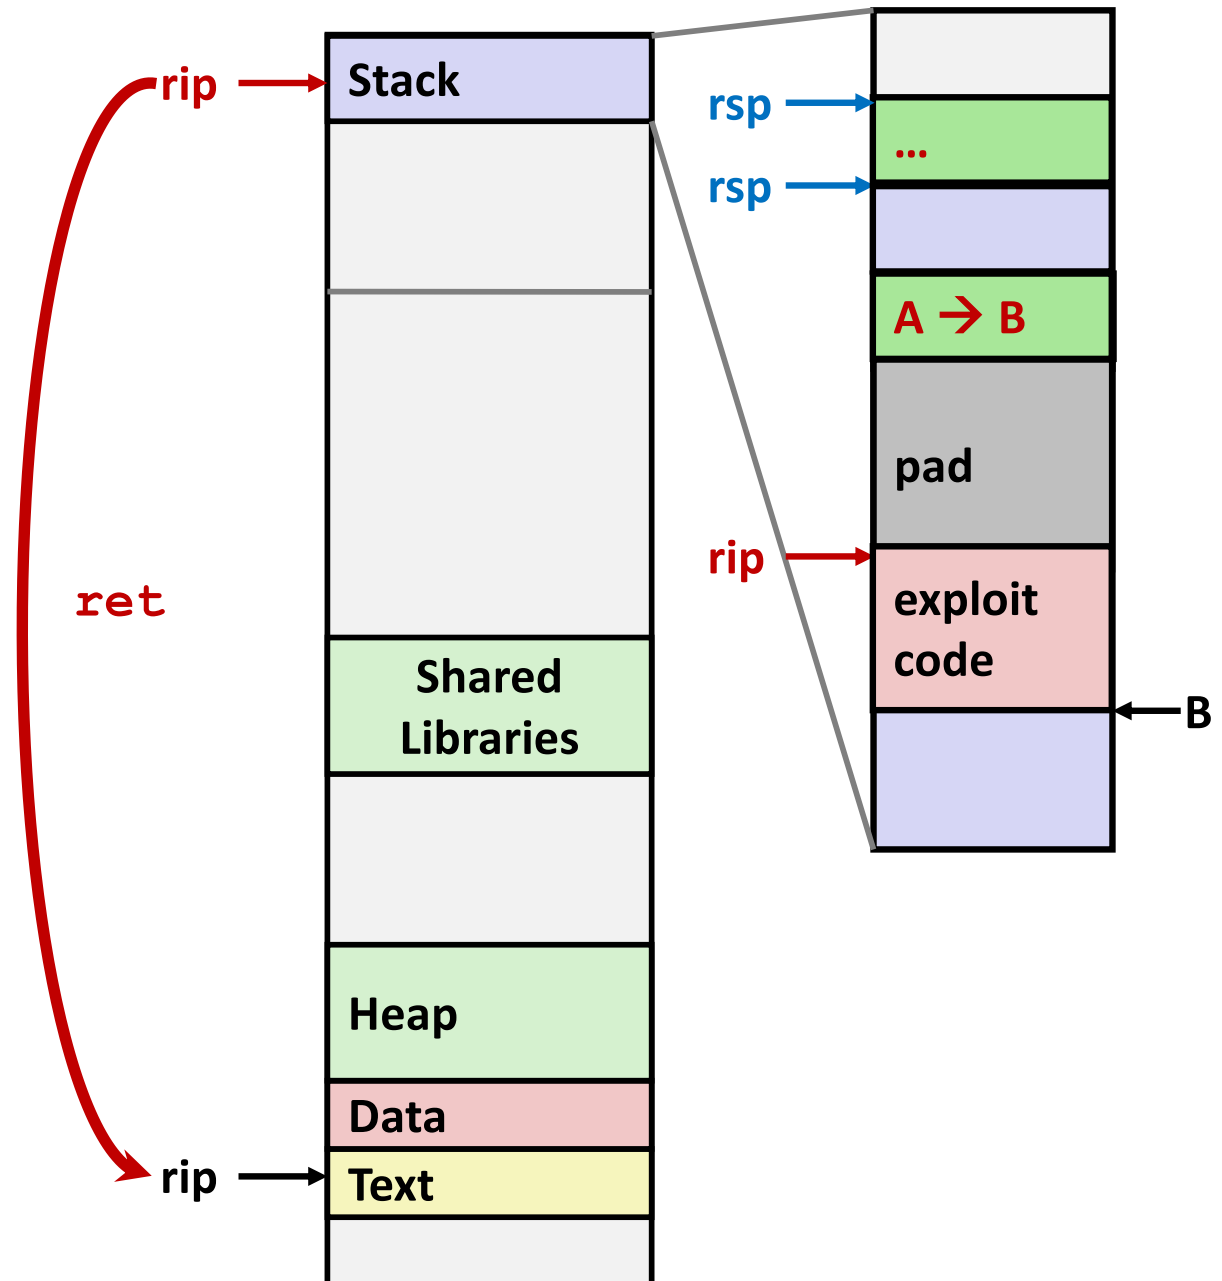


The Attacks Can Be Hidden

```
void P() {  
    Q();  
    ...  
}
```

```
int Q() {  
    char buf[64];  
    gets(buf); // A->B  
    ...  
    return ...;  
}
```

As part of the exploit code, move %rsp back to the next return address then call return.



Exploits Based on Buffer Overflows

- *Buffer overflow bugs can allow remote machines to execute arbitrary code on victim machines*
- **Distressingly common in real programs**
 - Programmers keep making the same mistakes ☹️
 - Recent measures make these attacks much more difficult
- **Examples across the decades**
 - Original “Internet worm” (1988)
 - “IM wars” (1999)
 - Twilight hack on Wii (2000s)
 - ... and many, many more
- **You will learn some of the tricks in attacklab**
 - Hopefully to convince you to never leave such holes in your programs!!

What to do about buffer overflow attacks

- **Avoid overflow vulnerabilities**
- **Employ system-level protections**
- **Have compiler use “stack canaries”**

1. Avoid Overflow Vulnerabilities in Code (!)


```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- For example, use library routines that limit string lengths
 - **fgets** instead of **gets**
 - **strncpy** instead of **strcpy**
 - Don't use **scanf** with **%s** conversion specification
 - Use **fgets** to read the string
 - Or use **%ns** where **n** is a suitable integer

2. System-Level Protections can help

■ Randomized stack offsets


- At start of program, allocate random amount of space on stack
- Shifts stack addresses for entire program
- Makes it difficult for hacker to predict beginning of inserted code
- E.g.: 5 executions of memory allocation code
 - Stack repositioned each time program executes



2. System-Level Protections can help

■ Nonexecutable code segments

- In traditional x86, can mark region of memory as either “read-only” or “writeable”
 - Can execute anything readable
- X86-64 added explicit “execute” permission
- Stack marked as non-executable



Any attempt to execute this code will fail

3. Stack Canaries can help

■ Idea

- Place special value (“canary”) on stack just beyond buffer
- Check for corruption before exiting function

■ GCC Implementation

- It is the default on some system; on linuxlabs, I enabled it using **-fstack-protector-all**

```
unix> ./bufdemo-sp  
Type a string: 0123456  
0123456
```

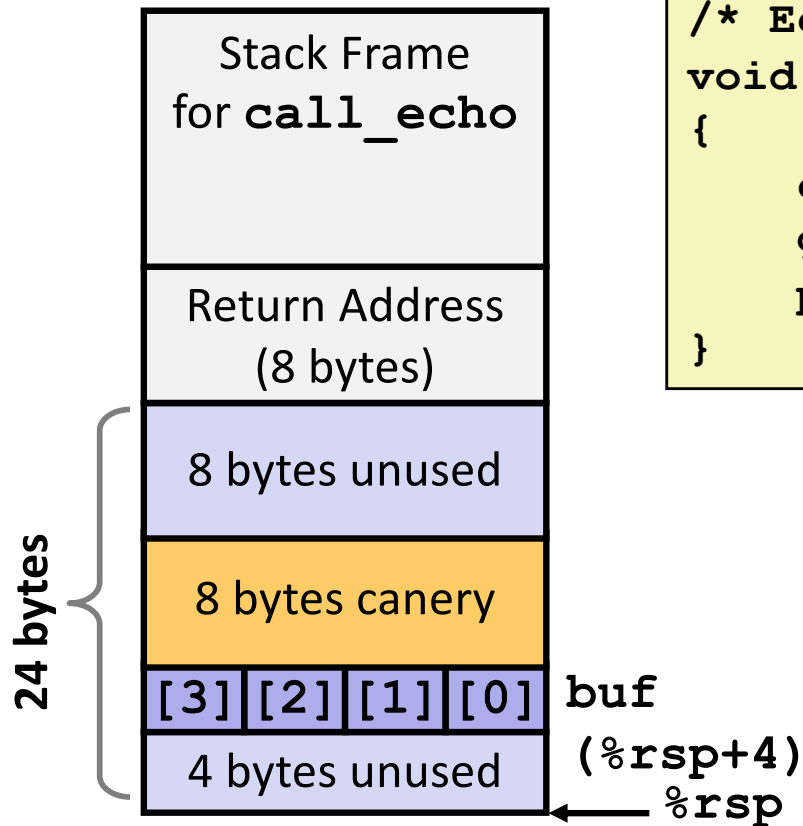
```
unix> ./bufdemo-sp  
Type a string: 01234567  
*** stack smashing detected ***
```

Protected Buffer Disassembly

echo:

```
400739:  sub    $0x18,%rsp
40073d:  mov     %fs:0x28,%rax
400746:  mov     %rax,0x8(%rsp)
40074b:  xor     %eax,%eax
40074d:  lea     0x4(%rsp),%rdi
400752:  callq   4006c6 <gets>
400757:  lea     0x4(%rsp),%rdi
40075c:  callq   400560 <puts@plt>
400761:  mov     0x8(%rsp),%rax
400766:  xor     %fs:0x28,%rax
40076f:  jne     400776 <echo+0x3d>
400771:  add     $0x18,%rsp
400775:  retq
400776:  callq   400570 <__stack_chk_fail@plt>
```

Setting Up Canary




```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

Before call to gets

```
echo:
    . . .
    movq    %fs:40, %rax    # Get canary
    movq    %rax, 8(%rsp)  # Place on stack
    xor     %eax, %eax      # Erase canary
    . . .
```

Checking Canary



```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

Input: 01234

After call to gets

echo:

```
. . .  
movq    8(%rsp), %rax    # Retrieve from stack  
xorq    %fs:40, %rax     # Compare to canary  
jne     .L6              # If not equal, jump  
                        # to __stack_chk_fail
```


Return-Oriented Programming Attacks

■ Challenge (for hackers)

- Stack randomization makes it hard to predict buffer location
- Marking stack nonexecutable makes it hard to insert binary code


■ Alternative Strategy

- Use existing code
 - E.g., library code from `stdlib`
- String together fragments to achieve overall desired outcome
- *Does not overcome stack canaries*

■ Construct program from *gadgets*

- Sequence of instructions ending in `ret`
 - Encoded by single byte `0xc3`
- Code positions fixed from run to run
- Code is executable

ROP Execution



- Trigger with `ret` instruction
 - Will start executing Gadget 1
- Final `ret` in each gadget will start next one

Gadget Example #1

```
long ab_plus_c  
    (long a, long b, long c)  
{  
    return a*b + c;  
}
```

```
00000000004004d0 <ab_plus_c>:  
4004d0: 48 0f af fe  imul %rsi,%rdi  
4004d4: 48 8d 04 17  lea (%rdi,%rdx,1),%rax  
4004d8: c3           retq
```

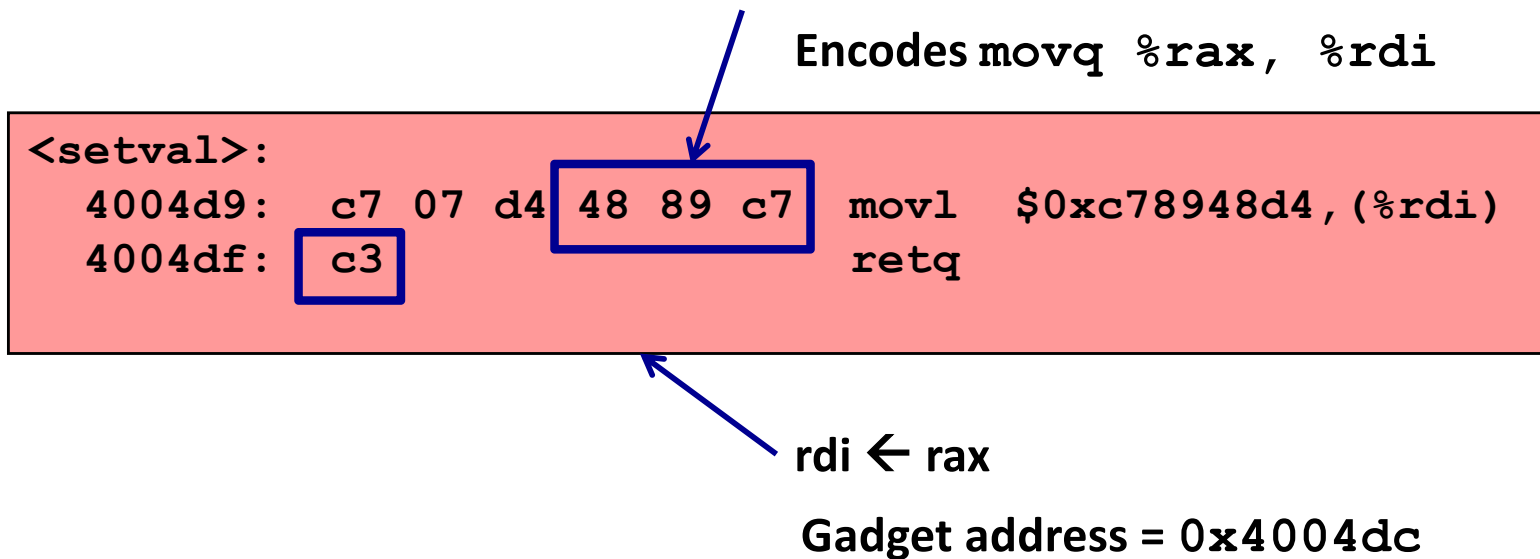
$\text{rax} \leftarrow \text{rdi} + \text{rdx}$

Gadget address = 0x4004d4

- Use tail end of existing functions


Gadget Example #2

```
void setval(unsigned *p) {  
    *p = 3347663060u;  
}
```



- Repurpose byte codes

ROP Execution



- Trigger with `ret` instruction
 - Will start executing Gadget 1
- Final `ret` in each gadget will start next one

Example: the original Internet worm (1988)

■ Exploited a few vulnerabilities to spread

- Early versions of the finger server (fingerd) used `gets ()` to read the argument sent by the client:
 - `finger droh@cs.cmu.edu`
- Worm attacked fingerd server by sending phony argument:
 - `finger "exploit-code padding new-return-address"`
 - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.


■ Once on a machine, scanned for other machines to attack

- invaded ~6000 computers in hours (10% of the Internet ☺)
 - see June 1989 article in *Comm. of the ACM*
- the young author of the worm was prosecuted...
- and CERT (Computer Emergency Response Team) was formed

Example 2: IM War

■ July, 1999

- Microsoft launches MSN Messenger (instant messaging system).
- Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



IM War (cont.)

■ August 1999

- Microsoft and AOL begin the IM war:
 - AOL changes server to disallow Messenger clients
 - Microsoft makes changes to clients to defeat AOL changes
 - At least 13 such skirmishes
- Mysteriously, Messenger clients can no longer access AIM servers
- What was really happening?
 - AOL had discovered a buffer overflow bug in their own AIM clients
 - They exploited it to detect and block Microsoft: the exploit code returned a 4-byte signature (the bytes at some location in the AIM client) to server
 - Interesting recounting of events by an engineer at Microsoft at the time: <https://nplusemag.com/issue-19/essays/chat-wars/>

Date: Wed, 11 Aug 1999 11:30:57 -0700 (PDT)
From: Phil Bucking <philbucking@yahoo.com>
Subject: AOL exploiting buffer overrun bug in their own software!
To: rms@pharlap.com

Mr. Smith,

I am writing you because I have discovered something that I think you might find interesting because you are an Internet security expert with experience in this area. I have also tried to contact AOL but received no response.

I am a developer who has been working on a revolutionary new instant messaging client that should be released later this year.

...

It appears that the AIM client has a buffer overrun bug. By itself this might not be the end of the world, as MS surely has had its share. **But AOL is now *exploiting their own buffer overrun bug* to help in its efforts to block MS Instant Messenger.**

....

Since you have significant credibility with the press I hope that you can use this information to help inform people that behind AOL's friendly exterior they are nefariously compromising peoples' security.

Sincerely,
Phil Bucking
Founder, Bucking Consulting
philbucking@yahoo.com

It was later determined that this email originated from within Microsoft!

Aside: Worms and Viruses

- **Worm: A program that**
 - Can run by itself
 - Can propagate a fully working version of itself to other computers

- **Virus: Code that**
 - Adds itself to other programs (i.e., infect another software)
 - Does not run independently

- **Both are (usually) designed to spread among computers and to wreak havoc**

Program Optimization

B&O Readings: 5

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Today

- **Overview**
- **Machine-Independent Optimizations**
 - Code motion/precomputation
 - Strength reduction
 - Common Subexpression Elimination
 - Removing unnecessary procedure calls
- **Optimization Blockers**
 - Procedure calls
 - Memory aliasing
- **Putting it all together**

Performance Realities

- *There's more to performance than asymptotic complexity*
- **Constant factors matter!**
 - Easily see 10:1 performance range depending on how code is written
 - Must optimize at multiple levels:
 - algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs are compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Compilers do 2 things:


■ Code generation

- Translate high-level language to machine instructions, 1 statement at a time

■ Optimization

- Preserve meaning but improve performance
- Active research area, but some standard optimizations

Bird's Eye View of Compiler Optimizations



- C: your ordinary C programs
- SSA IR: Static Single Assignment Intermediate Representation
- Obj: object files that are ISA dependent

Optimizing Compilers

- **Provide efficient mapping of program to machine**
 - register allocation
 - code selection and ordering (scheduling)
 - dead code elimination
 - eliminating minor inefficiencies
- **Don't (usually) improve asymptotic efficiency**
 - up to programmer to select best overall algorithm
 - big-O savings are (often) more important than constant factors
 - but constant factors also matter
- **Have difficulty overcoming “optimization blockers”**
 - potential memory aliasing
 - potential procedure side-effects

Generally Useful Optimizations

- **Optimizations that you or the compiler should do regardless of processor / compiler**
 - Code Motion
 - Strength Reduction
 - Common Subexpression Elimination
 - Tail Recursion Elimination

Code Motion

- Reduce frequency with which computation performed
 - If it will always produce same result
 - Especially moving code out of loop

```
void set_row(double *a, double *b,  
             long i, long n)  
{  
    long j;  
    for (j = 0; j < n; j++)  
        a[n*i+j] = b[j];  
}
```



```
long j;  
int ni = n*i;  
for (j = 0; j < n; j++)  
    a[ni+j] = b[j];
```

Compiler-Generated Code Motion (-O1)

```
void set_row(double *a, double *b,  
            long i, long n)  
{  
    long j;  
    for (j = 0; j < n; j++)  
        a[n*i+j] = b[j];  
}
```

```
long j;  
long ni = n*i;  
double *rowp = a+ni;  
for (j = 0; j < n; j++)  
    *rowp++ = b[j];
```

```
set_row:  
    testq    %rcx, %rcx                # Test n  
    jle      .L1                       # If 0, goto done  
    imulq    %rcx, %rdx                # ni = n*i  
    leaq     (%rdi,%rdx,8), %rdx        # rowp = A + ni*8  
    movl     $0, %eax                  # j = 0  
    .L3:  
    movsd    (%rsi,%rax,8), %xmm0      # t = b[j]  
    movsd    %xmm0, (%rdx,%rax,8)      # M[A+ni*8 + j*8] = t  
    addq     $1, %rax                  # j++  
    cmpq     %rcx, %rax                # j:n  
    jne      .L3                       # if !=, goto loop  
    .L1:  
    rep ; ret                          # done:
```

Reduction in Strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide
 - $16 * x \rightarrow x \ll 4$
 - Utility machine dependent
 - Depends on cost of multiply or divide instruction
- Recognize sequence of products

```
for (i = 0; i < n; i++) {  
    int ni = n*i;  
    for (j = 0; j < n; j++)  
        a[ni + j] = b[j];  
}
```



```
int ni = 0;  
for (i = 0; i < n; i++) {  
    for (j = 0; j < n; j++)  
        a[ni + j] = b[j];  
    ni += n;  
}
```

Share Common Subexpressions

- Reuse portions of expressions
- GCC will do this with `-O1`

```
/* Sum neighbors of i,j */
up =    val[(i-1)*n + j ];
down =  val[(i+1)*n + j ];
left =  val[i*n      + j-1];
right = val[i*n      + j+1];
sum = up + down + left + right;
```

3 multiplications: $i*n$, $(i-1)*n$, $(i+1)*n$

```
leaq    1(%rsi), %rax    # i+1
leaq    -1(%rsi), %r8    # i-1
imulq   %rcx, %rsi       # i*n
imulq   %rcx, %rax       # (i+1)*n
imulq   %rcx, %r8       # (i-1)*n
addq    %rdx, %rsi       # i*n+j
addq    %rdx, %rax       # (i+1)*n+j
addq    %rdx, %r8       # (i-1)*n+j
```

```
long inj = i*n + j;
up =    val[inj - n];
down =  val[inj + n];
left =  val[inj - 1];
right = val[inj + 1];
sum = up + down + left + right;
```

1 multiplication: $i*n$

```
imulq   %rcx, %rsi       # i*n
addq    %rdx, %rsi       # i*n+j
movq    %rsi, %rax       # i*n+j
subq    %rcx, %rax       # i*n+j-n
leaq    (%rsi,%rcx), %rcx # i*n+j+n
```

Tail Recursion Elimination

- Varies across languages
- Optional in C

```
int fact(int n) {  
    if(n <= 0)  
        return 1;  
    else {  
        int n1_fact = fact(n-1);  
        return n1_fact * n;  
    }  
}
```

-O1
→

```
fact:  
    pushq    %rbx  
    movl     %edi, %ebx  
    movl     $1, %eax  
    testl    %edi, %edi  
    jle      .L2  
    leal     -1(%rdi), %edi  
    call     fact  
    imull    %ebx, %eax  
  
.L2:  
    popq     %rbx  
    ret
```

-O2
↘

```
fact:  
    testl    %edi, %edi  
    movl     $1, %eax  
    jle      .L2  
  
.L3:  
    imull    %edi, %eax  
    subl     $1, %edi  
    jne      .L3  
  
.L2:  
    ret
```

With -O2:

- Compiler doesn't create stack frame for tail recursion
- No call to fact, just a jmp
- No longer need to store away caller-saved registers!