



Linking

B&O Readings: 7

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Today

- **Linking**
 - What does a linker do
 - Motivation --- why linkers?
 - How it works
- Static vs Dynamic libraries
- Case study: Library interpositioning

x86-64 Linux Memory Layout

not drawn to scale

0x00007FFFFFFF

■ Stack

- Runtime stack (8MB limit)
- E. g., local variables

■ Heap


- Dynamically allocated as needed
- When call `malloc()`, `calloc()`, `new()`

■ Data

- Statically allocated data
- E.g., global vars, **static** vars, string constants
- `.bss`, `.data`, `.rodata`

■ Text / Shared Libraries

- Executable machine instructions
- Read-only



Example C Program

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);
    return val;
}
```

main.c

```
int sum(int *a, int n)
{
    int i, s = 0;

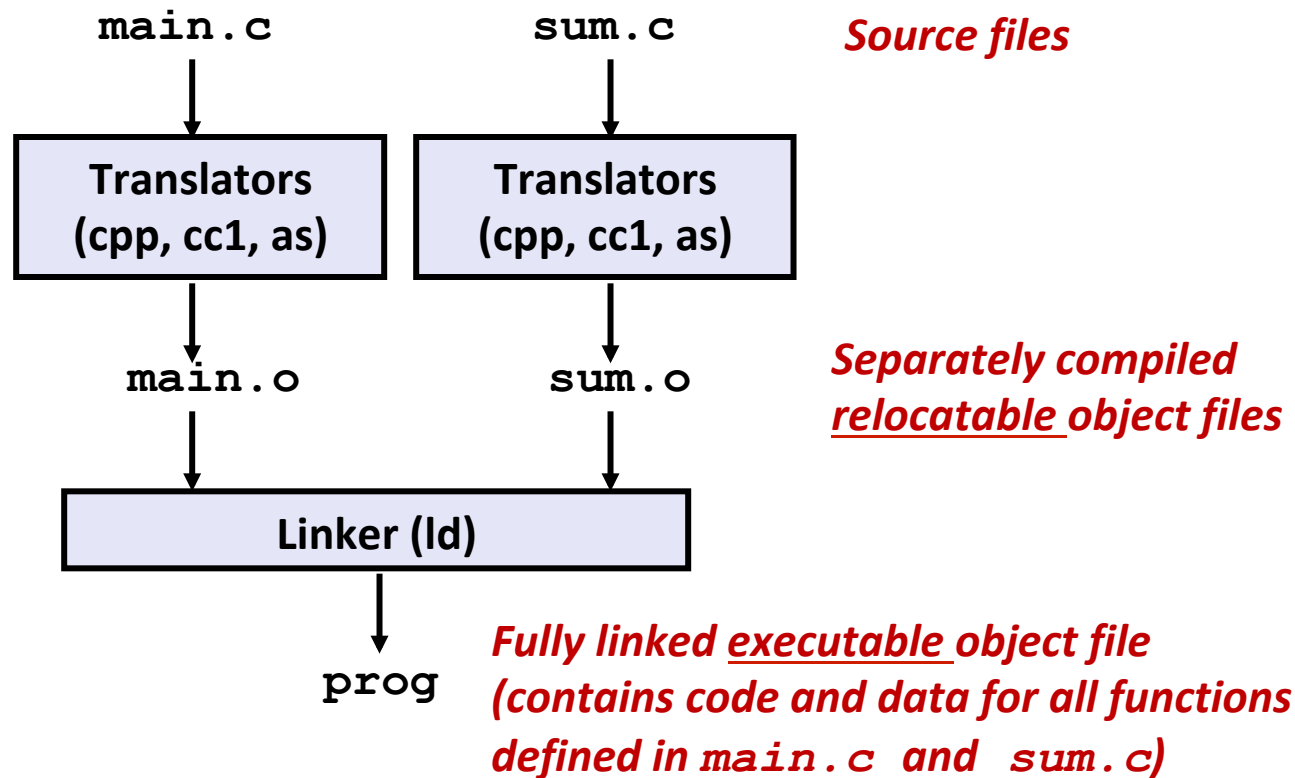
    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
```

sum.c

Static Linking

■ Programs are translated and linked using a *compiler driver*:

- `linux> gcc -Og -o prog main.c sum.c`
- `linux> ./prog`



Q: Where Does Each Variable Live?

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);
    return val;
}
```

main.c

```
int sum(int *a, int n)
{
    int i, s = 0;

    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
```

sum.c

- main: .text (function)
 - array: .data (global variable)
 - val: stack (local variable)
 - argc, argv: stack (args)
- sum: .text (function)
 - i, s: stack (local variables)
 - a, n: stack (args)
- When we compile main, how do we know where sum lives?
 - Where is array located?

Linker, at a Glance

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);
    return val;
}
```

main.c

```
int sum(int *a, int n)
{
    int i, s = 0;

    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
```

sum.c

- For every function / variable defined by your program that live in .text and .data, there is a **symbol** corresponds to it (naming the symbol = referring to its address).
- It's linker's job to figure out where each function / variable defined by your program goes and resolve references to these symbols.

Why Linkers?


■ Reason 1: Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass(/mess)
- Can build libraries of common functions
 - e.g., Math library, standard C library

Why Linkers? (cont)

■ Reason 2: Efficiency

- **Time:** Separate compilation
 - Change one source file, compile, and then relink
 - No need to recompile other source files



■ **Space:** Libraries

- Common functions can be aggregated into a single file...
- Executable files contain only code for the functions they actually use

Why Should You Understand How Linking Works

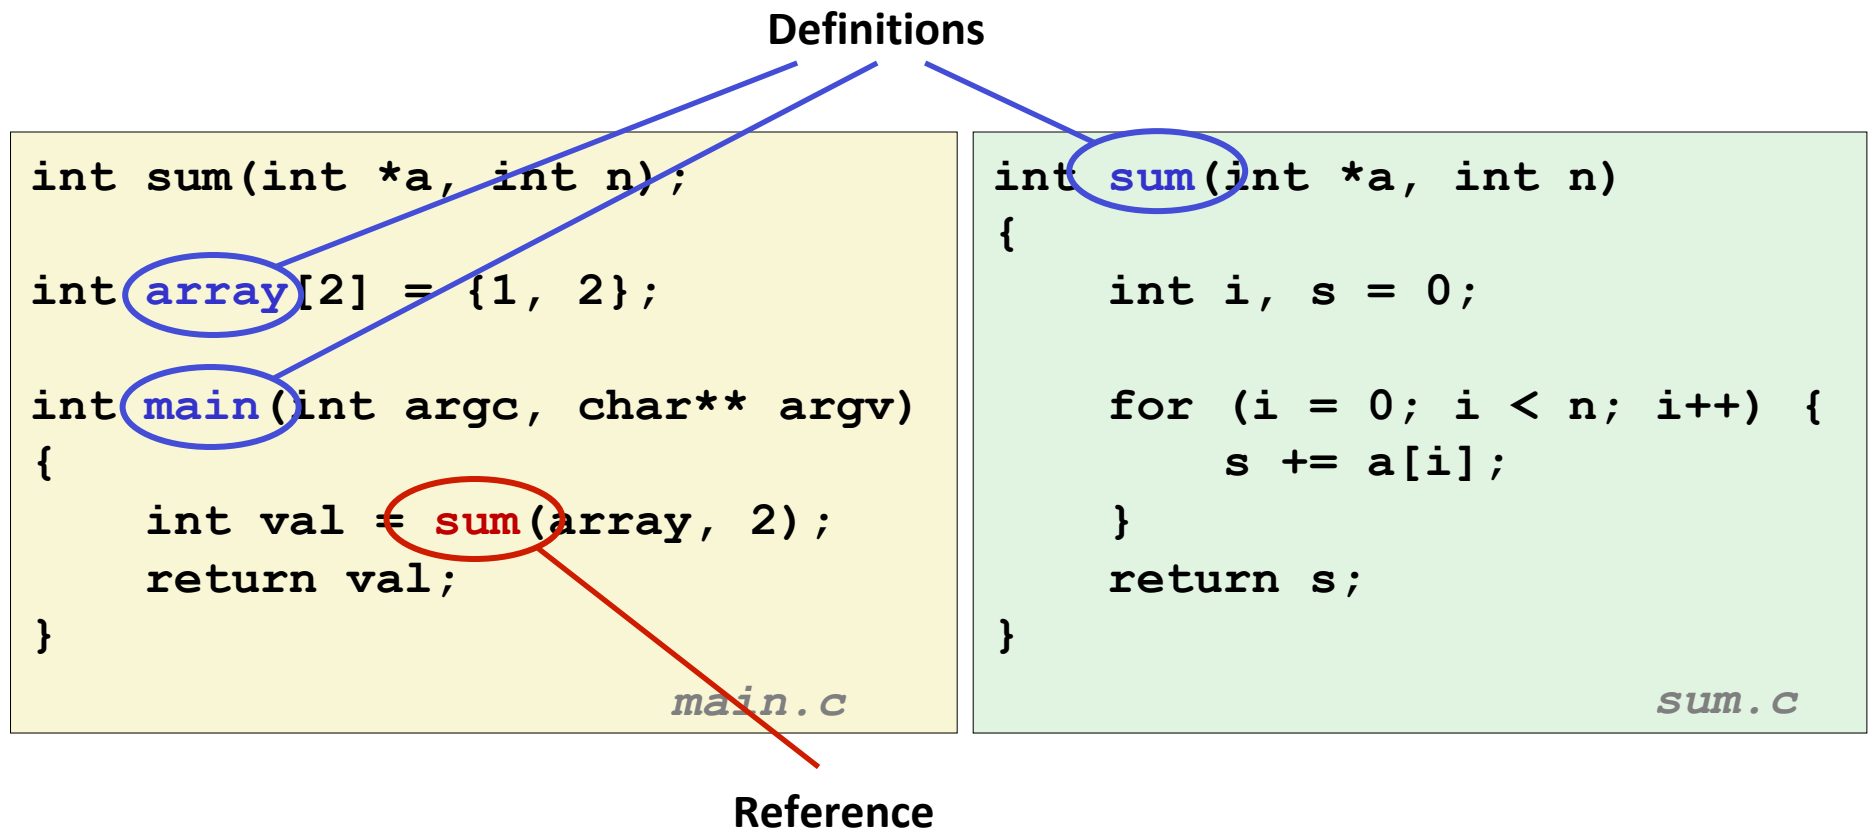
- Understanding linkers help you build large programs.
- Understanding linkers help you decipher whether an "compiling error" comes from the compiler vs. the linker errors.
- Understanding linkers help you avoid difficult-to-debug programming errors.
- Understanding linkers help you understand how scoping rules are implemented.

What Do Linkers Do? (1)

■ Step 1. Symbol resolution

- Programs define and reference *symbols* (variables, functions):
 - `void swap() {...} /* define symbol swap */`
 - `swap(); /* reference symbol a */`
 - `int *xp = &x; /* define symbol xp, reference x */`
- Symbol definitions are stored (by compiler) in *symbol table*
 - An array of structs
 - Each entry: name, size, location of symbol
- *Linker associates each symbol reference with exactly one symbol definition.*

Symbols in Example C Program



What Do Linkers Do? (2)

■ Step 2. Relocation

- **Merges** separate code and data sections into single sections (see picture on slide 15)
- **Relocates** symbols from their relative locations in the .o files to their final absolute memory locations in the executable
- **Updates** all references to these symbols to reflect their new positions

Three Kinds of Object Files (Modules)

■ Relocatable object file (.o file)

- Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source (.c) file

■ Executable object file (a.out file)

- Contains code and data in a form that can be copied directly into memory and then executed.

■ Shared object file (.so file)

- Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
- Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- **Standard binary format for object files**
- **One unified format for**
 - Relocatable object files (`.o`),
 - Executable object files (`a.out`)
 - Shared object files (`.so`)
- **Generic name: ELF binaries**

ELF Object File Format

■ Elf header

- Word size, byte ordering, file type (.o, exec, .so), machine type, etc.

■ .text section

- Code

■ .rodata section

- Read only data: jump tables, ...

■ .data section

- Initialized global variables

■ .bss section

- Uninitialized global variables or ones initialized with 0
- “Block Started by Symbol” (“**Better Save Space**”)
- Has section header but occupies no space

ELF header
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

0

ELF Object File Format (cont.)

- **.symtab section**
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- **.rel.text section**
 - Relocation info for **.text** section
 - Addresses of instructions that will need to be modified in the executable
 - Instructions for modifying.
- **.rel.data section**
 - Relocation info for **.data** section
 - Addresses of pointer data that will need to be modified in the merged executable
- **.debug section**
 - Info for symbolic debugging (**gcc -g**)
- **Section header table**
 - Offsets and sizes of each section

ELF header	0
.text section	
.rodata section	
.data section	
.bss section	
.symtab section	
.rel.txt section	
.rel.data section	
.debug section	
Section header table	

Linker Symbols

■ Global symbols

- Symbols defined by module m that can be referenced by other modules.
- E.g.: non-**static** C functions and non-**static** global variables.

■ External symbols

- Global symbols that are referenced by module m but defined by some other module.

■ Local symbols

- Symbols that are defined and referenced exclusively by module m .
- E.g.: C functions and global variables defined with the **static** attribute.
- **Local linker symbols are *not* local program variables**

Digression: Static Variables

```
static int x = 99;

int foo()
{
    static int y = 0;
    y++;
    printf("%d\n", y);
}

int bar()
{
    x++;
    printf("%d\n", x);
}
```

- The non-local static variable (e.g., `x`) is only visible within this file. Reference to it in a different file will trigger an error from the linker.
- The increment of `y` takes effect each time `foo` is called!

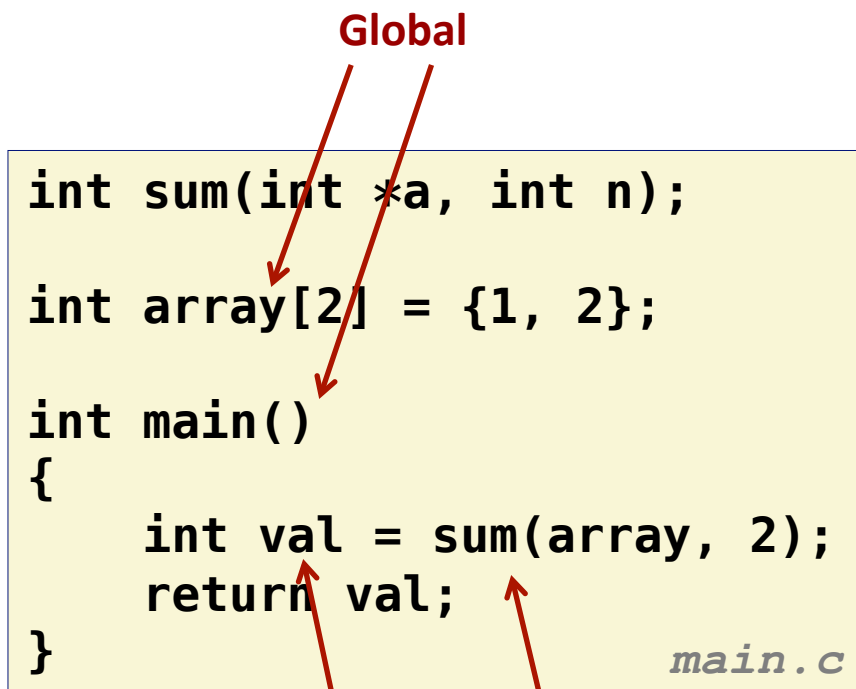
What does the `printf` in `foo` print ? 1, 2, 3, ...

What does the `printf` in `bar` print ? 100

Step 1: Symbol Resolution

```
int sum(int *a, int n);
int array[2] = {1, 2};

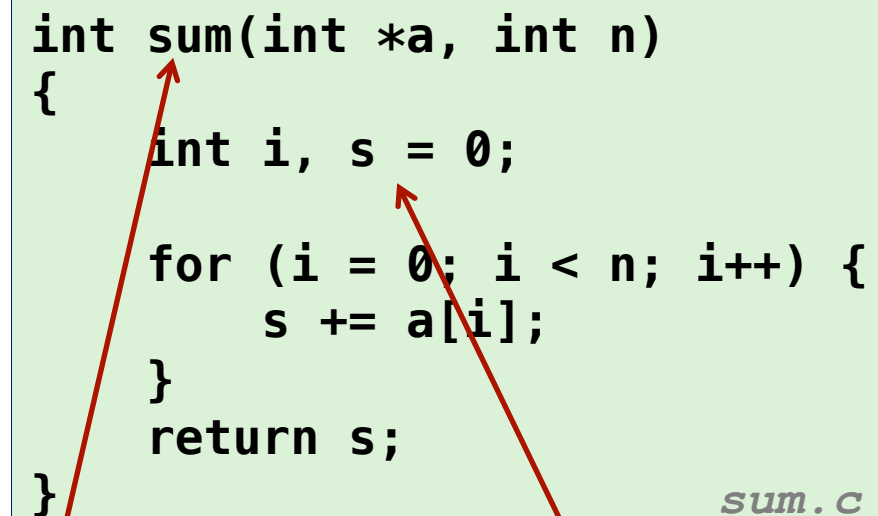
int main()
{
    int val = sum(array, 2);
    return val;
}
main.c
```



Linker knows
nothing of val

External

```
int sum(int *a, int n)
{
    int i, s = 0;
    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
sum.c
```



Global

Linker knows
nothing of i or s

Symbol Identification

Which of the following names will be in the symbol table of `symbols.o`?

Names:

`symbols.c`:

```
int time;

int foo(int a) {
    int b = a + 1;
    return b;
}

int main(int argc,
          char* argv[]) {
    printf("%d\n", foo(5));
    return 0;
}
```

- `time`
- `foo`
- `a`
- `argc`
- `argv`
- `b`
- `main`
- `printf`
- `"%d\n"`

Local Symbols

```
static int x = 99;
int f()
{
    static int x = 0;
    return x++;
}

int g()
{
    static int x = 1;
    return x++;
}
```

- The local static variable shadows the non-local static variable with the same name.
- Local non-static C variables vs. local static C variables
 - local non-static C variables: stored on the stack (linker knows nothing about these)
 - local static C variables: stored in either `.bss`, or `.data`

The linker allocates space in `.data` for each definition of `x` and creates local symbols in the symbol table with unique names, e.g., `x.1`, `x.2`, `x.3` ... etc.

Resolving Symbols

Which symbols are global? Which are external? Which are local?

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

main.c

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

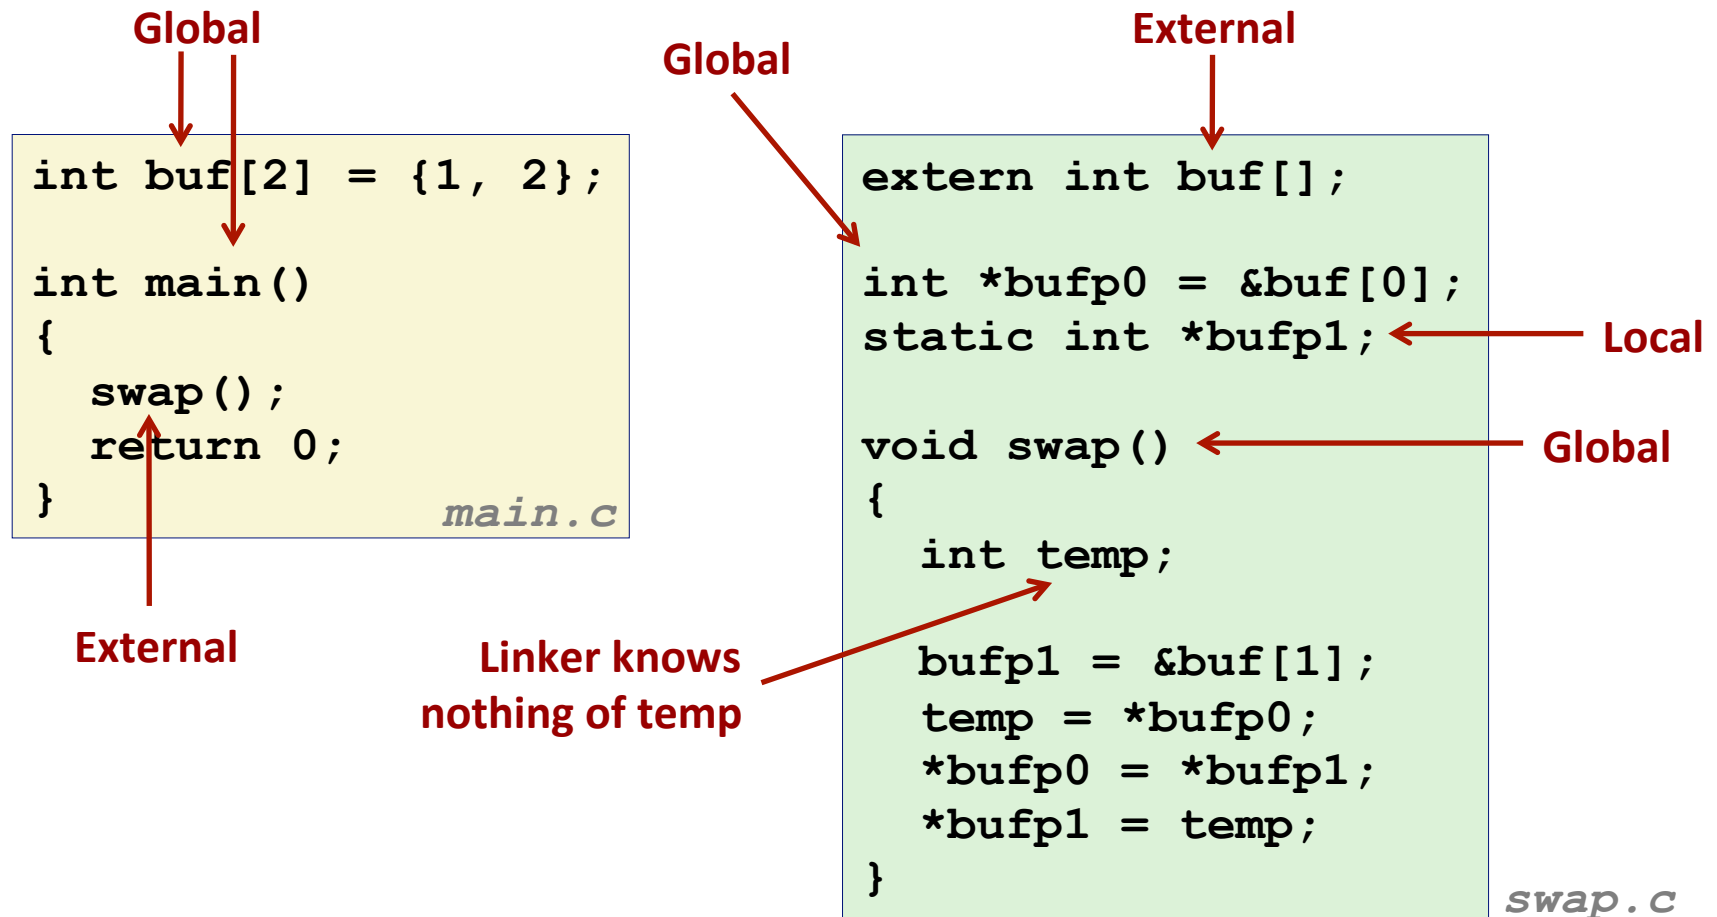
    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.c




Resolving Symbols

Which symbols are global? Which are external? Which are local?



How Linker Resolves Duplicate Symbol Definitions



- Program symbols are either ***strong*** or ***weak***
 - ***Strong***: procedures and initialized globals
 - ***Weak***: uninitialized globals

Linker's Symbol Rules

- **Rule 1: Multiple strong symbols are not allowed**
 - Each item can be defined only once
 - Otherwise: Linker error
- **Rule 2: Given a strong symbol and multiple weak symbols, choose the strong symbol**
 - References to the weak symbol resolve to the strong symbol
- **Rule 3: If there are multiple weak symbols, pick an arbitrary one**
 - Can override this with `gcc -fno-common`

Linker Puzzles

```
int x;  
p1() {}
```

```
p1() {}
```

```
int x;  
p1() {}
```

```
int x;  
p2() {}
```

```
int x;  
int y;  
p1() {}
```

```
double x;  
p2() {}
```

```
int x=7;  
int y=5;  
p1() {}
```

```
double x;  
p2() {}
```

```
int x=7;  
p1() {}
```

```
int x;  
p2() {}
```

What happens when you compile each of these pair of files?



Linker Puzzles

```
int x;  
p1() {}
```

```
p1() {}
```

Link time error: two strong symbols (**p1**)

```
int x;  
p1() {}
```

```
int x;  
p2() {}
```

References to **x** will refer to the same uninitialized int. Is this what you really want?

```
int x;  
int y;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to **x** in **p2** might overwrite **y**!
Evil!

```
int x=7;  
int y=5;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to **x** in **p2** will overwrite **y**!
Nasty!

```
int x=7;  
p1() {}
```

```
int x;  
p2() {}
```

References to **x** will refer to the same initialized variable.

What happens when you compile each of these pair of files?



Type Mismatch Example

```
long int x; /* Weak symbol */

int main(int argc,
          char *argv[]) {
    printf("%ld\n", x);
    return 0;
}
```

mismatch-main.c

```
/* Global strong symbol */
double x = 3.14;
```

mismatch-variable.c

- Compiles without any errors or warnings
- What gets printed?

Bit representation for double 3.14 interpreted as a long value.

Global Variables

- **Avoid if you can**
- **Otherwise**
 - Use **static** if you can
 - Initialize if you define a global variable
 - Use **extern** if you reference an external global variable
 - Treated as weak symbol
 - But also causes linker error if no strong symbol defined in some file.
 - Unfortunately, you still don't get an error if the extern declares the wrong type!

Use of extern in .h Files (#1)

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

global.h

```
extern int g;
int f();
```

c2.c

```
#include <stdio.h>
#include "global.h"

int g = 0;

int main(int argc, char argv[]) {
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

Use of extern in .h Files (#1)

c1.c

```
#include "global.h"  
  
int f() {  
    return g+1;  
}
```

```
extern int g;  
int f();
```

c2.c

```
#include <stdio.h>  
#include "global.h"  
  
int g = 0;  
  
int main(int argc, char argv[]) {  
    int t = f();  
    printf("Calling f yields %d\n", t);  
    return 0;  
}
```

A compiler warning on mismatched type will trigger.

```
extern int g;  
int f();
```

The compiler pre-processor simply dump dump the content of header file into the C file that includes it!

Use of .h Files (#2)

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

c2.c

```
#define INITIALIZE
#include <stdio.h>
#include "global.h"

int main(int argc, char** argv) {
    if (init)
        // do something, e.g., g=31;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

global.h

```
#ifndef INITIALIZE
    int g = 23;
    static int local_init = 1;
#else
    extern int g;
    static int local_init = 0;
#endif
```

Use of .h Files (#2)

c1.c

```
#include "global.h"  
  
int f() {  
    return g+1;  
}
```

global.h

```
extern int g;  
static int local_init = 0;
```

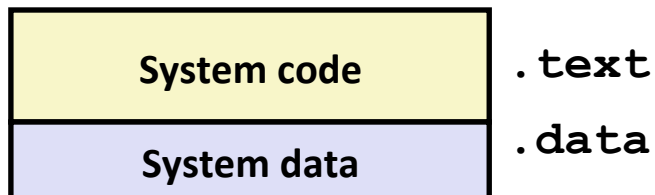
c2.c

```
#define INITIALIZE  
#include <stdio.h>  
#include "global.h"  
  
int main(int argc, char** argv) {  
    if (init)  
        // do something, e.g., g=31;  
    int t = f();  
    printf("Calling f yields %d\n", t);  
    return 0;  
}
```


```
int g = 23;  
static int local_init = 1;
```

Step 2: Relocation


Relocatable Object Files



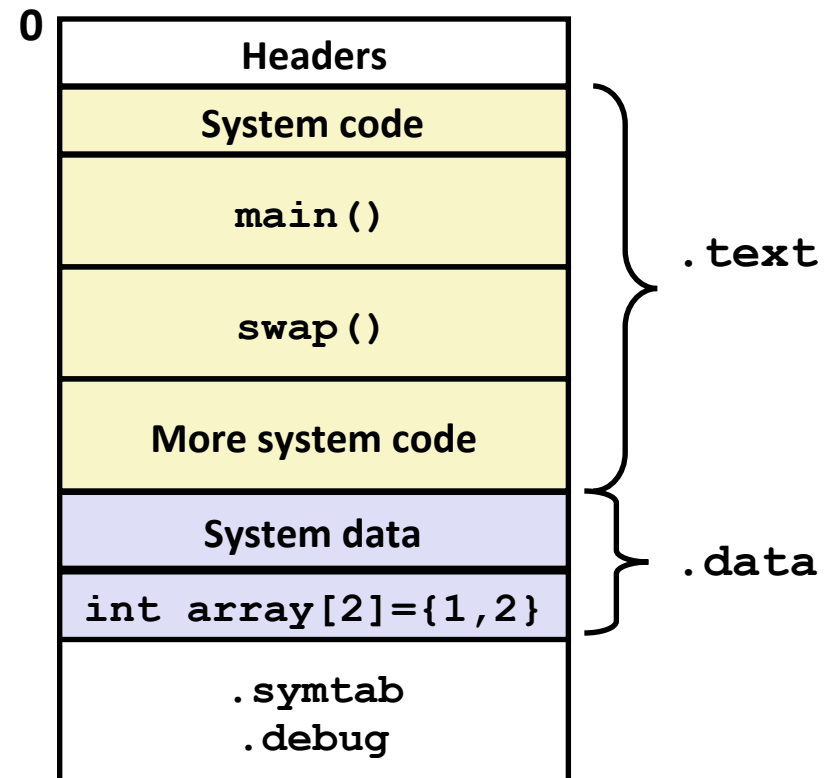
main.o



sum.o



Executable Object File



Relocation Entries

```
int array[2] = {1, 2};

int main()
{
    int val = sum(array, 2);
    return val;
}                                     main.c
```

```
0000000000000000 <main>:
 0:  48 83 ec 08                sub    $0x8,%rsp
 4:  be 02 00 00 00            mov    $0x2,%esi
 9:  bf 00 00 00 00            mov    $0x0,%edi    # %edi = &array
                        a: R_X86_64_32 array    # Relocation entry

 e:  e8 00 00 00 00            callq  13 <main+0x13> # sum()
                        f: R_X86_64_PC32 sum-0x4    # Relocation entry
13:  48 83 c4 08                add    $0x8,%rsp
17:  c3                        retq

                                           main.o
```

Relocated .text section

```
00000000004004d0 <main>:
 4004d0:      48 83 ec 08          sub    $0x8,%rsp
 4004d4:      be 02 00 00 00      mov    $0x2,%esi
 4004d9:      bf 18 10 60 00      mov    $0x601018,%edi # %edi = &array
 4004de:      e8 05 00 00 00      callq 4004e8 <sum>    # sum()
 4004e3:      48 83 c4 08          add    $0x8,%rsp
 4004e7:      c3                   retq

00000000004004e8 <sum>:
 4004e8:      b8 00 00 00 00      mov    $0x0,%eax
 4004ed:      ba 00 00 00 00      mov    $0x0,%edx
 4004f2:      eb 09              jmp     4004fd <sum+0x15>
 4004f4:      48 63 ca          movslq %edx,%rcx
 4004f7:      03 04 8f          add    (%rdi,%rcx,4),%eax
 4004fa:      83 c2 01          add    $0x1,%edx
 4004fd:      39 f2             cmp    %esi,%edx
 4004ff:      7c f3             jl     4004f4 <sum+0xc>
 400501:      f3 c3             repz   retq
```

Using absolute addressing for array: **0x601018** (notice the little endian)

Relocated .text section

```
00000000004004d0 <main>:
 4004d0:      48 83 ec 08          sub    $0x8,%rsp
 4004d4:      be 02 00 00 00      mov    $0x2,%esi
 4004d9:      bf 18 10 60 00      mov    $0x601018,%edi    # %edi = &array
 4004de:      e8 05 00 00 00      callq 4004e8 <sum>      # sum()
4004e3:      48 83 c4 08          add    $0x8,%rsp
 4004e7:      c3                   retq

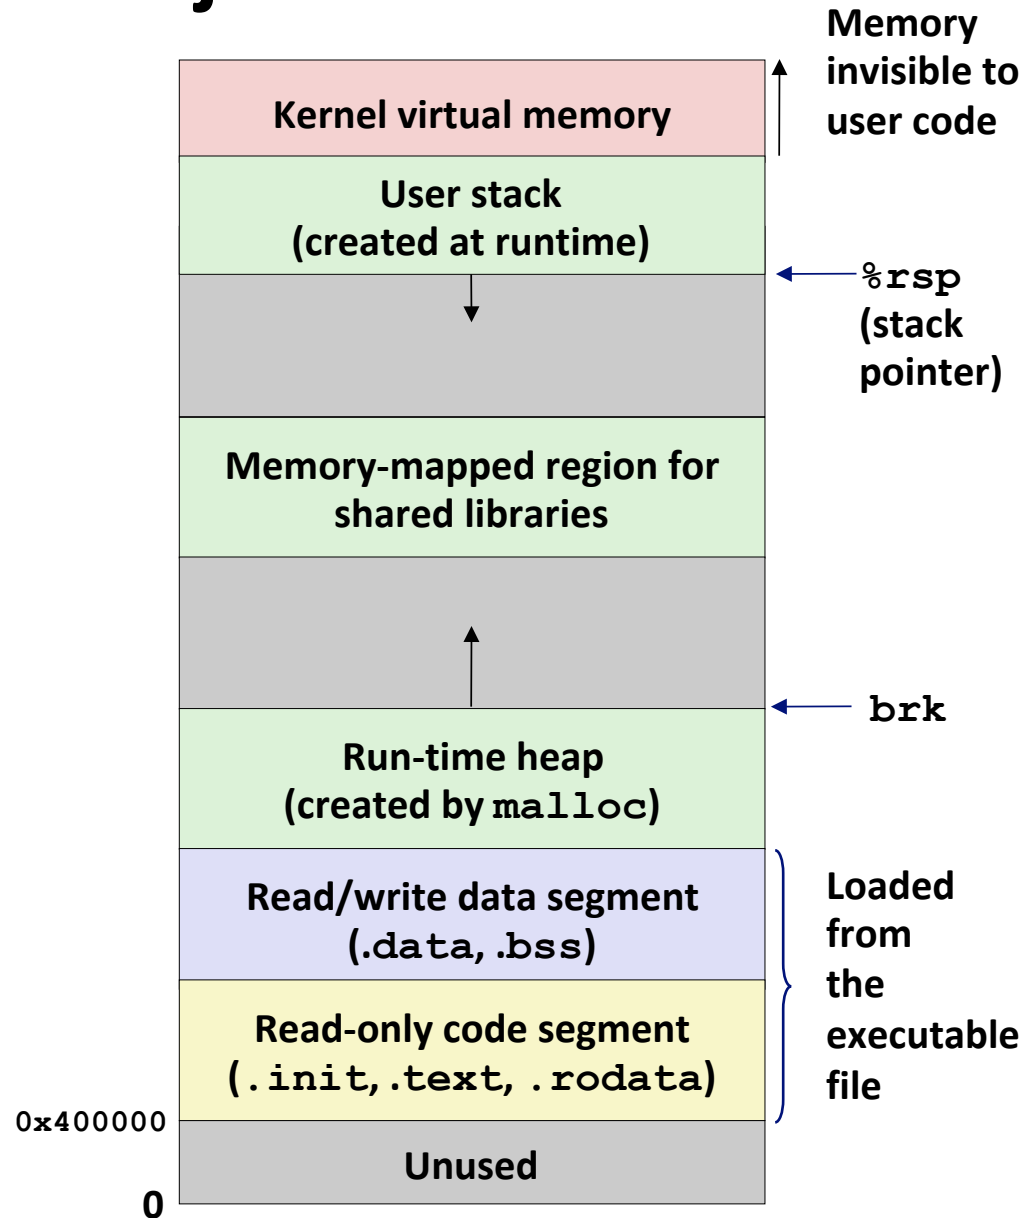
00000000004004e8 <sum>:
4004e8:      b8 00 00 00 00      mov    $0x0,%eax
 4004ed:      ba 00 00 00 00      mov    $0x0,%edx
 4004f2:      eb 09              jmp    4004fd <sum+0x15>
 4004f4:      48 63 ca          movslq %edx,%rcx
 4004f7:      03 04 8f          add    (%rdi,%rcx,4),%eax
 4004fa:      83 c2 01          add    $0x1,%edx
 4004fd:      39 f2              cmp    %esi,%edx
 4004ff:      7c f3              jl     4004f4 <sum+0xc>
 400501:      f3 c3              repz  retq
```

Using PC-relative addressing for sum(): $0x4004e8 = 0x4004e3 + 0x5$

Loading Executable Object Files

Executable Object File

0	ELF header
	Program header table (required for executables)
	.init section
	.text section
	.rodata section
	.data section
	.bss section
	.symtab
	.debug
	.line
	.strtab
	Section header table (required for relocatables)



Recap: What We Learned Thus Far

- **The role of a linker**
 - symbol resolution and relocation
- **Scoping rules for global / static / extern variables**
- **What program entities have a linker symbol associated with it**
 - YES: functions, global variables, static variables
 - NO: local variables, function arguments
- **How linker resolves duplicate symbols**
 - strong vs weak symbols
 - the linker's symbol rules
- **How to use header files to avoid pitfalls in the use of global / static variables.**