



Exceptional Control Flow: Processes and Signals (Cont.)

B&O Readings: 8.4-8.8

CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Safe Signal Handling

- **Proceed with caution: signal-handling is a tricky business**
 - The handler executes concurrently with the main program.
 - Shared data structures can become corrupted.
 - Can have nested signal handlers.
 - Certain library functions are **NOT SAFE** to invoke within a signal handler!

Guidelines for Writing Safe Handlers

- **G0: Keep your handlers as simple as possible**
 - e.g., Set a global flag and return
- **G1: Call only async-signal-safe functions in your handlers**
 - `printf`, `sprintf`, `malloc`, and `exit` are not safe!
- **G2: Save and restore `errno` on entry and exit**
 - So that this or other handlers don't overwrite the value of `errno` used by the application code
- **G3: Protect accesses to shared data structures by temporarily blocking all signals.**
 - To prevent possible corruption
- **G4: Declare global variables as `volatile`**
 - To prevent compiler from storing them in a register
- **G5: Declare global flags (integer) as `sig_atomic_t`**
 - Individual read / write will be atomic (e.g. `flag = 1`, not `flag++`)

Async-Signal-Safety

- Function is *async-signal-safe* if either reentrant (e.g., all variables stored on stack frame, CS:APP3e 12.7.2) or non-interruptible by signals.
- Posix guarantees 117 functions to be async-signal-safe
 - Source: `man 7 signal`
 - Popular functions on the list:
 - `_exit`, `write`, `wait`, `waitpid`, `sleep`, `kill`
 - Popular functions that are **not** on the list:
 - `printf`, `sprintf`, `malloc`, `exit`
 - Unfortunate fact: `write` is the only async-signal-safe output function

Safely Generating Formatted Output

- Use the reentrant SIO (Safe I/O library) from `csapp.c` in your handlers.

- `ssize_t sio_puts(char s[]) /* Put string */`
- `ssize_t sio_putl(long v) /* Put long */`
- `void sio_error(char s[]) /* Put msg & exit */`

```
void sigint_handler(int sig) /* Safe SIGINT handler */
{
    Sio_puts("So you think you can stop the bomb with ctrl-  
c, do you?\n");
    sleep(2);
    Sio_puts("Well...");
    sleep(1);
    Sio_puts("OK. :-)\n");
    _exit(0);
}
```

sigintsafe.c

Why Does the Execution Hang?

```
int ccount = 0;
void child_handler(int sig) {
    int olderrno = errno;
    pid_t pid;
    if ((pid = wait(NULL)) < 0)
        Sio_error("wait error");
    ccount--;
    Sio_puts("Handler reaped child ");
    Sio_putl((long)pid);
    Sio_puts(" \n");
    sleep(1);
    errno = olderrno;
}

void fork14() {
    pid_t pid[N];
    int i;
    ccount = N;
    Signal(SIGCHLD, child_handler);

    for (i = 0; i < N; i++) {
        if ((pid[i] = Fork()) == 0) {
            Sleep(1);
            exit(0); /* Child exits */
        }
    }
    while (ccount > 0) ; /* Parent spins */
}
```

forks.c

```
linux> ./forks 14
Handler reaped child 23240
Handler reaped child 23241
(hang)
```



```

int ccount = 0;
void child_handler(int sig) {
    int olderrno = errno;
    pid_t pid;
    if ((pid = wait(NULL)) < 0)
        Sio_error("wait error");
    ccount--;
    Sio_puts("Handler reaped child ");
    Sio_putl((long)pid);
    Sio_puts(" \n");
    sleep(1);
    errno = olderrno;
}

void fork14() {
    pid_t pid[N];
    int i;
    ccount = N;
    Signal(SIGCHLD, child_handler);

    for (i = 0; i < N; i++) {
        if ((pid[i] = Fork()) == 0) {
            Sleep(1);
            exit(0); /* Child exits */
        }
    }
    while (ccount > 0) ; /* Parent spins */
}

```

forks.c

Buggy Code

- Pending signals are not queued
 - For each signal type, one bit indicates whether or not signal is pending...
 - ...thus at most one pending signal of any particular type.

- You can't use signals to count events!

```

linux> ./forks 14
Handler reaped child 23240
Handler reaped child 23241
(hang)

```



Correct Signal Handling

- **Must wait for all terminated child processes**
 - Put `wait` in a loop to reap all terminated children

```
void child_handler2(int sig)
{
    int olderrno = errno;
    pid_t pid;
    while ((pid = waitpid(-1, NULL, WNOHANG)) > 0) {
        ccount--;
        Sio_puts("Handler reaped child ");
        Sio_putl((long)pid);
        Sio_puts(" \n");
    }
    if (errno != ECHILD)
        Sio_error("wait error");
    errno = olderrno;
}
```

```
linux> ./forks 15
Handler reaped child 23246
Handler reaped child 23247
Handler reaped child 23248
Handler reaped child 23249
Handler reaped child 23250
linux>
```


Blocking and Unblocking Signals

- **Why might we want to block / unblock signals?**
- **Implicit blocking mechanism**
 - Kernel blocks any pending signals of type currently being handled.
 - E.g., A SIGINT handler can't be interrupted by another SIGINT
- **Explicit blocking and unblocking mechanism**
 - `sigprocmask(int how, const sigset_t *set, sigset_t *oldset);`
- **Supporting functions**
 - `sigemptyset` – Create empty set
 - `sigfillset` – Add every signal number to set
 - `sigaddset` – Add signal number to set
 - `sigdelset` – Delete signal number from set

Temporarily Blocking Signals

```
sigset_t mask, prev_mask;

Sigemptyset(&mask);
Sigaddset(&mask, SIGINT);

/* Block SIGINT and save previous blocked set */
Sigprocmask(SIG_BLOCK, &mask, &prev_mask);

⋮    /* Code region that will not be interrupted by SIGINT */

/* Restore previous blocked set, unblocking SIGINT */
Sigprocmask(SIG_SETMASK, &prev_mask, NULL);
```

Example of a Simple Shell

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    char cmdline[MAXLINE];

    Signal(SIGCHLD, handler); /* Install SIGCHLD handler */
    initjobs(); /* Initialize the job list */
    Sigfillset(&mask_all);

    while (1) {
        get_cmd(cmdline, MAXLINE, stdin);
        if ((pid = Fork()) == 0) { /* Child */
            Execve(cmdline, argv, NULL);
        }
        Sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* In parent */
        addjob(pid); /* Add the child to the job list */
        Sigprocmask(SIG_SETMASK, &prev_all, NULL); /* Restore */
    }
}
```

procmask1.c

Example of a Simple Shell

```
void handler(int sig) { /* SIGCHLD handler */
    int olderrno = errno;
    sigset_t mask_all, prev_all;
    pid_t pid;

    Sigfillset(&mask_all);
    while ((pid = waitpid(-1, NULL, 0)) > 0) { /* Reap child */
        Sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        deletejob(pid); /* Delete the child from the job list */
        Sigprocmask(SIG_SETMASK, &prev_all, NULL);
    }
    if (errno != ECHILD)
        Sio_error("waitpid error");
    errno = olderrno;
}
```

procmask1.c

- By blocking the signals, we make sure that accesses to the job data structure is synchronized.

Example of a Simple Shell (Buggy)

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    char cmdline[MAXLINE];

    Sigfillset(&mask_all);
    Signal(SIGCHLD, handler); /* Install SIGCHLD handler */
    initjobs(); /* Initialize the job list */

    while (1) {
        get_cmd(cmdline, MAXLINE, stdin);
        if ((pid = Fork()) == 0) { /* Child */
            Execve(cmdline, argv, NULL);
        }
        Sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* In parent */
        addjob(pid); /* Add the child to the job list */
        Sigprocmask(SIG_SETMASK, &prev_all, NULL); /* Restore */
    }
}
```

procmask1.c

- Ooops. This code is buggy. Where is the bug?



Example of a Simple Shell (Buggy)

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, prev_all;
    char cmdline[MAXLINE];

    Sigfillset(&mask_all);
    Signal(SIGCHLD, handler); /* Install SIGCHLD handler */
    initjobs(); /* Initialize the job list */

    while (1) {
        get_cmd(cmdline, MAXLINE, stdin);
        if ((pid = Fork()) == 0) { /* Child */
            Execve(cmdline, argv, NULL);
        }
        Sigprocmask(SIG_BLOCK, &mask_all, &prev_all); /* In parent */
        addjob(pid); /* Add the child to the job list */
        Sigprocmask(SIG_SETMASK, &prev_all, NULL); /* Restore */
    }
}
```

procmask1.c



- There is a subtle synchronization error: what if the child finishes and the handler invoked before the addjob is called?

Corrected Shell Program without Race

```
int main(int argc, char **argv)
{
    int pid;
    sigset_t mask_all, mask_one, prev_one;

    Sigfillset(&mask_all);
    Sigemptyset(&mask_one);
    Sigaddset(&mask_one, SIGCHLD);
    Signal(SIGCHLD, handler);
    initjobs(); /* Initialize the job list */

    while (1) {
        Sigprocmask(SIG_BLOCK, &mask_one, &prev_one); /* Block SIGCHLD */
        if ((pid = Fork()) == 0) { /* Child process */
            Sigprocmask(SIG_SETMASK, &prev_one, NULL); /* Unblock SIGCHLD */
            Execve("/bin/date", argv, NULL);
        }
        Sigprocmask(SIG_BLOCK, &mask_all, NULL); /* In parent */
        addjob(pid); /* Add the child to the job list */
        Sigprocmask(SIG_SETMASK, &prev_one, NULL);
    }
}
```

See other examples
in textbook 8.5!

Block signals
before fork!

Remember to unblock
signals in child.

What's the Output of This Program?

```
pid_t pid;
int counter = 2;
void handler1(int sig) {
    counter = counter - 1;
    Sio_put(counter);
    fflush(stdout);
    exit(0);
}
int main() {
    Signal(SIGUSR1, handler1);
    printf("%d", counter);
    fflush(stdout);
    if ((pid = fork()) == 0) {
        while(1) {};
    }
    kill(pid, SIGUSR1);
    waitpid(-1, NULL, 0);
    counter = counter + 1;
    printf("%d", counter);
    exit(0);
}
```



What's the Output of This Program?

```
pid_t pid;
int counter = 2;
void handler1(int sig) {
    counter = counter - 1;
    Sio_put(counter);
    fflush(stdout);
    exit(0);
}
int main() {
    Signal(SIGUSR1, handler1);
    printf("%d", counter);
    fflush(stdout);
    if ((pid = fork()) == 0) {
        while(1) {};
    }
    kill(pid, SIGUSR1);
    waitpid(-1, NULL, 0);
    counter = counter + 1;
    printf("%d", counter);
    exit(0);
}
```

Answer: 213





Concurrent Programming and Synchronization

B&O Readings: Loosely based on Chp. 12 (12.1, 12.3-12.5)
CSE 361: Introduction to Systems Software

Instructor:

I-Ting Angelina Lee

Note: these slides were originally created by Markus Püschel at Carnegie Mellon University

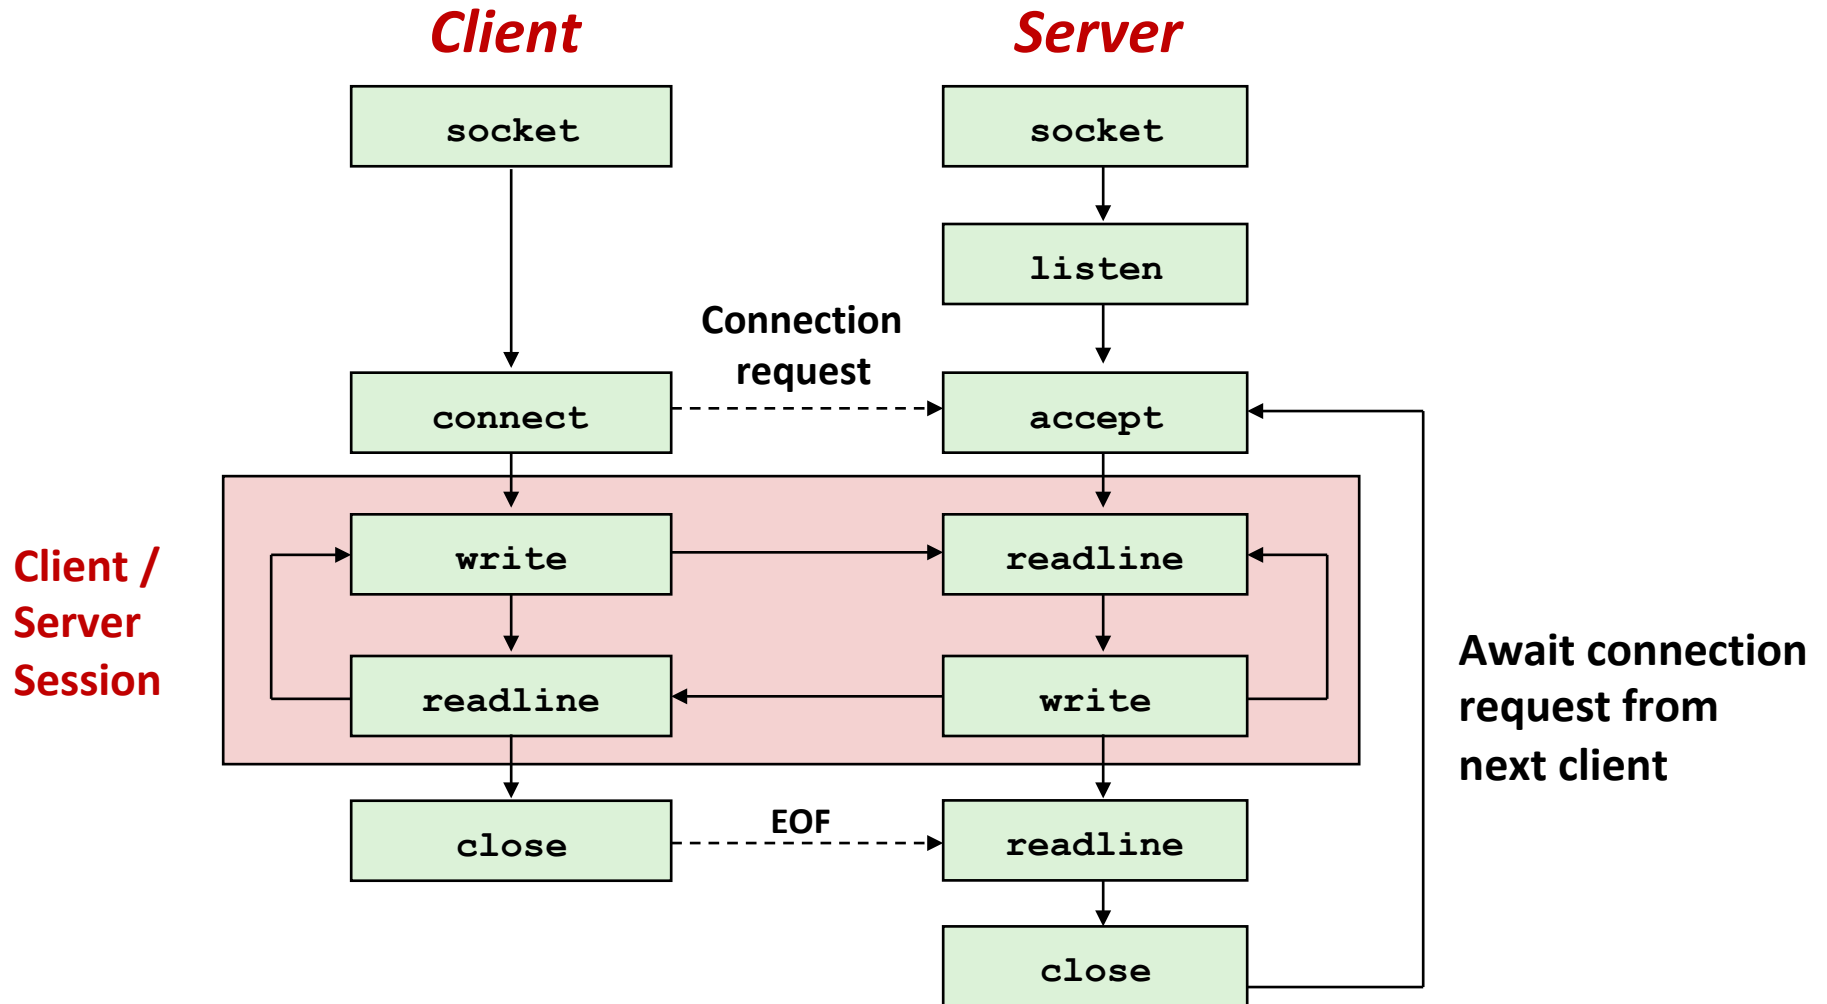
Today

- **Threads: a mechanism for concurrency**
- Synchronization in threaded programs

Different Forms of Concurrency


- Recall: logical control flows are *concurrent* if they overlap in time.
- Ex 1: concurrency among different applications: concurrent processes running different applications on your desktop.
- Ex 2: concurrency within a single application: signal handler and the main logical control of your shell program.

Motivation for Concurrent Programming: Iterative Echo Server



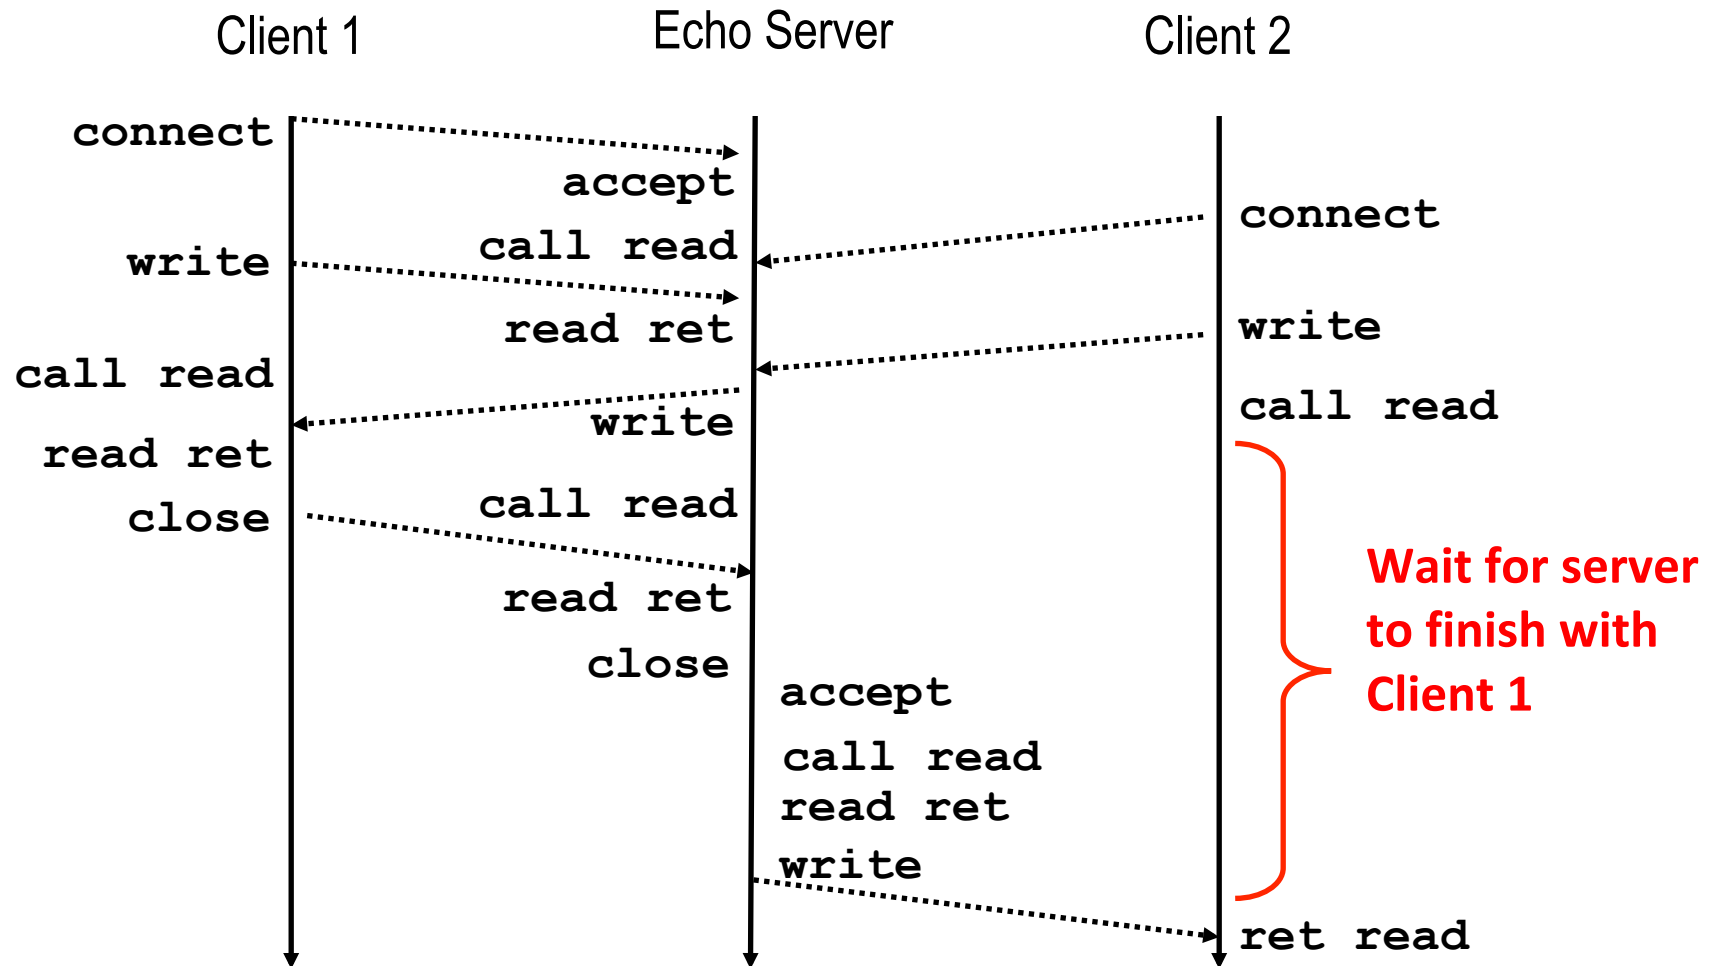
An Iterative Echo Server

- Iterative servers process one request at a time




An Iterative Echo Server

- Iterative servers process one request at a time



- **Iterative servers process one request at a time**



- **Solution:** use a *concurrent server*, which uses multiple concurrent flows to serve multiple clients at the same time.

Approaches for Writing A Concurrent Server

Allow server to handle multiple clients concurrently

1. Process-based

- Kernel automatically interleaves multiple logical flows
- Each flow has its own private address space

2. Event-based (we won't cover this)

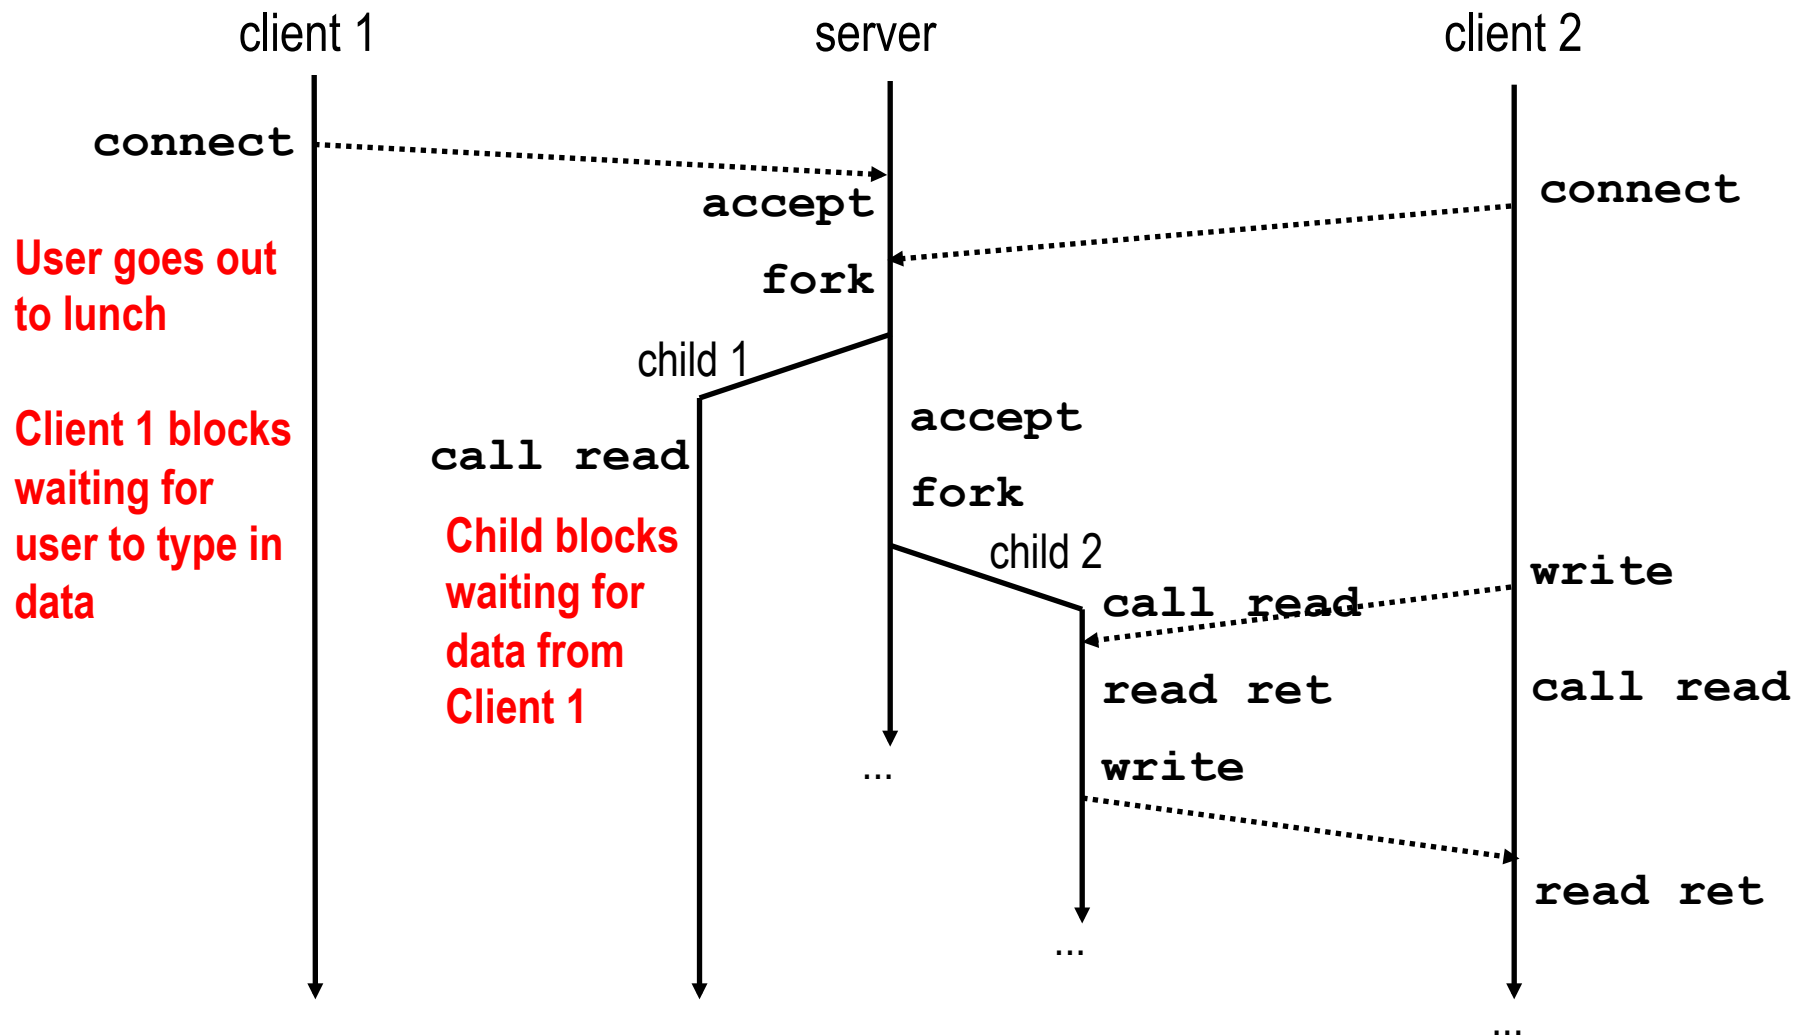
- Programmer manually interleaves multiple logical flows
- All flows share the same address space
- Uses technique called *I/O multiplexing*.

3. Thread-based

- Kernel automatically interleaves multiple logical flows
- Each flow shares the same address space
- Hybrid of of process-based and event-based.

Approach #1: Process-based Servers

- Spawn separate process for each client



Pros and Cons of Process-based Servers

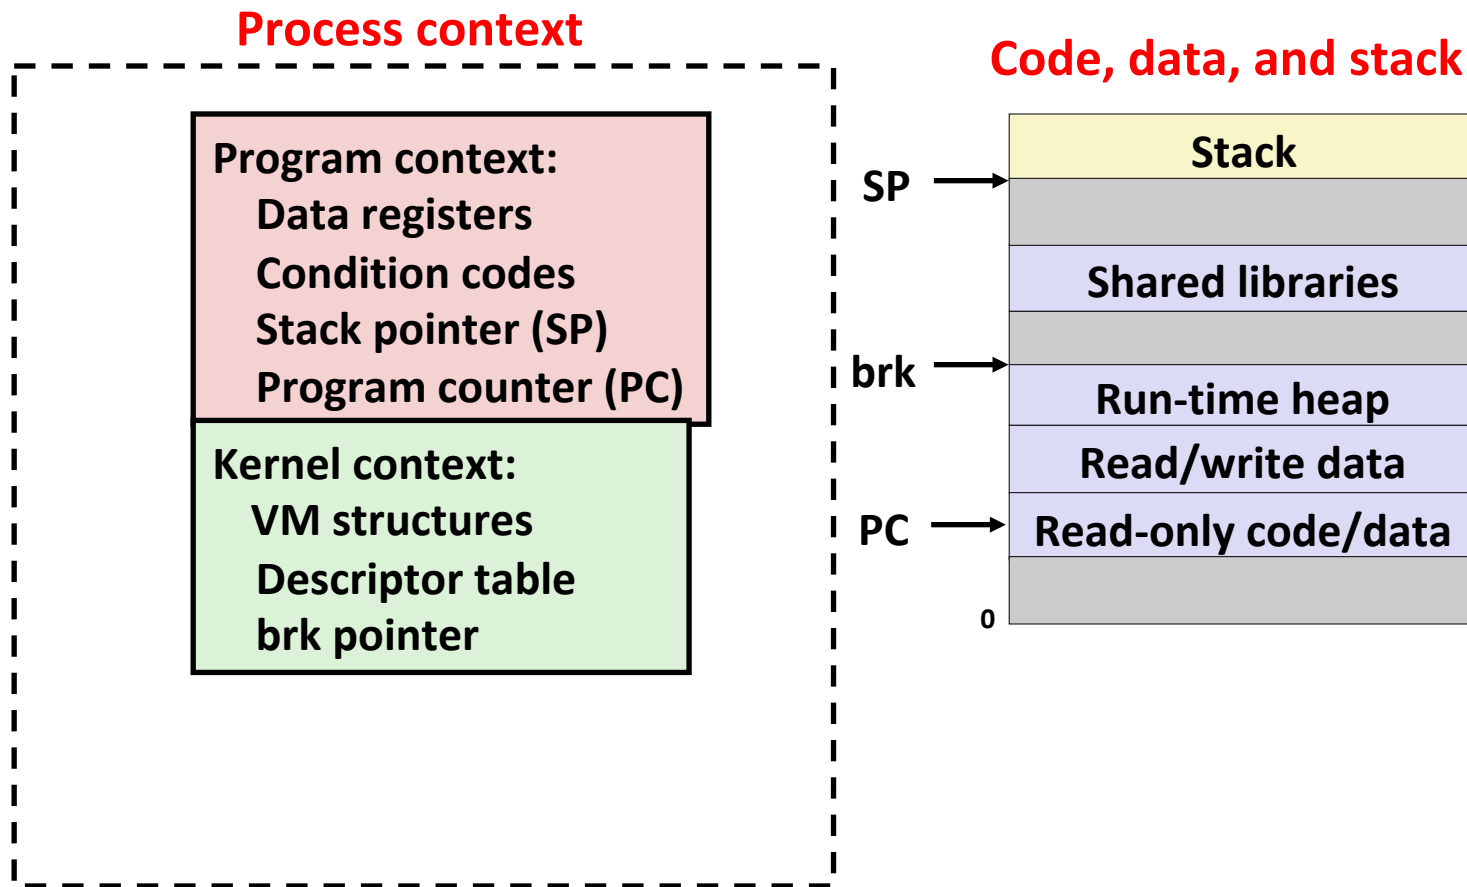
- **+ Handle multiple connections concurrently**
- **+ Clean sharing model**
 - separate address space
- **+ Simple and straightforward**
- **– Additional overhead for process control**
- **– Nontrivial to share data between processes**
 - Requires IPC (interprocess communication) mechanisms
 - Pipes, signals, explicitly share memory via mmap ... etc.

Approach #3: Thread-based Servers

- **Very similar to approach #1 (process-based)**
 - ...but using threads instead of processes

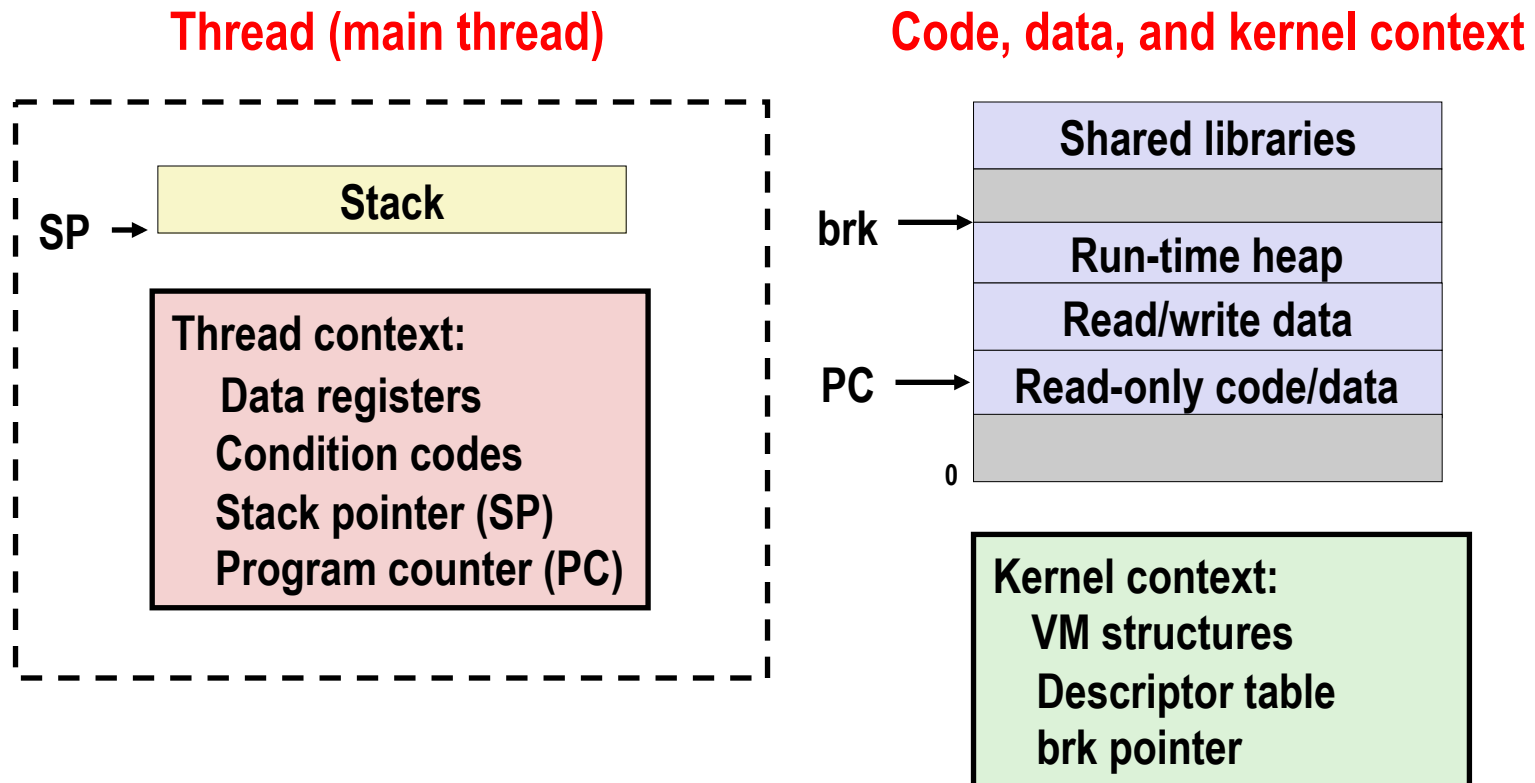
Traditional View of a Process

- Process = process context + code, data, and stack



Alternate View of a Process

- Process = thread + code, data, and kernel context



A Process With Multiple Threads

- **Multiple threads can be associated with a process**
 - Each thread has its own logical control flow
 - Each thread has its own stack for local variables
 - but not protected from other threads
 - Each thread has its own thread id (TID)
 - All threads share the same code, data, and kernel context

Thread 1 (main thread)

Thread 2 (peer thread)


Shared code and data

stack 1

stack 2

Thread 1 context:
Data registers
Condition codes
SP1
PC1

Thread 2 context:
Data registers
Condition codes
SP2
PC2

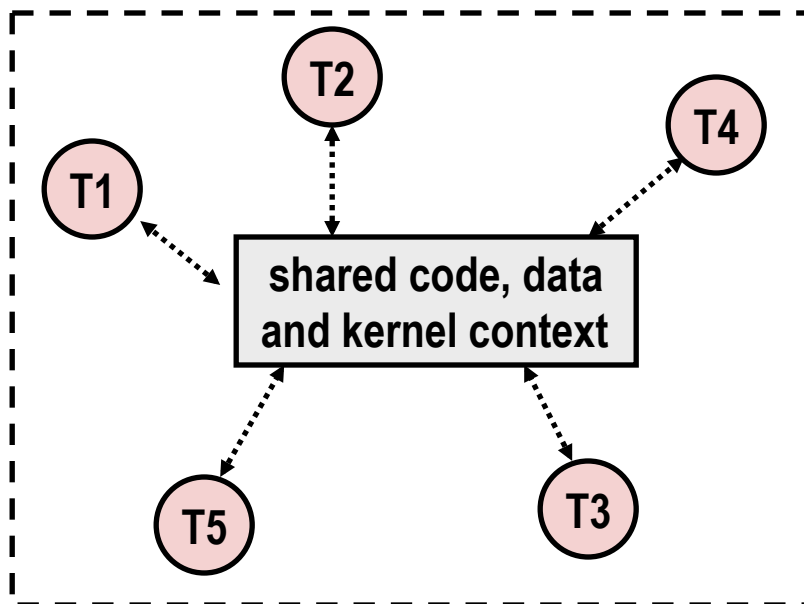


Kernel context:
VM structures
Descriptor table
brk pointer


Logical View of Threads

- **Threads associated with process form a pool of peers**
 - Unlike processes which form a tree hierarchy

Threads associated with process foo



Process hierarchy




Concurrent Threads

- Two threads are *concurrent* if their flows overlap in time
- Otherwise, they are sequential

- **Examples:**


- Concurrent: A & B, A&C
- Sequential: B & C



Concurrent Thread Execution

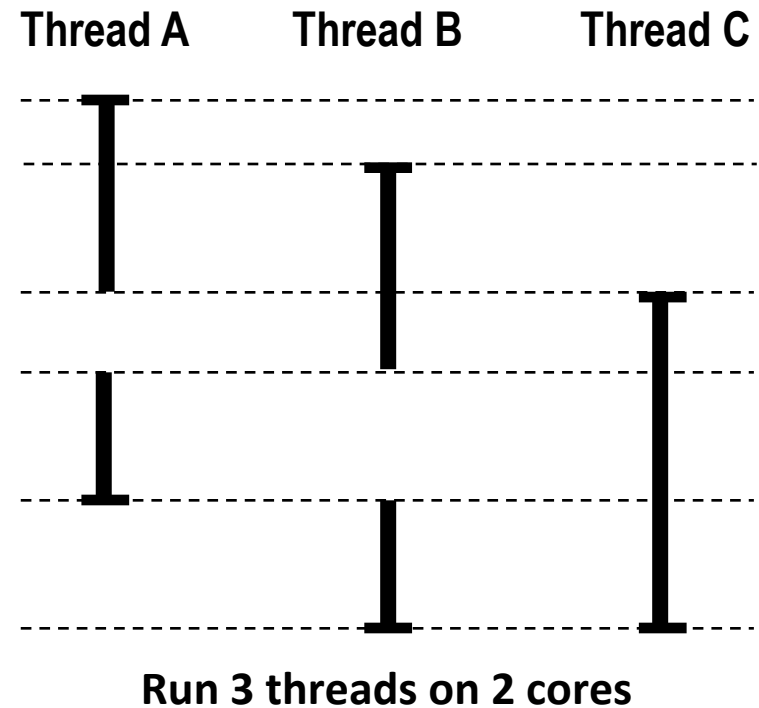
■ Single Core Processor

- Simulate concurrency by time slicing



■ Multi-Core Processor

- Can have true concurrency



Threads vs. Processes

■ How threads and processes are similar

- Each has its own logical control flow
- Each can run concurrently with others (possibly on different cores)
- Each is context switched

■ How threads and processes are different

- Threads share all code and data (except local stacks)
 - Processes (typically) do not
- Threads are somewhat less expensive than processes
 - Process control (creating and reaping) twice as expensive as thread control
 - Linux numbers:
 - ~20K cycles to create and reap a process
 - ~10K cycles (or less) to create and reap a thread

Posix Threads (Pthreads) Interface

- ***Pthreads***: Standard interface for ~60 functions that manipulate threads from C programs
 - Creating and reaping threads
 - `pthread_create()`
 - `pthread_join()`
 - Determining your thread ID
 - `pthread_self()`
 - Terminating threads
 - `pthread_cancel()`
 - `pthread_exit()`
 - `exit()` [terminates all threads], `RET` [terminates current thread]
 - Synchronizing access to shared variables
 - `pthread_mutex_init`
 - `pthread_mutex_[un]lock`

The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```

hello.c

Thread ID

Thread attributes
(usually NULL)

Thread routine

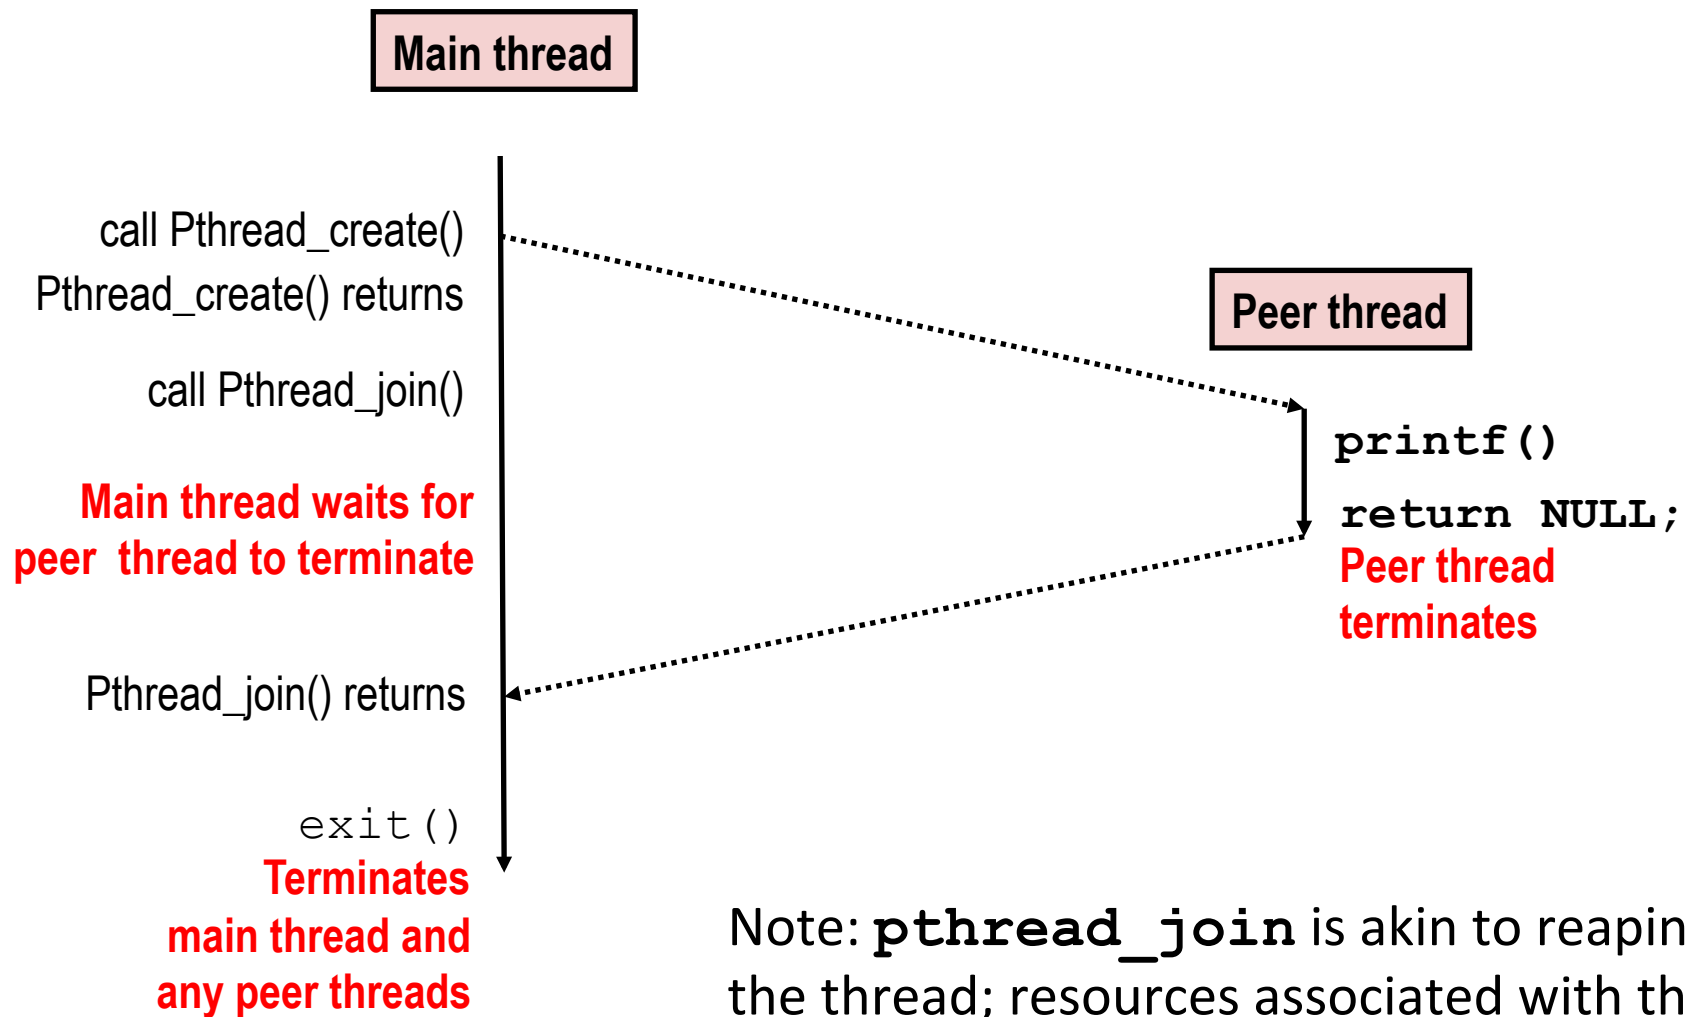
Thread arguments
(void *p)

Return value
(void **p)

```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

hello.c

Execution of Threaded “hello, world”



Note: `pthread_join` is akin to reaping the thread; resources associated with the thread is not freed until it's reaped.

Pros and Cons of Thread-Based Designs

- **+ Easy to share data structures between threads**
 - e.g., logging information, file cache
- **+ Threads are more efficient than processes**
- **– Unintentional sharing can introduce subtle and hard-to-reproduce errors!**
 - The ease with which data can be shared is both the greatest strength and the greatest weakness of threads
 - Hard to know which data shared & which private
 - Hard to detect by testing
 - Probability of bad race outcome very low, but nonzero!
 - All functions called by threads must be thread-safe.

Concurrent Programming is Hard!

- **Need to reason about concurrent events occurring at the same time.**
 - Events interact through shared state.
 - Reasoning about the correctness of your code involves reasoning about all possible interleaving of events.
 - Ex: job list in your shell program
 - The human mind tends to be sequential.
- **Thinking about all possible sequences of events in a computer system is at least error prone and frequently impossible**

Classical problems of concurrent programs

■ ***Races:***

- improper coordination when accessing shared resources concurrently
- outcome depends on arbitrary scheduling decisions
- Example: who gets the last seat on the airplane?


■ ***Deadlock:***

- improper resource allocation prevents forward progress
- Example: traffic gridlock

■ ***Livelock / Starvation / Fairness:***

- external events and/or system scheduling decisions can prevent sub-task progress
- Example: people always jump in front of you in line

Data Race



Classical problems of concurrent programs

■ ***Races:***

- improper coordination when accessing shared resources concurrently
- outcome depends on arbitrary scheduling decisions
- Example: who gets the last seat on the airplane?


■ ***Deadlock:***

- improper resource allocation prevents forward progress
- Example: traffic gridlock

■ ***Livelock / Starvation / Fairness:***

- external events and/or system scheduling decisions can prevent sub-task progress
- Example: people always jump in front of you in line

Deadlock



Classical problems of concurrent programs

■ ***Races:***

- improper coordination when accessing shared resources concurrently
- outcome depends on arbitrary scheduling decisions
- Example: who gets the last seat on the airplane?


■ ***Deadlock:***

- improper resource allocation prevents forward progress
- Example: traffic gridlock

■ ***Livelock / Starvation / Fairness:***

- external events and/or system scheduling decisions can prevent sub-task progress
- Example: people always jump in front of you in line

Starvation



- Yellow must yield to green
- Continuous stream of green cars
- Overall system makes progress, but some individuals wait indefinitely

Classical problems of concurrent programs

■ ***Races:***

- improper coordination when accessing shared resources concurrently
- outcome depends on arbitrary scheduling decisions
- Example: who gets the last seat on the airplane?

■ ***Deadlock:***

- improper resource allocation prevents forward progress
- Example: traffic gridlock

■ ***Livelock / Starvation / Fairness:***

- external events and/or system scheduling decisions can prevent sub-task progress
- Example: people always jump in front of you in line

Many aspects of concurrent programming are beyond the scope of our course...